Stochastic Processes

STAT-150 Lecture Series, Jim Pitman

(version 1.6, updated 10/09/2019)

Fall 2019

University of California, Berkeley

This document is intended to capture lecture material into a referential and well-prepared resource for the cohort of students taking the STAT-150 Stochastic Processes course taught by Jim Pitman. Many thanks to all the student contributors and a special thank-you to John-Michael Laurel and Daniel Suryakusuma.

Contents

1 Stochastic Processes and Markov Chains **1**
- 1.1 Markov Chains . . . 2
- 1.2 Specifying Joint Probabilities . . . 2
- 1.3 Transition Mechanism . . . 3
  - 1.3.1 Absorbing States . . . 4
- 1.4 Constructing the Joint Distribution . . . 4
  - 1.4.1 Notation . . . 5
- 1.5 Simulating a Markov Chain . . . 5
- 1.6 Gambler's Ruin Chain . . . 6

2 Transition Mechanism **7**
- 2.1 Action of a Transition Matrix on a Row Vector . . . 7
  - 2.1.1 Conclusion . . . 8
- 2.2 Action of a Transition Matrix on a Column Vector . . . 8
- 2.3 Two Steps . . . 9
  - 2.3.1 Review: Matrix Multiplication . . . 9
- 2.4 Techniques for Finding P^n for some P . . . 10
- 2.5 First Example . . . 10
- 2.6 Second Example: More Challenging . . . 11

3 Hitting Times, Strong Markov Property, and State Classification **15**
- 3.1 Hitting Times . . . 15
- 3.2 Iterating . . . 16
- 3.3 State Classification . . . 18
  - 3.3.1 Constructing ρ_y . . . 18
- 3.4 Lemma 1.3 . . . 19

4 Exchangeability, Stationary Distributions, and Two State Markov Chains **23**
- 4.1 Sampling without Replacement . . . 23
- 4.2 Exchangeability and Reversibility . . . 24
- 4.3 Stationary Distributions . . . 27
- 4.4 Two State Transition Matrices . . . 30
- 4.5 Two State Transition Diagrams . . . 31

5 Recurrence Classes, x-Blocks, and Limit Theorem **37**
- 5.1 Key Points for Homework . . . 37

5.2 Positive and Null Recurrence . . . 37
5.3 Review: Mean Return Time . . . 38
5.4 Example: Symmetric Random Walk . . . 38
5.4.1 Recurrence versus Transience . . . 39
5.5 Notion of x-Blocks of a Markov chain . . . 39
5.5.1 Example: x-Blocks . . . 40
5.6 Positive Recurrent Chains (P irreducible) . . . 40
5.6.1 Explanation of the Key Formula . . . 41
5.7 Limit Theorem . . . 42
5.8 Review and Audience Questions . . . 42

6 First Step Analysis and Harmonic Equations 43
6.1 Hitting Places . . . 43
6.2 Method of Solution (Durret p.54) . . . 45
6.3 Canonical Example: Gambler's Ruin for a Fair Coin . . . 48
6.3.1 Gambler's Ruin with a Biased Coin . . . 49

7 First Step Analysis Continued 51
7.1 First Step Analysis: Continued . . . 51
7.1.1 Example: Mean Hitting Times . . . 51
7.1.2 Application: Duration of a Fair Game . . . 52
7.1.3 Summarizing our Findings . . . 54
7.2 Conditioning on other variables . . . 55
7.2.1 Runs in Independent Bernoulli(p) Trials . . . 55
7.2.2 Conditioning on the First Zero . . . 56

8 Infinite State Spaces and Probability Generating Functions 59
8.1 Infinite State Spaces . . . 59
8.2 Review of Mathematics: Power Series . . . 59
8.2.1 Binomial Theorem . . . 59
8.3 Probability Generating Functions . . . 60
8.4 Probability Generating Functions and Random Sums . . . 63
8.5 Application: Galton-Watson Branching Process . . . 64

9 Potential Theory (Green Matrices) 69
9.1 Potential Theory (Green Matrices) . . . 69
9.1.1 Example . . . 70
9.2 Escape Probability . . . 73
9.3 More Formulas for Simple Random Walks (SRW) . . . 73
9.4 Green's Matrix for Finite State Space S . . . 74
9.4.1 Return to Gambler's Ruin . . . 76
9.4.2 Conclusion . . . 78

10 The Fundamental Matrix of a Positive Recurrent Chain 79
10.1 Comments on Homework 5 . . . 79
10.2 Renewal Generating Functions . . . 80

10.3 Variance of Sums Over a Markov chain 83
10.3.1 The Mean 84
10.3.2 The Variance 84
10.3.3 The Central Limit Theorem 85
10.4 Further Applications of the Fundamental Matrix 86
10.4.1 Stopping times 86
10.4.2 Occupation Measures for Markov chains 89
10.4.3 Positive recurrent chains: the ergodic theorem 92
10.4.4 Occupation measures for recurrent chains 93
10.4.5 The fundamental matrix of a positive recurrent chain 96
10.4.6 Exercises. 101
10.5 References 102

11 Poisson Processes, Part 1 **103**
11.1 Introduction: Poisson Processes 103
11.2 Sum of Independent Poissons 105
11.3 Poissonization of the Multinomial 105
11.4 Poisson Point Processes (PPP) 107
11.4.1 PPP Strips 108
11.5 Applications 112
11.6 Secret Method 112

12 Poisson Processes, Part 2 **113**
12.1 Theorem 2.10 113
12.2 Generalization to a Stopping Time N 114
12.2.1 Wald's Identities 114
12.3 Poisson Thinning 116
12.3.1 Poisson Thinning for a General Region 117
12.3.2 Poisson Thinning for Two General Regions 118
12.4 General Measures 119

Bibliography **121**

LECTURE 1

Stochastic Processes and Markov Chains

A *stochastic process* is a collection of random variables (R.V.s) that is indexed by a *parameter set* $\mathcal{I}$. We shall use the following notation

$$(X_i, i \in \mathcal{I}) = (X_i)_{i \in \mathcal{I}} \tag{1.1}$$

The parameter set commonly represents a set of times, but can extend to e.g. space or space-time. Underlying (1.1) there is always a *probability measure* $\mathbb{P}$ on some outcome space Ω with $\mathbb{P}$ a function of subsets F of Ω, ranging over a suitable collection of subsets $\mathcal{F}$ called *events*. Use the command `\mathbb{P}` or `\mathds{P}` to produce $\mathbb{P}$ or $\mathbb{P}$ in LaTeX respectively.

In the *canonical* setup, Ω is a *product space*

$$\Omega = \prod_{i \in \mathcal{I}} \mathcal{S}_i$$

where $\mathcal{S}_i$ is a space of values of X_i, and the X_i are just coordinate maps on this product space, and $\mathbb{P}$ is a probability measure on the product space, with $\mathcal{F}$ the product σ-field. So $\omega = (x_i, i \in \mathcal{I}) \in \Omega$ and $X_i(\omega) = x_i$. But

$$(\Omega, \mathcal{F}, \mathbb{P})$$

could be any *probability space*, and each X_i a random variable defined as a function on that space. All of the italicized terms here are standard. Their definitions can be found in `Wikipedia`.

Here's a conversion table on notation used in the course text Durrett. *Essentials of Stochastic Processes* and that of these lecture notes.

Pitman	Durrett	description
$\mathbb{P}$	P	probability measure
P	p	probability transition matrix

1.1 Markov Chains

For a countable *state space* $\mathcal{S}$, for instance $\mathcal{S} = \mathbb{N}_0 := \{0, 1, \dots\}$, we construct a sequence of evolving discrete R.V.s

$$(X_0, X_1, \dots, X_{n-1}, X_n) = (X_i)_{i \in \mathbb{N}_0} \tag{1.2}$$

where

$$\begin{aligned} X_0 &:= \text{ initial state at time } 0 \\ X_1 &:= \text{ state of process after time } 1 \\ &\vdots \\ X_n &:= \text{ state of process after time } n \end{aligned}$$

and call $(X_i)_{i \in \mathbb{N}_0}$ a *Markov Chain* if it satisfies the *Markov Property*.

Markov Property

A stochastic process is a Markov Chain $(X_i)_{i \in \mathcal{S}}$ if it satisfies

$$\mathbb{P}\left(X_{n+1} = x_{n+1} \;\middle|\; \bigcap_{i=0}^{n} \{X_i = x_i\} \right) = \mathbb{P}(X_{n+1} = x_{n+1} \,|\, X_n = x_n)$$

known as the *Markov Property*. In words,

Past and future are conditionally independent given the present.

So given the present value $X_n = x_n$, the past values $X_0, \dots, X_{n-1}$ become irrelevant for predicting values of X_{n+1}.

1.2 Specifying Joint Probabilities

To specify a stochastic process, you must describe the joint distribution of its variables. For example, we know for R.V.s $X_0, X_1, X_2 \in \mathbb{N}_0$

$$\mathbb{P}(X_0 \leq 3, X_1 \leq 5, X_0 \leq 7) = \sum_{x=0}^{3} \sum_{y=0}^{5} \sum_{z=0}^{7} p(x, y, z)$$

So to specify the joint distribution of the three variables X_0, X_1, X_2 it is enough to specify their *joint probability function* $p(x, y, z)$. This must be some non-negative

function which sums to 1 over all triples (x, y, z). Now for any sequence of three R.V.s

$$\mathbb{P}(X_0 = x, X_1 = y, X_2 = z) = \\ \mathbb{P}(X_0 = x)\mathbb{P}(X_1 = y \,|\, X_0 = x)\mathbb{P}(X_2 = z \,|\, X_1 = y, X_0 = x)$$

For a Markov Chain, this reduces to

$$\mathbb{P}(X_0 = x, X_1 = y, X_2 = z) = \\ \mathbb{P}(X_0 = x)\mathbb{P}(X_1 = y \,|\, X_0 = x)\mathbb{P}(X_2 = z \,|\, X_1 = y)$$

Now *transition matrices* P_1 and P_2 can be defined by

$$\mathbb{P}(X_1 = y \,|\, X_0 = x) = P_1(x, y)$$
$$\mathbb{P}(X_2 = z \,|\, X_1 = y) = P_2(y, z)$$

Then

$$\mathbb{P}(X_0 = x, X_1 = y, X_2 = z) = \mathbb{P}(X_0 = x)P_1(x, y)P_2(y, z) \tag{1.3}$$

Most commonly we assume that the chain has *homogeneous* transition probabilities. That means $P_1 = P_2 = P$ for a single transition matrix P. We deal with this idea formally in the next section,

1.3 Transition Mechanism

For finite $\mathcal{S}$, we can build a *transition matrix*

$$P = P(x, y) \tag{1.4}$$

a "set of rules" or "mechanism" for moving between different states in $\mathcal{S}$. Rules of probability imply that (1.4) is a *stochastic matrix*.

Stochastic Matrix

A *stochastic matrix* is a non-negative matrix with all row sums equal to 1. With the usual convention of x indexing rows and y indexing columns, we say P is stochastic if it satisfies

$$P(x, y) \geq 0 \quad \text{and} \quad \sum_y P(x, y) = 1$$

It should make intuitive sense that for a fixed row, summing its elements yields one. Given the present state x, you're bound to go somewhere.

For now, our Markov chains will possess *homogeneous transition probabilities*. All that means is we use the same matrix P at each step in time. To make this more

concrete, observe for a Markov Chain

$$\begin{aligned}&\mathbb{P}(X_0 = x_0, X_1 = x_1, X_2 = x_2, X_3 = x_3)\\ &\qquad = \mathbb{P}(X_0 = x_0)P_1(x_0, x_1)P_2(x_1, x_2)P_3(x_2, x_3)\end{aligned}$$

and by time homogeneity $P_1, P_2, P_3 = P$ and we have

$$\begin{aligned}&\mathbb{P}(X_0 = x_0, X_1 = x_1, X_2 = x_2, X_3 = x_3)\\ &\qquad = \mathbb{P}(X_0 = x_0)P(x_0, x_1)P(x_1, x_2)P(x_2, x_3)\end{aligned}$$

1.3.1 Absorbing States

When $P(x, x) = 1$, we say that x is an *absorbing state*. If you start in an absorbing state, you are sure to be there next step, and there again after two steps, and so on. Put another way, once you arrive at an absorbing state, you never leave it.

1.4 Constructing the Joint Distribution

To get the chain going, we initialize the process with an assigned initial distribution λ for X_0. That is

$$\mathbb{P}(X_0 = x_0) = \lambda(x_0)$$

Again, rules of probability require λ is a probability distribution on the state space.

$$\lambda(x_0) \geq 0 \quad \text{and} \quad \sum_{x_0} \lambda(x_0) = 1$$

We finally have everything we need to completely specify the joint distribution of a Markov chain with homogeneous transition probabilities.

Prescription for the Joint Distribution

Let the Markov chain $(X_i)_{i \in \mathcal{S}}$ have a finite state space $\mathcal{S}$, assigned initial distribution λ, and transition probability matrix P. Then for sequences of length $n + 1$, the joint distribution

$$(X_0 = x_0, X_1 = x_1, \ldots, X_{n-1} = x_{n-1}, X_n = x_n)$$

has the following prescription

$$\mathbb{P}\left(\bigcap_{i=0}^{n} \{X_i = x_i\}\right) = \lambda(x_0) \prod_{i=0}^{n-1} P(x_i, x_{i+1}) \tag{1.5}$$

This construction is a proper assignment of a joint distribution, according to the rules of probability. One can check this by verifying 1) all the joint probabilities are

non-negative (which is trivial) and 2) all the probabilities sum to one, that is

$$\sum_{x_0}\sum_{x_1}\cdots\sum_{x_{n-1}}\sum_{x_n}\lambda(x_0)\prod_{i=0}^{n-1}P(x_i,x_{i+1})=1 \tag{1.6}$$

which one can show using mathematical induction on n. e.g. assuming it is true for $n-1$ instead of n, in going from $n-1$ to n there is just one more summation, which simplifies as it should using row sums of the transition matrix equal 1.

1.4.1 Notation

$\mathbb{P}_\lambda$ is used to show that $\mathbb{P}$ makes X_0 have distribution λ.

$\mathbb{P}_x$ is used to show that $\mathbb{P}$ makes $X_0 = x$, i.e. $\lambda(x) = 1$ and $\lambda(x_0) = 0$ for $x_0 \neq x$. So under $\mathbb{P}_x$ the chain starts in state x.

1.5 Simulating a Markov Chain

We can simulate a Markov chain $(X_i)_{i\in\mathbb{N}}$, with a supply of uniform R.V.s

$$U_0, U_1, \ldots \sim \textbf{Uniform}(0,1)$$

For the initial state X_0 to have distribution λ, that is $X_0 \sim \lambda$, let

$$X_0 = \begin{cases} 0, & \text{if } 0 \le U_0 < \lambda(0) \\ 1, & \text{if } \lambda(0) \le U_0 < \lambda(0)+\lambda(1)] \\ 2, & \text{if } \lambda(0)+\lambda(1) \le U_0 < \lambda(0)+\lambda(1)+\lambda(2)] \\ & \vdots \end{cases}$$

Now, if $X_0 = x_0$, define

$$X_1 = \begin{cases} 0, & \text{if } 0 \le U_1 < P(x_0,0) \\ 1, & \text{if } P(x_0,0) \le U_1 < P(x_0,0)+P(x_0,1)] \\ 2, & \text{if } P(x_0,0)+P(x_0,1) \le U_1 < P(x_0,0)+P(x_0,1)+P(x_0,2)] \\ & \vdots \end{cases}$$

and so on. Hence, given $X_0 = x_0$ and $X_1 = x_0$, we create intervals using the elements $P(x_1, \cdot)$. It is easy to implement this simulation using a language like `R`, `Python`, etc.

Often the row $P(x, \cdot)$ is a standard distribution, e.g. uniform or binomial or Poisson or geometric with parameters depending on x. Then there are built in packages for generating such variables which can be used instead of the crudest scheme indicated above.

1.6 Gambler's Ruin Chain

A gambler has a fortune of $\$a$, where $0 \le a \le N$. At each play the gambler wins a \$1 with probability p and loses \$1 with probability $q = 1 - p$. The gambler plays until $X_n = N$ (quitting with a gain) or until $X_n = 0$ (ruined). Here

$$X_n := \text{ the gambler's capital after } n \text{ plays}$$

The transition probabilities for the edge cases are

$$P(0,0) = 1 \quad \text{and} \quad P(N,N) = 1$$

So states 0 and N are *absorbing*. Reference. [1] Durrett's, *Essentials of Stochastic Processes* Section 1.1. Here's a depiction of the gambler's chain.

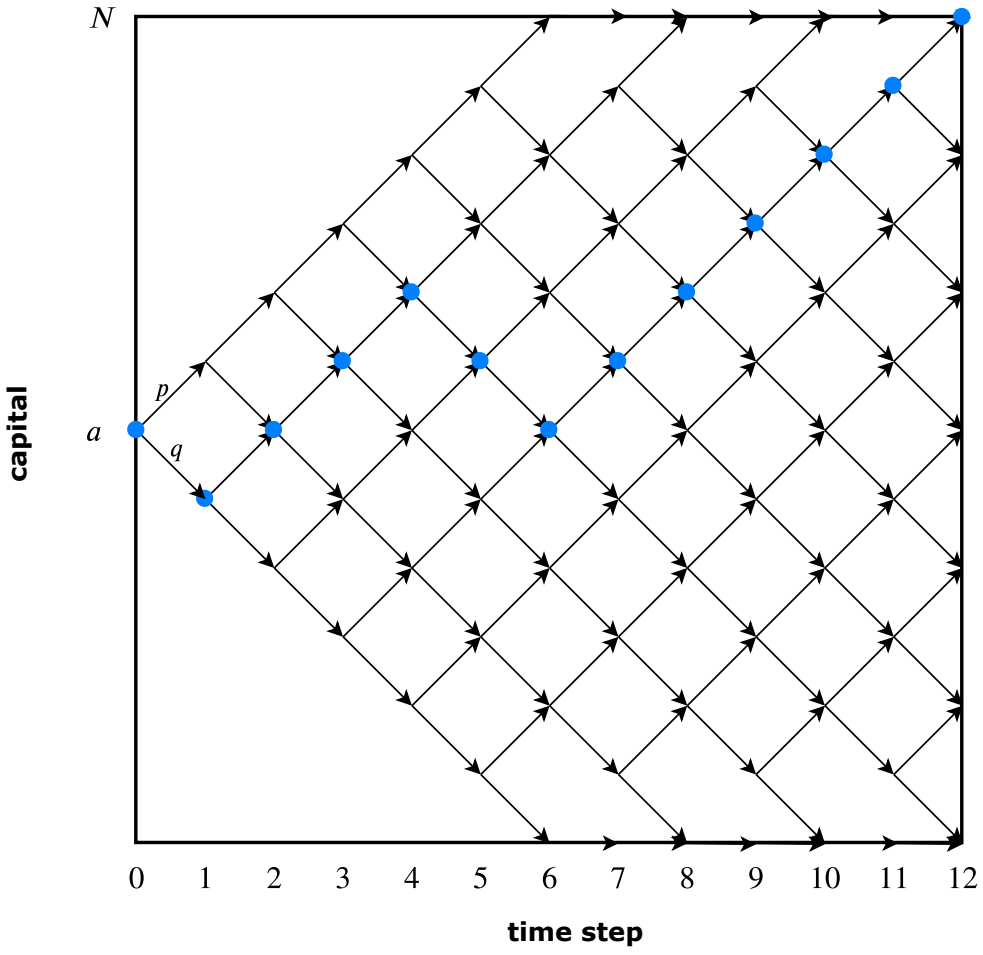

LECTURE 2

Transition Mechanism

Pitman reminds us that Wikipedia serves as a valuable resource for clarifying most basic definitions in this course.
Recall from Lecture 1 we worked with a *transition matrix* P with rows x and columns y. The xth row and yth column entry is $P(x, y)$. All entries are non-negative. All row sums are 1.
For the first step in the Markov chain, we have:

$$P(x, y) = \mathbb{P}(X_1 = y \,|\, X_0 = x).$$

With many steps, and homogeneous transition probabilities, also

$$P(x, y) = \mathbb{P}(X_{n+1} = y \,|\, X_n = x).$$

Pitman notes that the first problem on homework 1 is very instructional, which gets us to think about what exactly is the Markov property.

2.1 Action of a Transition Matrix on a Row Vector

Take an initial distribution $\lambda(x) = \mathbb{P}(X_0 = x)$. If we write $P(x, \cdot)$, we're taking the row of numbers in the matrix. With N states we can simply consider sequences of length N rather than N-dimensional space. To ensure we really know what's going on here, consider 2 steps (indexed 0 and 1). What is the distribution of X_0? Trivially, it's λ. Now what is the distribution of X_1? We need to do a little more. We don't know how we started, and we want to think of all the ways we could have ended up at our final state X_1.
To do this, we use the **law of total probability**, which gives:

$$\boxed{\mathbb{P}(X_1 = y) = \sum_{x \in S} \mathbb{P}(X_0 = x, X_1 = y).}$$

Now it just takes a little bit of calculation to go forward. Conditioning on X_0 (turning a joint probablity into a marginal for the first and a conditional given the

first) gives:

$$\begin{aligned}
\mathbb{P}(X_1 = y) &= \sum_{x \in S} \mathbb{P}(X_0 = x, X_1 = y) \\
&= \sum_{x \in S} \mathbb{P}(X_0 = x) \cdot \mathbb{P}(X_1 = y \,|\, X_0 = x) \\
\mathbb{P}(X_1 = y) &= \sum_{x \in S} \lambda(x) \underbrace{P(x, y)}_{\text{matrix entry}} \\
&= (\lambda P)(y) \text{ or equivalently, } = (\lambda P)_y
\end{aligned}$$

To have this fit with matrix multiplication, we must take $\lambda(x)$ to be a ROW VECTOR. Back in our picture going from one step to the next of a Markov chain, we use x the (nth state) to index the row of the matrix and y (the $n+1$th state) to index the column of the matrix $P(x, y)$.

2.1.1 Conclusion

There is a happy coincidence between the rules of probability and the rules of matrices, which implies that if a Markov Chain has $X_0 \sim \lambda$ (meaning random variable X_0 has distribution λ), then at the next step we have the following distribution

$$re\boxed{X_1 \sim \lambda P}$$

where argument y is hidden. If you evaluate the row vector λP at entry y, you get $(\lambda P)_y = \mathbb{P}(X_1 = y)$. Although this may not be terribly exciting, Pitman notes this is fundamental and important to understand the connection between linear algebra and rules of matrices with probability. We will maintain and strengthen this connection throughout the course.

2.2 Action of a Transition Matrix on a Column Vector

Suppose f is a function on S. Think of it as a **reward** in that if $X_1 = x$, then you get $\$f(x)$ (random monetary reward $f(X_1)$ where $X_1 \in S$ as an abstract object; these can be partitions or permutations or something very abstract, not necessarily numerical). Pitman notes some applications of Markov chains to Google's PageRank with a very big state space of web pages. Without being scared about the potential size of the **state space**, we open to some abstraction in our immediate example. Consider the Markov chain step from X_0 to X_1 and the conditional expectation:

$$\mathbb{E}\big[f(X_1) \,|\, X_0 = x\big] = \sum_y \underbrace{P(x, y)}_{\text{matrix}} \underbrace{f(y)}_{\text{col.vec.}}$$

where we could make some concrete financial definitions to apply our abstract problem if we wish.

Starting at state x, we move to the next state according to the row $P(x, \cdot)$. Recognize this as a matrix operation and we have, for the above:

$$\mathbb{E}\left(f(X_1) \,|\, X_0 = x\right) = (Pf)(x)$$

Remark: The function or column vector f can be signed (there is no difficulty if we are losing money as opposed to gaining); it is more difficult to interpret the action on a signed row vector λ. But easy to interpret λP for a probability measure λ.

2.3 Two Steps

Now consider two steps in time

$$X_0 \xrightarrow{P} X_1 \xrightarrow{Q} X_2$$

Assume the Markov property. Now let's discuss the probability of X_2, knowing $X_0 = x$. That is,

$$\mathbb{P}(X_2 = z \,|\, X_0 = x)$$

where we have some mystery intermediate X_1. The row out of the matrix which we use for the intermediate is random.
We condition upon what we don't know in order to reach a solution. It should become instinctive to us soon to do such a thing: condition on X_1. This gives

$$\begin{aligned}\mathbb{P}(X_2 = z \,|\, X_0 = x) &= \sum_y \mathbb{P}\left(X_1 = y, X_2 = z \,|\, X_0 = x\right) \\ &= \sum_y P(x, y)Q(y, z)\end{aligned}$$

where in the homogeneous case, $P = Q$; however, here we prefer the more clear notation as above. Pitman jokes that generations of mathematicians developed a surprisingly compact form for this, which is matrix multiplication. If P, Q are matrices, this is simply:

$$\sum_y P(x, y)Q(y, z) = PQ(x, z)$$

where we take the x, zth element of the resulting matrix PQ.

2.3.1 Review: Matrix Multiplication

Assuming P, Q, R are $S \times S$ matrices, where S is the label set of indices, then Pitman notes that indeed,

$$PQR := (PQ)R = P(QR)$$

via the associativity of matrix multiplication. This is true for all finite matrices. As a side comment, this is also true for infinite matrices, provided they are nonnegative ≥ 0 (of course, if we have signed things, then summing infinite arrays in different

orders may cause issues). For our purposes, all our entries are nonnegative, so we have no issues.
Now, recall that typically, matrix multiplication is not commutative; that is,

$$PQ \neq QP$$

However, one easy (and highly relevant) case:
If our chain has homogeneous transition probabilities: P, P, P, P. If Pitman asks us what is the probability that $X_n = z$ if we knew $X_0 = x$, then we iterate what we found for 2 steps

$$\mathbb{P}(X_n = z \,|\, X_0 = x) = \underbrace{PPP\cdots P}_{n \text{ times}}(x, z) =: \boxed{P^n(x, z)}$$

Again, Pitman notes we have a very happy 'coincidink' (coincidence): If we take an n-step transition matrix (TM) of a markov chain (MC) with homogeneous probabilities P, this is equivalent to simply P^n, the nth power of matrix P. We can bash this out with computers, but Pitman notes there are techniques of diagonalizing and spectral theory to perform high powers of matrices. Realize that every technique here has an **immediate application** to Markov chains (with very many steps).
Note the *Chapman-Kolmogorov equations*

$$P^{m+n} = P^m P^n = P^n P^m$$

So powers of a single matrix do commute. These equations are easily justified either by algebra, or by probabilistic reasoning. See text Section 1.2 for details of the probabilistic reasoning.

2.4 Techniques for Finding P^n for some P

Pitman wants to warn us that these ideas will be coming and eventually will be useful for this course. Especially, we consider matrices P related to sums of independent random variables. The most basic example is a **Random Walk** on $\mathbb{N}_0 := \{0, 1, 2, \dots\}$.
In this problem one usually writes S_n for the state instead of X_n. Our basic X has $X_0, X_1, X_2, \dots$ i.i.d. according to some P. This is truly a trivial MC. All rows of P are equal to some $p = (p_0, p_1, \dots)$ We consider:

$$S_n = X_0 + X_1 + \cdots + X_n = \text{ cumulated winnings in a gambling game}$$

(Ignore costs or losses for convenience, so natural state space of S_n is $\mathbb{N}_0$).

2.5 First Example

Let $p \sim$ **Bernoulli**(p) where values $0, 1$ have probabilities q, p, respectively. Then $S_n := X_0 + X_1 + \cdots + X_n$.

This admits the following (infinite) matrix:

$$\begin{bmatrix} * & 0 & 1 & 2 & 3 & 4 & 5 & \cdots \\ 0 & q & p & 0 & 0 & 0 & 0 & \cdots \\ 1 & 0 & q & p & 0 & 0 & 0 & \cdots \\ 2 & 0 & 0 & q & p & 0 & 0 & \cdots \\ 3 & 0 & 0 & 0 & q & p & 0 & \cdots \\ 4 & 0 & 0 & 0 & 0 & q & p & \cdots \\ 5 & 0 & 0 & 0 & 0 & 0 & q & \cdots \\ \vdots & & & & & & & \end{bmatrix}$$

Because we can only win \$1 at a time, we fill in the first row trivially.
Pitman asks us now to write down a formula for P^n. As a hint, he says to start with the top row.

$$P^n(0,k) = \mathbb{P}(\underbrace{X_1 + \cdots + X_n}_{n \text{ iid } \textbf{Bernoulli}(()p)} = k)$$

If this doesn't come quickly to us (the answer is trivial according to Pitman), then we should re-visit our 134 probability text (which for me happens to be by Pitman). To find P^n, we note $n = 1$ is known, so taking $n = 2$ for a state space of X_0, X_1, X_2 gives the probabilities:

$$\begin{aligned} P^2(0,0) &= q^2 \\ P^2(0,2) &= p^2 \\ P^2(0,1) &= 2pq, \end{aligned}$$

and this is the familiar **binomial distribution**. Our formula is:

$$\begin{aligned} P^n(0,k) &= \mathbb{P}(\underbrace{X_1 + \cdots + X_n}_{n \text{ iid } bern(p)} = k) \\ &= \boxed{\binom{n}{k} p^k q^{n-k}}. \end{aligned}$$

Now being at an initial fortune i, we have:

$$P^n(i,k) = \binom{n}{k-i} p^{k-i} q^{n-(k-i)}.$$

2.6 Second Example: More Challenging

Now consider the same problem, same setup, but now with $X_1, X_2, \dots$ are i.i.d. with the distribution $(p_0, p_1, p_2, \dots)$ (perhaps all strictly positive) instead of $(q, p, 0, 0, 0, \dots)$. We are interested in the distribution of our Markov Chain after n steps. Taking the same method, it's enough to discuss the distribution of $S_n = X_1 + \cdots + X_n$, because we just shift by i to $S_0 = i$.

Our matrix is now:

$$\begin{bmatrix} * & 0 & 1 & 2 & 3 & 4 & 5 & \cdots \\ 0 & p_0 & p_1 & p_2 & \cdots & \cdots & \cdots & \cdots \\ 1 & 0 & p_0 & p_1 & p_2 & \cdots & \cdots & \cdots \\ 2 & 0 & 0 & p_0 & p_1 & p_2 & \cdots & \cdots \\ 3 & 0 & 0 & 0 & p_0 & p_1 & p_2 & \cdots \\ 4 & 0 & 0 & 0 & 0 & p_0 & p_1 & \cdots \\ 5 & 0 & 0 & 0 & 0 & 0 & p_0 & \cdots \\ \vdots & & & & & & & \end{bmatrix}$$

Again, to get closer to induction, we take $n = 1$ to $n = 2$ steps (with $S_0 = 0$). In matrix notation, we have:

$$\mathbb{P}_0(S_2 = k) = \sum_{j=0}^{k} P(0, j)P(j, k),$$

where we stop at k because we are only adding nonnegative variables. And in probability notation, where we start with j and need to get to k (so we move $k - j$) we have:

$$\mathbb{P}_0(S_2 = k) = \sum_{j=0}^{k} \mathbb{P}(X = j)\mathbb{P}(X = k - j),$$

and either way (of the above two), this ends up being equal to:

$$\mathbb{P}_0(S_2 = k) = \sum_{j=0}^{k} P_j P_{k-j}.$$

So in conclusion, we have found

$$P^2(0, k) = \sum_{j=0}^{k} P_j P_{k-j}$$

where we want to know the name of this operation: **discrete convolution** (so that we know what to look up!). This gets us from a distribution of random variables to the distribution of their sum. There is a "brilliant idea" (as given by Pitman) Consider the power series (of the generating function) $G(z) := \sum_{n=0}^{\infty} p_n z^n$, where taking

$$\left(p_0 + p_1 z + p_2 z^2 + \cdots\right)\left(p_0 + p_1 z + p_2 z^2 + \cdots\right)$$

yields that $\sum_{j=0}^{k} P_j P_{k-j}$ is simply the coefficient of a particular term. Pitman gives us a slick notation:

$$\begin{aligned} P^2(0, k) &= \sum_{j=0}^{k} P_j P_{k-j} \\ &= \left[z^k\right] \underbrace{\left(\sum_{n=0}^{\infty} p_n z^n\right)^2}, \end{aligned}$$

which is just the coefficient of z^k in the under-braced expression.
Repeating this convolution, we move forward from $n = 2$, by induction on n if you want to be careful:

$$P^n(0,k) = [z^k][G(z)]^n$$

Example: Pitman asks us to evaluate via Wolfram Alpha dice rolls $(p_0, p_1, \dots) = \left(\underbrace{\frac{1}{6}, \frac{1}{6}, \dots, \frac{1}{6}}_{6}, 0, \dots\right)$ and want to find: $P^4(0,5)$ for dice rolls $= \mathbb{P}(S_4 = 5)$.
We have

$$\begin{aligned}\left[\frac{1}{6}(z+z^2+z^3+z^4+z^5+z^6)\right]^4 = \frac{1}{6^4}(&z^{24}+4z^{23}+10z^{22}+20z^{21}\\ &+35z^{20}+56z^{19}+80z^{18}\\ &+104z^{17}+125z^{16}+140z^{15}\\ &+146z^{14}+140z^{13}+125z^{12}\\ &+104z^{11}+80z^{10}+56z^9+35z^8\\ &+20z^7+10z^6+\underbrace{4z^5}+z^4)\end{aligned}$$

which implies

$$P^4(0,5) = \frac{4}{6^4}$$

where we took the coefficient of the under-braced term (power of 5). Pitman credits the invento of this method, Laplace. This is unusually simple but demonstrates the general method. Of course, $P^4(0,5) = \frac{4}{6^4}$ is rather trivial because you can count the number of dice patterns on one hand. But the evaluations of $P^4(0,k)$ for all $4 \le k \le 24$ above are not so trivial. This method can be used to prove all the familiar properties of sums of independent discrete variables, e.g. sums of Poissons are Poisson. You should try it for that purpose.

LECTURE 3

Hitting Times, Strong Markov Property, and State Classification

3.1 Hitting Times

This discussion follows quite closely §1.3 of the text. See text for further developments and details. Consider a Markov Chain with fixed transition matrix P and state space $\mathcal{S}$. Consider states $x, y \in \mathcal{S}$ [1]. We are interested in the *first hitting time* or *first passage time*

$$T_B := \min\{n \geq 1 : X_n \in B\}$$

for some target set of states B. In words, the first time at or after time 1 that the chain hits the set of states B. An immediate pedantic issue with is what if the chain never reaches B, that is $X_n \notin B \ \forall n$? In this case, we need to make the following (very useful) convention

$$\min\{\varnothing\} = \inf\{\varnothing\} := \infty$$

where ∞ is a conventional element assumed to be greater than every positive integer.

Strong Markov Property (SMP)

Start with $X_0, X_1, X_2, \dots$ which is a Markov chain with transition matrix P, and any initial distribution for X_0. Conditionally given $T_B = n < \infty$ and $X_n = y \in B$, the following process

$$(X_n, X_{n+1}, X_{n+1}, \dots)$$

is a copy of the original Markov Chain with transition matrix P conditioned to start in state y.

[1] sometimes $i, j \in \mathcal{S}$

In particular, the distribution of $X_{n+1} \,|\, X_n = y$ is $P(y, \cdot)$, that is, for any state z

$$\mathbb{P}(X_{n+1} = z \,|\, T_y = n, X_n = y) = P(y, z), \text{ also}$$
$$\mathbb{P}(X_{n+1} = z, X_{n+2} = w \,|\, T_y = n, X_n = y) = P(y, z)P(z, w)$$

and so on, an infinite list of equations.

Proof. See Durett page 14. □

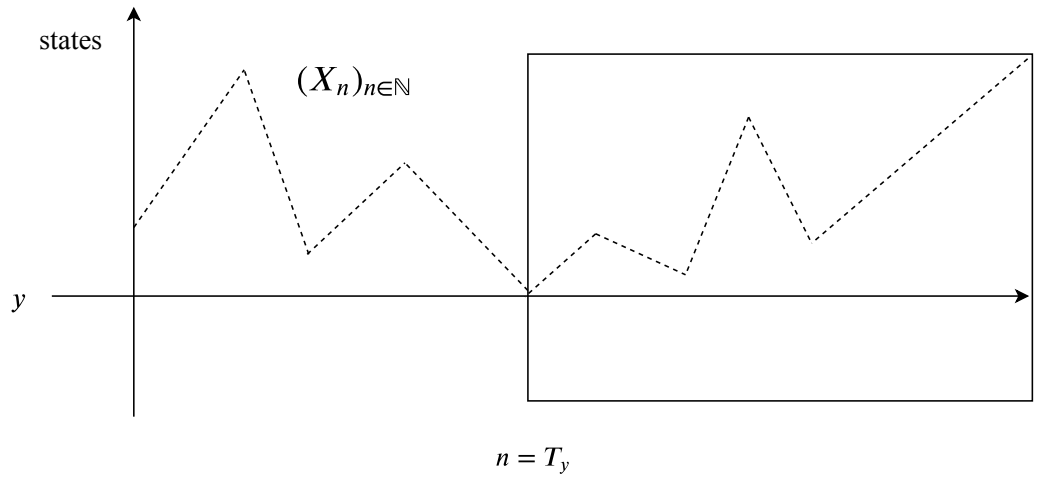

Remark: In discrete time (even with a general state space), *all* Markov chains with a homogeneous transition mechanism have the Strong Markov Property. Now, we can use the SMP to discover and prove things about Markov chains.

3.2 Iterating

From $T_y^0 = 0$, for $k \geq 1$

$$T_y^k := \min\{n > T_y^{k-1} : X_n = y\}.$$

Note T_y^k is a k^{th} iterate of the scheme for defining $T_y = T_y^1$, not the k^{th} power of T_y. Suppose we have a path that hits the state y a finite number of times $n \geq 1$, say exactly four times: T_y, T_y^2, T_y^3, T_y^4. Then by our convention, we say that $T_y^5 = \infty$. Consider the random variable which is the total number of hits of y, at any time $n \geq 1$, deliberately not counting a hit at time 0 if $X_0 = y$:

$$N_y := \sum_{n=1}^{\infty} \mathbb{1}(X_n = y).$$

The possible values of N_y are $\{0, 1, 2, \ldots, \infty\}$, an infinite time horizon. By the logic of the definitions, there is the identity of events:

$$(N_y = 0) = (T_y = \infty).$$

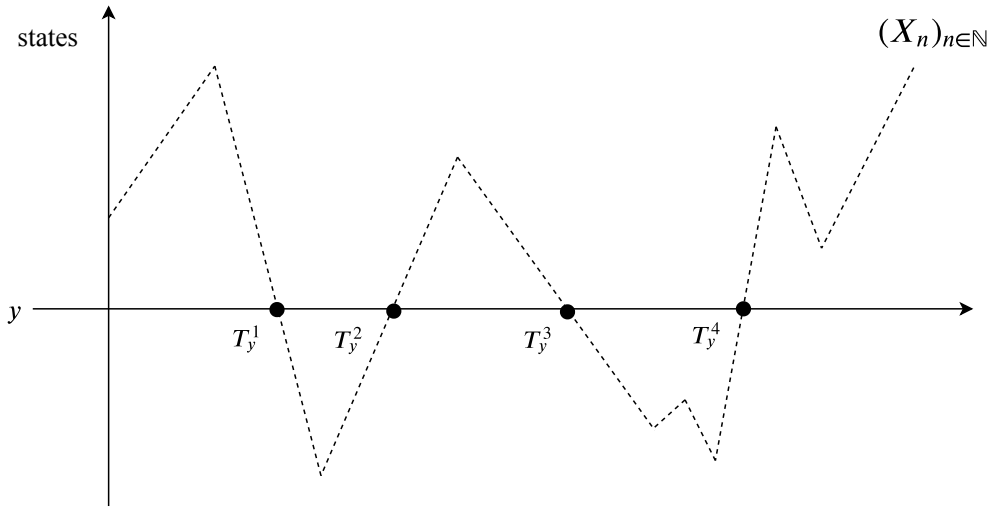

As another example, consider $(N_y \geq 1)$, the complement of $(N_y = 0)$ because we include ∞ as a part of $(N_y \geq 1)$. Hence

$$(N_y \geq 1) = (T_y < \infty).$$

Recall $N_y := \sum_{n=1}^{\infty} \mathbb{1}(X_n = y)$ is simply counting the number of hits on y. Pitman asks the audience to find expressions in terms of T_y^k for the left hand side of

$$\begin{aligned}
(N_y \geq 3) &= (T_y^3 < \infty) \\
(N_y = 3) &= (T_y^3 < \infty, T_y^4 = \infty) \\
(N_y \geq k) &= (T_y^k < \infty) \\
(N_y = k) &= (T_y^k < \infty, T_y^{k+1} = \infty).
\end{aligned}$$

Now let's discuss the probabilities. The SMP gives

$$\mathbb{P}_y(T_y^k < \infty) = \rho_y^k,$$

where k on the RHS is a power, and k on the LHS is an index. Now taking $k = 1$, we have the definition of ρ_y:

$$\mathbb{P}_y(T_y < \infty) = \rho_y$$

called the *first return probability* of state y. Now how to get from ρ_y to ρ_y^2? Basically, this is by the SMP. Observe

$$(T_y^k < \infty) = (N_y \geq k).$$

which tells us that the probability of hitting y at least $k \in \mathbb{N}_0$ times is

$$\mathbb{P}_y(N_y \geq k) = \rho_y^k$$

If we want to find the point probability that $N_y = k$, we take

$$\begin{aligned}\mathbb{P}_y(N_y = k) &= \mathbb{P}_y(N_y \geq k) - \mathbb{P}_y(N_y \geq k+1) \\ &= \rho_y^k - \rho_y^{k+1} \\ &= \boxed{\rho_y^k(1-\rho_y)}\end{aligned}$$

Now *either* $\rho_y = 1$ and this probability is 0 for all $k < \infty$, pushing all the probability to $\mathbb{P}_y(N_y = \infty) = 1$, *or* $\rho_y < 1$ in which case the probability distribution (starting at y) of $N_y := \sum_{n=1}^{\infty} \mathbb{1}(X_n = y)$ is geometric(p) on $\{0, 1, 2, \ldots\}$ with parameter

$$p = 1 - \rho_y = \mathbb{P}_y(T_y = \infty) = \mathbb{P}_y(N_y = 0).$$

Notice, via the tail-sum formula for $\mathbb{E}$ of a non-negative integer valued random variable

$$\mathbb{E}_y N_y = \sum_{k=1}^{\infty} \mathbb{P}_y(N_y \geq k) = \sum_{k=1}^{\infty} \rho_y^k = \frac{\rho_y}{1-\rho_y} = \frac{q}{p}$$

for $q = \rho_y$ and $p = 1 - \rho_y$, in agreement with the standard formula for $\mathbb{E}$ of a geometric(p) variable.

3.3 State Classification

There are two cases to consider.

(1) y is *transient* : $0 \leq \rho_y < 1$. This implies that our expected number of visits is:

$$\mathbb{E}_y N_y = \frac{\rho_y}{1-\rho_y} < \infty$$

which implies

$$\mathbb{P}_y(N_y < \infty) = 1,$$

which says that if we have a transient state, then we only return to y a finite number of times. In other words, after some point, the Markov chain never visits y again.

(2) y is *recurrent* : $\rho_y = 1$. In other words, $\mathbb{P}_y(N_y = \infty) = 1$ in that given any number of hits, we are sure to hit y again.

3.3.1 Constructing ρ_y

Here is an explicit a formula:

$$\begin{aligned}\rho_y &= \mathbb{P}_y(T_y = 1) + \mathbb{P}_y(T_y = 2) + \mathbb{P}_y(T_y = 3) + \cdots \\ &= P(y,y) + \sum_{y_1 \neq y} P(y,y_1)P(y_1,y) + \sum_{y_1 \neq y} \sum_{y_2 \neq y} P(y,y_1)P(y_1,y_2)P(y_2,y) + \cdots\end{aligned}$$

But this is not so nice to work with.

Exercise: Show that for $n \geq 2$ the nth term $\mathbb{P}_y(T_y = n)$ can be expressed in matrix notation as $P(y, \cdot)K^{n-2}P(\cdot, y)$ for a suitable matrix K to be determined. Note that K is *sub-stochastic* with non-negative entries and row sums ≤ 1.

3.4 Lemma 1.3

Reference. Durrett's, *Essentials of Stochastic Processes* Page 16. Take B to be a set of states. Hypothesis: Suppose the probability starting at x that $T_B \leq k$ is at least $\alpha > 0$ for some fixed k and all x:

$$\mathbb{P}_x(T_B \leq k) \geq \alpha > 0 \text{ for all states } x$$

Then

$$\mathbb{P}_x(T_B > nk) \leq (1-\alpha)^n.$$

As an example where the hypothesis is obviously satisfied, consider the Gambler's Ruin chain with state 0 and N as absorbing states. That is, $B = \{0, N\}$, with $P(i, i+1) = p$ and $P(i, i-1) = q$ for $0 < i < N$. Then this condition holds with $k = N$ and

$$\alpha = p^N + q^N > 0$$

because no matter where you start away from the boundary states, a sequence of either at most N consecutive up steps or N consecutive down steps will get you to the boundary.

Proof. The conclusion is obvious by taking complements if $n = 1$: $\mathbb{P}_x(T_B > k) \leq 1 - \alpha$ for all x. Now by induction on n. Observe that

$$\begin{aligned}
\mathbb{P}_x(T_B > (n+1)k) &= \mathbb{P}_x(T_B > nk \text{ and after time } nk \text{ before time } (n+1)k \text{ still don't hit } B) \\
&= \sum_{y \notin B} \mathbb{P}_x(T_B > nk, X_{nk} = y, \text{ do not hit } B \text{ before time } (n+1)k) \\
&= \sum_{y \notin B} \mathbb{P}_x(T_B > nk, X_{nk} = y)\mathbb{P}_y(T_B > k) \\
&\leq \sum_{y \notin B} \mathbb{P}_x(T_B > nk, X_{nk} = y)(1-\alpha) \\
&= \mathbb{P}_x(T_B > nk)(1-\alpha) \\
&\leq (1-\alpha)^n(1-\alpha) = (1-\alpha)^{n+1}.
\end{aligned}$$

□

Pitman mentions Kai Lai Chung of Stanford who presents this Lemma in terms of a pedestrian repeatedly crossing the road. Depending on visibility and weather conditions, (current state), there is a varying chance of making it to the other side in k steps. But suppose that no matter how favorable the visibility and weather conditions, there is always at least a small chance α that the pedestrian gets killed while crossing the road. Then, if the pedestrian repeatedly attempts to cross the road, eventually they will get killed, with a geometric bound as above on how long that takes. In the language of Markov chains, if there is always at least a some strictly postive chance of the chain reaching a boundary B in the next k steps, no matter where it starts, eventually the chain will hit such a boundary set. Observe that in standard real numbers,

$$0 \leq \mathbb{P}_x(T_B = \infty) \leq \mathbb{P}_x(T_B > nk) \leq (1-\alpha)^n, \ \forall_n$$

implies

$$\mathbb{P}_x(T_B = \infty) = 0$$

Back to Gambler's Ruin. Suppose $0 < p < 1$. In the language of transient and recurrent states, every state $x \notin \{0, N\}$ is transient! Moreover, $x \in \{0, N\}$ is recurrent. Here's a Gambler's Ruin chain for $N = 4$.

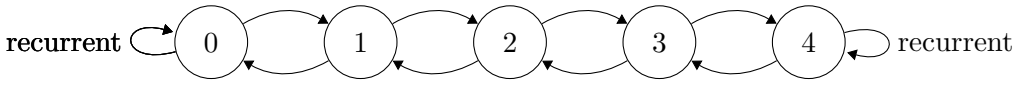

Irreducible Matrix

We say that a matrix P is **irreducible** if

$$\forall \ x, y \in \mathcal{S}, \ \exists n : P^n(x, y) > 0$$

In words, for every pair of states x, y, it is possible to get from x to y in some number n of steps. Here $n = n(x, y)$ is a function of x, y. If matrix P is irreducible, then either

all states are recurrent **or** all states are transient

We then say the matrix P is "recurrent" or "transient", meaning that it drives a chain all of whose states are recurrent or transient, as the case may be.

This and other properties of states of a chain with irreducible transition matrix P, which hold for one state iff they hold for all states, are called *solidarity properties*. Other examples are the conditions that $\mathbb{E}_x T_x < \infty$, and that state x is *aperiodic* as discussed in the text, or that state x has a particular period d.

Easy fact. (Pigeon hole principle: with a finite state space, and infinitely many steps, some state must be hit infinitely often): Suppose $\mathcal{S}$ is finite and P is irreducible. Then P is recurrent. Notice the Gambler's Ruin chain exhibits a matrix that is **not** irreducible, which can be seen via the definition above and the requirement that there exists some n where $P^n(x, y) > 0$.

LECTURE 4

Exchangeability, Stationary Distributions, and Two State Markov Chains

4.1 Sampling without Replacement

Consider $(X_1, \ldots, X_N)$ an exhaustive random sample without replacement from a box of $N = A + B$ tickets, with A labeled 1 (success) and B labeled 0 (fail). Let $S_0 := 0$ and $S_n := X_1 + \cdots + X_n$ the number of 1s and $\bar{S}_n := n - S_n$ the number 0s in the first n places of the sample. Then $((\bar{S}_n, S_n), 0 \le n \le N)$ is a Markov chain with transition matrix

$$P((f,s),(f+1,s)) = \frac{B-f}{A+B-f-s} \tag{4.1}$$

$$P((f,s),(f,s+1)) = \frac{A-s}{A+B-f-s} \tag{4.2}$$

and all other entries 0. In the following diagrams, with Cartesian coordinates (f, s), the the horizontal scale counts the number of failures f, the vertical scale counts the number successes s, and the sum of coordinates is the number of draws $n = f + s$. The chain $W_n := (\bar{S}_n, S_n)$ starts at the origin $(\bar{S}_0, S_0) = (0, 0)$ at time $n = 0$, and terminates at (B, A) at time $n = N = A + B$.

- each step right in this chain increments the first component f of (f, s) to $f+1$ for a failure (0) in the sequence $(X_1, \ldots, X_N)$.
- each step up in this chain increments of the second component s of (f, s) to $s + 1$ for a success (1) in the sequence $(X_1, \ldots, X_N)$.

Altogether there are $N = A + B$ steps, with A steps up and B steps right. All $\binom{A+B}{A}$ possible paths of the chain are equally likely.

Figure 4.1: Transition probabilities for sampling without replacement. Here $A = 3, B = 7$. The only possible transitions are one step up or one step right, following arrows on the grid of possible states (f, s), with $0 \le f \le 7$ the number of failures (0) and $0 \le s \le 3$ the number of successes, after $f + s$ draws without replacement from 7 values 0 and 3 values 1. The pair of transition probabilities out of each state (f, s) is represented by a vector with tail (f, s) and head $(f + q, s + p)$ for $q = P((f, s), (f + 1, s)$ and $p = P((f, s), (f, s + 1)$ as above. The head of each vector is the conditional mean of the random vector $(\bar{S}_{n+1}, bS_{n+1})$ given $(\bar{S}_n = f, S_n = s)$ with $n = f + s$. All the transition vectors point towards the terminal state of the chain at $(B, A) = (7, 3)$ after $n = 10$ draws.

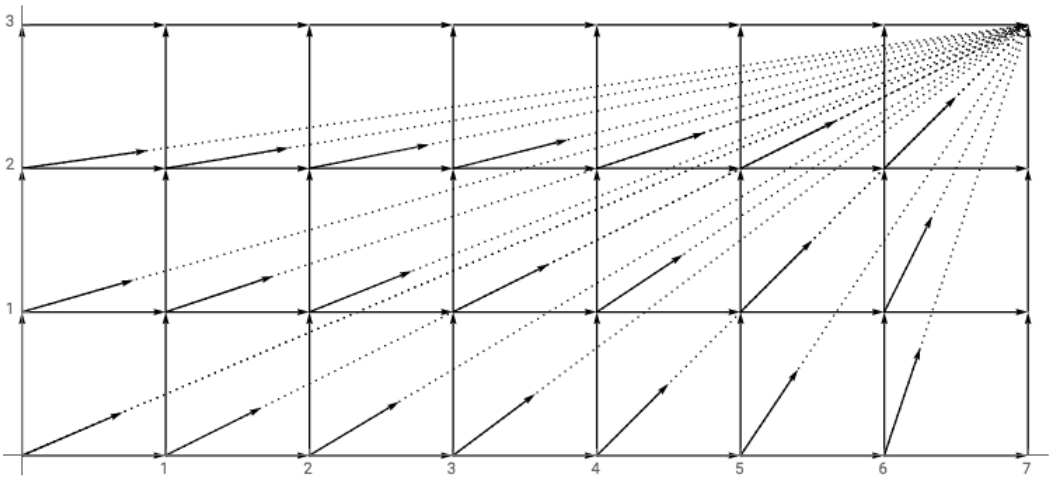

4.2 Exchangeability and Reversibility

It is an important general property of a sample without replacement $(X_1, \ldots, X_N)$ that these random variables are *exchangeable*, meaning that for every permutation σ of $[N] := \{1, \ldots, N\}$

$$(X_{\sigma(1)}, \ldots, X_{\sigma(N)}) \stackrel{d}{=} (X_1, \ldots, X_N) \tag{4.3}$$

where $\stackrel{d}{=}$ denotes equality in distribution. [2] See Pitman *Probability* Section 3.6. In particular, this holds for the sample $(X_1, \ldots, X_N)$ of A ones and B zeros considered here. Except in degenerate cases, the sequence $(X_1, \ldots, X_N)$ is not Markov: given $X_1, \ldots, X_n$ the conditional probability that $X_{n+1} = 1$ is $(A - S_n)/(N - n)$ which is typically not just a function of X_n, but involves all of the previous values $X_1, \ldots, X_n$ through their sum S_n, the number of 1s in the first n draws without replacement. However, in the model of sampling without replacement from A values 1 and B values 0, it is instructive to study the common joint distribution of every pair of draws

$$(X_{\sigma(1)}, X_{\sigma(2)}) \stackrel{d}{=} (X_1, X_2) \qquad (\sigma(1) \ne \sigma(2)). \tag{4.4}$$

The joint probability function of this pair of draws is obtained by assuming that all $(A + B)(A + B - 1)$ possible pairs of different tickets are equally likely to appear on

the first and second draws. By counting pairs of different tickets

$$\mathbb{P}(X_1 = 0, X_2 = 0) = \frac{B(B-1)}{(A+B)(A+B-1)} \tag{4.5}$$

$$\mathbb{P}(X_1 = 1, X_2 = 1) = \frac{A(A-1)}{(A+B)(A+B-1)} \tag{4.6}$$

$$\mathbb{P}(X_1 = 0, X_2 = 1) = \mathbb{P}(X_1 = 1, X_2 = 0) = \frac{AB}{(A+B)(A+B-1)} \tag{4.7}$$

The last equality of the two off-diagonal probabilities is important. In counting pairs of different tickets, this comes from $BA = AB$: the number of different ways to get 1 followed by 0 from the first two draws is equal to the number of different ways to get 0 followed by 1. In probabilistic terms, the equality (4.7) of off-diagonal probabilities gives the equality in distribution

$$(X_2, X_1) \stackrel{d}{=} (X_1, X_2) \tag{4.8}$$

Such a pair of random variables (X_1, X_2) is called either *reversible* or *exchangeable*. In terms of a joint distribution table of numerical random variables X_1 and X_2, displayed in Cartesian coordinates with values of X_1 horizontal and values of X_2 vertical, such a distribution is symmetric with respect to reflection accross the set of diagonal values $(X_1 = X_2)$:

$$\mathbb{P}(X_2 = x, X_1 = y) = \mathbb{P}(X_1 = x, X_2 = y) \tag{4.9}$$

for all possible values x and y. If $x = y$ this identity is trivial. If X_1 and X_2 have only two possible values 0 and 1, there are only two possible off-diagonal pairs $(0, 1)$ and $(1, 0)$. So for indicator variables, (X_1, X_2) is reversible iff (4.9) holds for the single pair $(x, y) = (0, 1)$, as it does in (4.7).

In general, for $N \geq 2$, a sequence of random variables $(X_1, \ldots, X_N)$ is called *reversible* if (4.3) holds just for the single permutation σ which reverses the order of indices, that is

$$(X_N, \ldots, X_1) \stackrel{d}{=} (X_1, \ldots, X_N). \tag{4.10}$$

For a random vector of length $N = 2$, reversible is the same as exchangeable, because there are only two permutations of $\{1, 2\}$, the identity permutation, for which there is nothing to check, and the permutation which switches 1 and 2. For a random vector of length $N \geq 3$ exchangeable implies reversible, but not conversely. For instance, if $N = 3$ there are $3! - 1 = 5$ permutations besides the identity, and reversibility only involves an identity in distribution for just one of these 5 permutations. See also further discussion below.

In sampling without replacement from B values 0 and A values 1, the first variable X_1 has distribution $\mathbb{P}(X_1 = i) = \pi_i$ given by

$$(\pi_0, \pi_1) = \frac{(B, A)}{A + B} \tag{4.11}$$

and the step from X_1 to X_2 is made according to the transition probability matrix

$$P = \begin{bmatrix} P_{00} & P_{01} \\ P_{10} & P_{11} \end{bmatrix} = \frac{1}{A+B-1} \begin{bmatrix} B-1 & A \\ B & A-1 \end{bmatrix} \tag{4.12}$$

This description of the joint distribution of (X_1, X_2) is logically equivalent to the previous description of the joint probability function (4.5)-(4.7) by four applications of the product rule $\mathbb{P}(CD) = \mathbb{P}(C)\mathbb{P}(D \,|\, C)$. Either description implies $(X_1, X_2) \stackrel{d}{=} (X_2, X_1)$ and hence $X_2 \stackrel{d}{=} X_1$. More algebraically, the distribution of X_2 is determined by

$$\begin{aligned} \mathbb{P}(X_2 = 1) &= \mathbb{P}(X_1 = 0, X_2 = 1) + \mathbb{P}(X_1 = 1, X_2 = 1) \\ &= \mathbb{P}(X_1 = 0)\mathbb{P}(X_2 = 1 \,|\, X_1 = 0) + \mathbb{P}(X_1 = 1)\mathbb{P}(X_2 = 1 \,|\, X_1 = 1) \\ &= \frac{B}{(A+B)} \frac{A}{(A+B-1)} + \frac{A}{(A+B)} \frac{(A-1)}{(A+B-1)} \\ &= \frac{A(A+B-1)}{(A+B)(A+B-1)} = \frac{A}{A+B} \end{aligned}$$

The probability $\mathbb{P}(X_2 = 0)$ can be found similarly, or by

$$\mathbb{P}(X_2 = 0) = 1 - \mathbb{P}(X_1 = 1)$$

since the only possible values of X_2 are 0 and 1. The simple algebraic structure of this joint distribution of the pair of indicator variables (X_1, X_2) derived from sampling without replacement from A values 1 and B values 0 is worth understanding thoroughly, especially the reversibility $(X_1, X_2) \stackrel{d}{=} (X_2, X_1)$ which implies $X_2 \stackrel{d}{=} X_1$. The algebraic structure of this joint distribution of dependent indicators (X_1, X_2), with two parameters A and B, in terms of which the algebra comes out very nicely, turns out to be shared by every pair of exchangeable indicator variables X_1 and X_2 which are not independent.

Remark. There is a simple construction of dependent indicator variables (X_1, X_2) with various levels of dependence, which may be helpful. Draw a Venn diagram with two regions V_1 and V_2, each of area p, for some fixed $0 < p < 1$. Let X_i be the indicator of V_i. Now move the regions around in the diagram to vary their overlap. There is no loss of generality in making each region a rectangle of height 1 over some base interval of length p. Now each V_i may be identified with a subinterval of $[0, 1]$ of length p, and it is just a matter of moving around two subintervals of $[0, 1]$, each of length p, and considering the possible length of overlap of these two intervals V_i. You can treat each X_i as a function $X_i(\omega)$ of $\omega \in [0, 1]$ with value $X_i(\omega) = 1$ if ω falls in some interval V_i of length p, and value 0 if $\omega \notin V_i$. So $\mathbb{E}X_i = 1 \times p + 0 \times (1-p) = p$. The most the two intervals can overlap is if they are identical, which makes $\mathbb{E}X_1X_2 = p$ and $\mathbb{P}(X_1 = X_2) = 1$, with correlation 1. By an elementary argument (Boole's inequality)

$$0 \leqq \mathbb{E}(1 - X_1)(1 - X_2) = 1 - 2p + \mathbb{E}(X_1X_2)$$

the least they can overlap is if

$$\mathbb{P}(V_1V_2) = \mathbb{E}(X_1X_2) = (2p-1)_+$$

which is 0 if $0 \leqq p \leqq 1/2$, and $2p-1$ if $1/2 < p \leqq 1$. This bound is achieved by $V_1 = [0,p]$ and $V_2 = [1-p,1]$. Any value of $\mathbb{P}(V_1V_2)$ in this range $[(2p-1)_+, p]$ determines a possible exchangeable joint distribution of (X_1, X_2) with $\mathbb{E}X_1 = \mathbb{E}X_2 = p$, which is realized on $[0,1]$ by any two intervals of length p with the assigned overlap. And for an allowed value of $\mathbb{E}X_1X_2$, Every exchangeable joint law of a pair of indicators (X_1, X_2) is completely determined by the common value of $p := \mathbb{E}X_i$ and the value of $\mathbb{E}X_1X_2$ in $[(2p-1)_+, p]$. Always included in the range of possible values of $\mathbb{E}(X_1X_2)$ is the value p^2 for independent X_i. Thus

$$(2p-1)_+ < p^2 < p \text{ for } 0 < p < 1$$

as you should check by sketching graphs of all three functions of p over $[0,1]$. Indicators X_1 and X_2 are called *positively dependent* or *negatively dependent* according to the sign of $\mathrm{Cov}(X_1, X_2) := \mathbb{E}X_1X_2 - \mathbb{E}X_1\mathbb{E}X_2$.

4.3 Stationary Distributions

For a pair of discrete random variables (X_1, X_2), write either

$$X_1 \sim \pi \text{ and } (X_2 \,|\, X_1) \sim P(X_1, \cdot\,)$$

or

$$\mathbb{P}(X_1 \in \cdot) = \pi(\cdot) \text{ and } \mathbb{P}(X_2 \in \cdot \,|\, X_1) = P(X_1, \cdot)$$

to mean that X_1 has distribution π, and the conditional distribution of X_2 given $X_1 = x$ is given by the row $P(x, \cdot)$ of some transition probability matrix P, for every possible value x of X_1. This prescription of a distribution π for X_1 and the conditional distribution $P(X_1, \cdot)$ for X_2 given X_1 uniquely determines the joint distribution of X_1 and X_2, and is equivalent to the formula for the joint probability function of (X_1, X_2)

$$\mathbb{P}(X_1 = x, X_2 = y) = \pi(x)P(x,y)$$

as x and y range over all possible values of X_1 and X_2 respectively. The distribution of X_2 is then determined by the matrix operation $X_2 \sim \pi P(\,\cdot\,)$:

$$\begin{aligned}\mathbb{P}(X_2 = y) &= \sum_x \mathbb{P}(X = x)\mathbb{P}(Y = y \,|\, X = x) \\ &= \sum_x \lambda(x)P(x,y) = (\pi P)(y).\end{aligned}$$

In particular, for X_1 and X_2 with the same set of possible values,

$$X_1 \stackrel{d}{=} X_2 \iff \pi = \pi P \qquad \text{meaning} \tag{4.13}$$

$$\sum_x \pi(x)P(x,y) = \pi(y) \text{ for all states } y. \tag{4.14}$$

Then π is called a *stationary* (or *invariant* or *equilibrium* or *steady state*) *distribution* for the transition matrix P. This condition $X_1 \stackrel{d}{=} X_2$ is implied by the stronger *reversibility condition*

$$(X_1, X_2) \stackrel{d}{=} (X_2, X_1) \iff \pi(x)P(x,y) = \pi(y)P(y,x) \text{ for all } x, y \tag{4.15}$$

when π is called a *reversible equilibrium distribution* for the transition matrix P. The equations in (4.14) are called *balance equations* while those in (4.15) are called *detailed balance equations*. If there are $|S| = N$ states, there are N different balance equations, and $\binom{N}{2}$ different detailed balance equations. For a prescribed transition matrix P, to solve either system of equations to obtain a stationary probability distribution π you must add the constraint $\sum_x \pi(x) = 1$. Issues of existence and uniqueness of solutions of these balance equations are treated in the text and will be discussed further in following lectures. It is often easy to see directly that some distribution π provides a reversible equilibrium for a particular transition matrix P. This just involves checking $\pi(x)P(x,y) = \pi(y)P(y,x)$ for $x \neq y$, which was already noticed above in the case of sampling without replacement, by counting outcomes. No summations were involved.

Exercise: The text on page 22 has a nice *sand metaphor* for the balance equations. Explain the meaning of detailed balance in terms of the sand metaphor.
Here are some easy consequences of these definitions, all of which you should be able to derive for yourself without consulting any text::

- If $(X_0, X_1, X_2, \dots,)$ is a Markov chain with homogeneous transition matrix P and $X_0 \sim \pi$ with $\pi P = \pi$, then for all positive integers n and N

  $$(X_0, \dots, X_N) \stackrel{d}{=} (X_n, \dots, X_{n+N}).$$

  A stochastic process $(X_0, X_1, X_2, \dots,)$ with this property is called *stationary*. In words: the finite dimensional distributions of a stationary process are invariant with respect to a shift in time.

- If a distribution π solves the detailed balance equations for P , then π also solves the balance equations for P;

- If X_1 and X_2 are random variables, each with only two possible values, then $X_1 \stackrel{d}{=} X_2$ iff $(X_1, X_2) \stackrel{d}{=} (X_2, X_1)$;

- If X_1 and X_2 have three or more possible values, it is possible to have $X_1 \stackrel{d}{=} X_2$ without $(X_1, X_2) \stackrel{d}{=} (X_2, X_1)$. An example on three states is (X_1, X_2) with transition matrix

  $$P = \begin{bmatrix} 0 & 1 & 0 \\ 0 & 0 & 1 \\ 1 & 0 & 0 \end{bmatrix}$$

  corresponding to deterministic rotation by one step around three states $0, 1, 2$ arranged in a circle. The unique equilibrium distribution is the uniform distribution $\pi = (1, 1, 1)/3$, and this equilibrium is not reversible. This is an

example of a *periodic chain* with *period* 3. In general for transition matrix P, the *period of a state* x is the greatest common divisor $d(x)$ of the set of positive integers n such that $P^n(x,x) > 0$. For an irreducible matrix P, Lemma 1.17 of the text shows that $d(x) \equiv d$ for some positive integer d, called the *period of* P.

- The above transition matrix P on 3 states is *doubly stochastic*, meaning that all its row sums are 1 and all its column sums are 1. For an $S \times S$ matrix P with S finite, the uniform distribution π on S is P-invariant iff P is doubly stochastic. (Text Theorem 1.14)

- If π is P-invariant, then π is P^n-invariant for every positive integer n. Here P^n is the nth iterate of the transition matrix P, which is the n-step transition matrix for a Markov chain with homogeneous transition matrix P.

- In terms of a Markov chain $(X_0, X_1, \ldots)$ with $X_0 \sim \pi$ and homogeneous transition matrix P, an equilibrium π for P is reversible iff for every $N \geq 1$ there is the equality in distribution

$$(X_0, \ldots, X_N) \stackrel{d}{=} (X_N, \ldots, X_0).$$

See the text §1.4 and §1.5 for further discussion and many examples. Two less obvious but very important facts, treated in the text in §1.6, 1.7 and 1.8 are :

- if P is irreducible with a finite number of states, then there is a unique stationary distribution π for P, specifically

$$\pi_j = \frac{1}{\mathbb{E}_j T_j} \tag{4.16}$$

where $\mathbb{E}_j T_j$ is the mean return time of state j. This is also true more generally if P is irreducible and *positive recurrent*, meaning that $\mathbb{E}_j T_j < \infty$ for some (hence all) states j,

- If P is irreducible and positive recurrent and *aperiodic*, meaning that some (and hence every) state x has period 1, then

$$\lim_{n\to\infty} P^n(i,j) = \pi_j \tag{4.17}$$

as above for all states j.

- Consequently, for any Markov chain (X_n) with countable state space S and such a transition matrix P, no matter what the distribution of X_0, there is the convergence in distribution $X_n \stackrel{d}{\to} \pi$ as $n \to \infty$, meaning

$$\lim_{n\to\infty} \mathbb{P}(X_n = j) = \pi_j \qquad (j \in S).$$

Exercise. Explain exactly how each of these results can be deduced from specific theorems in the text.

4.4 Two State Transition Matrices

Suppose A and B are positive integers with $A+B = N \geq 3$, and consider $(X_1, \ldots, X_N)$ an exhaustive sample without replacement from A values 1 and B values 0. In the sample (X_1, X_2, X_3) of size 3, every pair of variables has the reversible joint distribution of (X_1, X_2) displayed in (4.5) - (4.7).

Exercise. Check, by calculations like (4.5) - (4.7) that this sequence (X_1, X_2, X_3) of exchangeable indicators is not Markovian.
For P the transition matrix of (X_1, X_2) displayed in (4.12), with parameters (B, A), the iterates P^n of P have no obvious meaning in terms of the exhaustive sample $(X_k, 1 \leq k \leq N)$. In particular, the conditional distribution of X_3 given X_1 is provided by P, not by P^2, as you can see from the formula for P^2 for a two state Markov matrix P (Homework 2). Rather, this example of (X_1, X_2, X_3) derived from sampling without replacement is non-Markovian and exchangeable. The single transition matrix P provides the conditional distribution of X_i given X_j for every $i \neq j$.
Observe that the matrix P defined by (4.12), with two parameters A and B, has row sums 1 not only for all positive integers A and B, but also for any choice of real parameters A and B with $A + B - 1 \neq 0$. In fact, this construction generates every 2×2 transition matrix P except for the relatively uninteresting *Bernoulli*(p) *matrices*

$$\begin{bmatrix} q & p \\ q & p \end{bmatrix} \qquad (0 \leq p \leq 1, p + q = 1).$$

For $p = A/(A+B)$ the matrix above is associated with sampling without replacement from a population of A ones and B zeros. For general $0 \leq p \leq 1$, the Bernouli(p) matrix corresponds to an unlimited sequence of independent Bernoulli(p) trials. You can easily check the following proposition:

Proposition 0.1. *Let P be a 2×2 transition matrix with $P_{01} \neq P_{11}$. Then P is of the algebraic form* (4.12)*, that is*

$$\begin{bmatrix} P_{00} & P_{01} \\ P_{10} & P_{11} \end{bmatrix} = \frac{1}{A + B - 1} \begin{bmatrix} B - 1 & A \\ B & A - 1 \end{bmatrix} \tag{4.18}$$

for a unique pair of real parameters (B, A):

$$\begin{bmatrix} B - 1 & A \\ B & A - 1 \end{bmatrix} = \frac{1}{P_{01} - P_{11}} \begin{bmatrix} P_{00} & P_{01} \\ P_{10} & P_{11} \end{bmatrix} \tag{4.19}$$

Assume further that $P_{01} + P_{10} > 0$, to exclude the trivial case $P = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$ corresponding to $(B, A) = (0, 0)$. Then the unique stationary distribution for P is $\pi = (\pi_0, \pi_1)$ defined by

$$(\pi_0, \pi_1) = \frac{(P_{10}, P_{01})}{P_{10} + P_{01}} = \frac{(B, A)}{A + B}; \tag{4.20}$$

This π is a reversible equilibrium for P: if $X_1 \sim \pi$ then the joint distribution of X_1 and X_2 is given by the formulas (4.5)–(4.7) for sampling without replacement, without the requirement that A and B are positive integers. This makes

$$Cov\,(X_1, X_2) := \mathbb{E}X_1X_2 - (\mathbb{E}X_1)(\mathbb{E}X_2) = \frac{-AB}{(A+B)^2(A+B-1)}. \tag{4.21}$$

The range of parameters (A, B) in this construction has two connected components:

- *$A \geq 1$ and $B \geq 1$, when X_1 and X_2 are negatively dependent;*
- *$A = -a \leq 0$ and $B = -b \leq 0$, with $a + b > 0$, when X_1 and X_2 are positively dependent; then in terms of $a = -A \geq 0$ and $b = -B \geq 0$*

$$\begin{bmatrix} P_{00} & P_{01} \\ P_{10} & P_{11} \end{bmatrix} = \frac{1}{a+b+1} \begin{bmatrix} b+1 & a \\ b & a+1 \end{bmatrix}$$

The transition matrix (4.4) of independent Bernoulli(p) trials is recovered as the limit case when either A and B both tend to $+\infty$, or A and B both tend to $-\infty$, with $A/(A+B) \to p$.

Exercise. Check that the formula (4.20) for the stationary distribution π of a two state chain agrees with the general formula $\pi_i = 1/\mathbb{E}_i T_i$ in (4.16), by directly evaluating $\mathbb{E}_i T_i = \sum_{n=1}^{\infty} \mathbb{P}_i(T_i \geq n)$ for the two state chain.

4.5 Two State Transition Diagrams

To fully understand the dependence between the variables X_1 and X_2 in a two-state Markov chain, and how this affects the distribution of $X_1 + X_2$, regard each X_i as the indicator of success on trial i in a pair of dependent trials. Let $S_2 := X_1 + X_2$ be the number of successes in the two trials, and

$$\bar{S}_2 := (1 - X_1) + (1 - X_2) = 2 - S_2$$

the number of failures in the two trials. With $(S_0, \bar{S}_0) := (0, 0)$ and $(S_1, \bar{S}_1) := (X_1, 1 - X_1)$, the pair of dependent indicators (X_1, X_2) is now encoded in the sequence

$$W_0 := (\bar{S}_0, S_0); \qquad W_1 := (\bar{S}_1, S_1); \qquad W_2 := (\bar{S}_2, S_2).$$

So (W_0, W_1, W_2) is a Markov chain with 6 states:

- $(0, 0)$ is the initial state of $W_0 = (\bar{S}_0, S_0)$,
- $(0, 1)$ and $(1, 0)$ are the two possible states of $W_1 = (1 - X_1, X_1)$, corresponding to $(X_1 = 1)$ and $(X_1 = 0)$ respectively
- $(0, 2)$ and $(1, 1)$ and $(2, 0)$ are the three possible states of $W_2 = (\bar{S}_2, S_2)$, corresponding to the events

$$\begin{aligned} (W_2 = (0, 2)) &= (S_2 = 2) = (X_1 = 1, X_2 = 1) \\ (W_2 = (1, 1)) &= (S_2 = 1) = (X_1 = 0, X_2 = 1) \cup (X_1 = 1, X_2 = 0) \\ (W_2 = (2, 0)) &= (S_2 = 0) = (X_1 = 0, X_2 = 0). \end{aligned}$$

There are two motivations for this proliferation of states:

- the distribution of $S_2 = X_1 + X_2$, the number of successes in the two dependent trials, is naturally of interest; this is encoded in the distribution of W_2.
- each of the $2 \times 2 = 4$ possible values of (X_1, X_2) corresponds to two consecutive transitions of the chain (W_0, W_1, W_2); vectors representing probabilities of these transitions are easily displayed graphically, as in Figure 4.1 for the cumulative counts in sampling without replacement.

For the transition matrix P as in (4.12) derived from (X_1, X_2) a sample of size 2 without replacement from A values 1 and B values 0, the transition diagram of $(\bar{S}_n, S_n)$ for $0 \le n \le 2$ is just the bottom left corner of the larger diagram already displayed in Figure 4.1 for $A = 3$ and $B = 7$, involving just the first two steps away from $(0, 0)$. See Figure 4.2. The special feature of this transition diagram, that lines through the various probability vectors all pass through the point (B, A), is essentially an algebraic property of the transition rules for sampling without replacement. Remarkably, this algebraic property extends to the the setting of the above proposition, as follows:

Corollary 0.1. *Let a 2×2 transition probability matrix P with $P_{01} \neq P_{11}$ be represented in the form (4.18) for a pair of real parameters (B, A). Consider a Cartesian plane of pairs of real numbers $w = (f, s)$, with the six pairs indexed by non-negative integers f and s with $f + s \le 2$ representing possible states of the chain $W_i := (\bar{S}_i, S_i)$ for $i \in \{0, 1, 2\}$, derived as above from a pair of indicator variables (X_1, X_2) with transition matrix (4.18). For each of the states $w = (0, 0)$ or $(0, 1)$ or $(1, 0)$ represent the two transition probabilities of the W-chain out of state w by a vector pointing from w to $w + v(w)$, where $v(w)$ is the following probability vector:*

$$v(0,0) = \lambda(\cdot) = (\lambda_0, \lambda_1) \text{ is the distribution of } X_1 \tag{4.22}$$

$$v(0,1) = P(1,\cdot) = (P_{10}, P_{11}) \text{ is the distribution of } X_2 \text{ given } X_1 = 1 \tag{4.23}$$

$$v(1,0) = P(0,\cdot) = (P_{00}, P_{01}) \text{ is the distribution of } X_2 \text{ given } X_1 = 0. \tag{4.24}$$

Regard these three probability vectors, together with the probability vector λP representing the unconditional distribution of X_2, and the vector with components (B, A), as five points in the (f, s)-plane. Then:

(i) (B, A) is the unique point of intersection of the lines through w and $w + v(w)$ for $w = (0, 1)$ and $w = (1, 0)$.

(ii) For each initial distribution λ of X_1, the line through λ in direction λP passes through (B, S).

(iii) The point

$$\lambda + P\lambda = \mathbb{E}(\bar{S}_2, S_2) \tag{4.25}$$

is the point of intersection of the upsloping line through λ and (B, A) and the downsloping line $\{(f, s) : f + s = 2\}$.

(iv) For $(B, A) \neq (0, 0)$, the unique stationary distribution π for P is the point $(\pi_0, \pi_1) = (B, A)/(A + B)$ where the line from $(0, 0)$ to (B, S) intersects the line $\{(f, s) : f + s = 1\}$.

Proof. Part ((i)) is implied by the cases $\lambda = (0, 1)$ and $\lambda = (1, 0)$ of part ((ii)), and parts ((iii)) and ((iv)) also follow easily from part ((ii)). So it suffices to check part ((ii)). By the assumption that λ is a probability vector, $\lambda_0 + \lambda_1 = 1$. So the representation (4.18) of P in terms of A and B makes

$$(\lambda P)_1 = \frac{\lambda_0 A + \lambda_1 (A - 1)}{A + B - 1} = \frac{A - \lambda_1}{A + B - 1} \tag{4.26}$$

$$(\lambda P)_0 = \frac{\lambda_0 (B - 1) + \lambda_1 B}{A + B - 1} = \frac{B - \lambda_0}{A + B - 1} \tag{4.27}$$

In Cartesian coordinates (f, s) with horizontal coordinate f counting failures, that is values $X_i = 0$, and vertical coordinate s counting successes, that is values $X_i = 1$, the slope of the probability vector representing the distribution λP of X_2 is therefore

$$\frac{(\lambda P)_1}{(\lambda P)_0} = \frac{A - \lambda_1}{B - \lambda_0}$$

which is the slope of the line through the points $\lambda = (\lambda_0, \lambda_1)$ and (B, A). □

Remarks. Parts ((i)) and ((iv)) of the above corollary were pointed out by Bruno de Finetti in his study of exchangeable sequences of random variables. I do not know a reference for parts ((ii)) and ((iii)). It is a curious feature of this geometric construction of the vector $\lambda P(\,\cdot\,)$ from vectors representing $P(0, \cdot)$ and $P(1, \cdot)$, that the basic decomposition $\lambda P(\,\cdot\,) = \lambda_0 P(0, \,\cdot\,) + \lambda_1 P(1, \,\cdot\,)$ is not very apparent from the geometry. This makes it hard to give a comparably simple construction of the two inverse probability vectors $\mathbb{P}(X_1 \in \cdot \,|\, X_2 = j)$ for $j = 0, 1$ which are given by Bayes' rule:

$$\mathbb{P}(X_1 = i \,|\, X_2 = j) = \frac{\lambda_i P(i, j)}{(\lambda P)_j}. \tag{4.28}$$

Exercise. Show that if the probability vectors $P(0, \,\cdot\,)$ and $P(1, \cdot)$ are both drawn emanating from $(0, 0)$ (rather than from $(1, 0)$ and $(0, 1)$ as in Figure 4.3) , so the tips of both $P(0, \,\cdot\,)$ and $P(1, \,\cdot\,)$ fall on the downsloping line $\{(f, s) : f + s = 1\}$, then $\lambda P(\,\cdot\,)$ is the vector emanating from $(0, 0)$ whose tip is on the same downsloping line, a fraction λ_1 of the way along the directed line segment from $P(0, \,\cdot\,)$ to $P(1, \,\cdot\,)$. Embellish this diagram by making $\lambda P(\,\cdot\,)$ the top right corner of a parallelogram with two sides which are initial segments of the vectors $P(0, \,\cdot\,)$ and $P(1, \,\cdot\,)$. Each of the four terms $\lambda_i P(i, j)$ should now be apparent as a length on or other of the two axes.

Problem. How best to visualize Bayes' rule geometrically?

Figure 4.2: Transition vector diagram for a sample of size 2 without replacement.

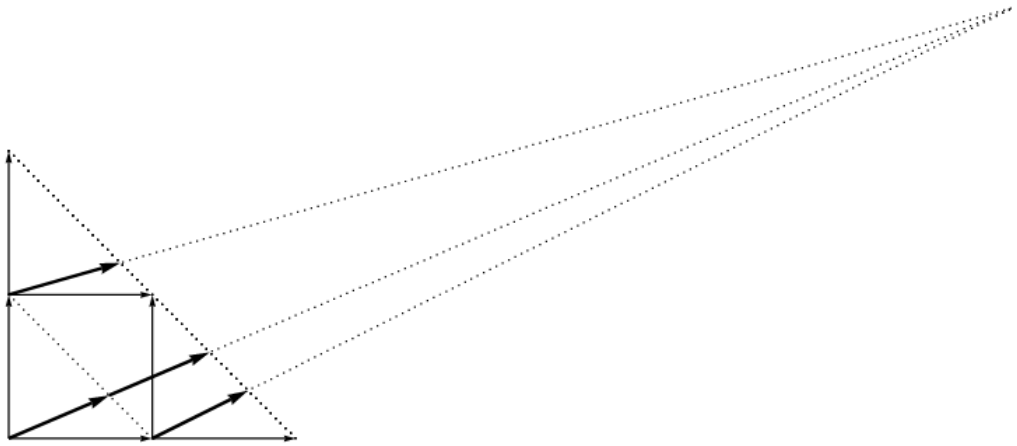

This diagram just amplifies the bottom left corner of the transition diagram of Figure 4.1 for sampling without replacement. The chain (W_0, W_1, W_2) makes 2 steps through 6 states (f, s) for non-negative counts f and s of failures and successes with $f + s \leq 2$, driven by (X_1, X_2) a sample of size 2 without replacement from 3 ones (successes) and 7 zeros (failures). Starting from $W_0 = (0, 0)$ the first transition is to $W_1 = (0, 1)$ (up) or $W_1 = (1, 0)$ (right) according to whether $X_1 = 1$ or 0. The next step to $W_2 = (X_1 + X_2, 1 - X_1 + 1 - X_2)$ is up or right according to whether $X_2 = 1$ or 0. After these two steps, W_2 is in one the states $(0, 2), (1, 1)$ or $(2, 0)$.

The transition probability vector

- out of $(0, 0)$ gives the stationary distribution $(7/10, 3/10)$ for X_1.
- out of $(0, 1)$ gives the distribution $(P_{10}, P_{11}) = (7/9, 2/9)$ of X_2 given $X_1 = 1$.
- out of $(1, 0)$ gives the distribution $(P_{00}, P_{01}) = (6/9, 3/9)$ of X_2 given $X_1 = 0$.

The distribution of X_2, which is identical to the distribution of X_1, is represented by copy of the vector for the stationary distribution of X_1, added to the tip of that vector. Observe that all the probability vectors point towards the point $(B, A) = (7, 3)$, representing the total numbers of failures and successes if the process of sampling without replacement is continued to an exhaustive sample of size 10.

Figure 4.3: Transition vector diagrams for two-state Markov chains.

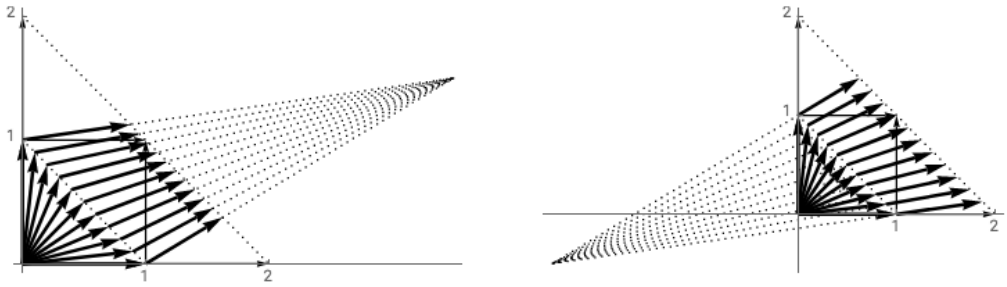

The left hand diagram shows the (f, s)-Cartesian plane for an indicator chain with $P_{01} = 3/8$ and $P_{11} = 1/8$ corresponding to $(B, A) = (7, 3)/2$. The geometric structure is very similar to that of Figure 4.2 for (X_1, X_2) a sample of size 2 without replacement from a population of 7 zeros and 3 ones. Now B and A are no longer integers, but the algebraic prescription of transition probabilities (4.18) still defines a 2×2 transition probability matrix. Here

- the transition vector out of $(1, 0) \longleftrightarrow (X_1 = 0)$ adds $P(0, \cdot) = (5, 3)/8$
- the transition vector out of $(0, 1) \longleftrightarrow (X_1 = 1)$ adds $P(1, \cdot) = (7, 1)/8$.

In accordance with Corollary 0.1, these transition vectors point to $(B, A) = (7, 3)/2$. The diagram shows the 11 initial probability vectors $\lambda = (i, 10-i)/10$ for $0 \leq i \leq 10$, emanating from the origin. Added to the tip of each of these vectors λ is the corresponding probability vector λP, which always points from λ to (B, A). The stationary probability vector is $\pi = (7, 3)/10$, the unique vector such that both λ and λP point directly to $(B, A) = (7, 3)/2$. The right hand diagram is the corresponding geometric description of the two state indicator chain with $P_{01} = 1/8$ and $P_{11} = 3/8$ corresponding to $(B, A) = (-5, -1)/2$. This diagram for positively dependent (X_1, X_2) is similar to the left hand diagram for negatively dependent (X_1, X_2), except that each vector λP added to λ points away from (B, A) instead of towards (B, A). Now the stationary probability vector is $\pi = (B, A)/(A + B) = (5, 1)/6$, which does not equal any of the displayed initial probability vectors $\lambda = (\lambda_0, \lambda_1)$, with λ_1 ranging over a multiples of $1/10$ as in the left hand diagram.

LECTURE 5

Recurrence Classes, x-Blocks, and Limit Theorem

5.1 Key Points for Homework

Pitman gives a few key pointers (which are from the textbook) that may help with finishing the homework due tonight.

- Recall the definition of an *irreducible* chain. That is,

$$\forall x, y \in \mathcal{S},\ \exists n\ : P^n(x, y) > 0$$

  This forbids a random walk on a graph with 2 or more components (closed classes). Most of the chains we commonly deal with (and in our homework) are irreducible.

- Fact: (See Theorem 1.7 in Durett). If P is irreducible and if there is a stationary probability vector π for P (that is, we can solve $\pi P = \pi$ where $\sum_x \pi(x) = 1, \pi(x) \geq 0$), then all the states are positive recurrent, i.e. the chain is positive recurrent.

5.2 Positive and Null Recurrence

Positive Recurrence

We say that an irreducible chain (or transition matrix) is *positive recurrent* when, for some or for all x

$$\mathbb{E}_x T_x < \infty$$

Note that

$$\mathbb{E}_x T_x = \sum_{n=1}^{\infty} \mathbb{P}_x(T_x \geq n).$$

You should check that if $\mathbb{E}_x T_x < \infty$ for some x and P is irreducible, then

$$\mathbb{E}_x T_x < \infty, \quad \forall x \in \mathcal{S}$$

This is closely related to the formula $\pi(x) = 1/\mathbb{E}_x T_x$

Null Recurrence

If a state is recurrent, but not positive recurrent (i.e. $\mathbb{P}_x(T_x < \infty) = 1$, but $\mathbb{E}_x T_x = \infty$), then we say that x is *null recurrent*.

5.3 Review: Mean Return Time

Pitman reminds us that there is a formula relating the mean return time and the stationary probability (Theorem 1.21 Durett):

$$\pi(x) = \frac{1}{\mathbb{E}_x T_x}$$

As a simple corollary, this formula directly implies that π is unique. There is no doubt about this for a stationary measure in terms of the mean recurrence time. If we discuss a system of countably infinite space, our traditional linear algebra may fail. This result provides an interpretation beyond a system of finitely many equations and unknowns.

Conversely, if P is irreducible and positive recurrent, then there exists this π. This is almost trivial, but of course we have to check that π is a stationary probability.

5.4 Example: Symmetric Random Walk

Consider a simple (symmetric) random walk with equal probability of going either direction on $\mathbb{Z} = \{\dots, -2, -1, 0, 1, 2, \dots\}$. We take the usual notation S_n for the walk. Start at $x = 0$, so that

$$S_n := \Delta_1 + \Delta_2 + \cdots + \Delta_n$$

where Δ_k is $+1$ or -1, each with probability $\frac{1}{2}$. This gives

$$P^n(0,0) = \begin{cases} 0 & \text{, if } n \text{ is odd} \\ \binom{2m}{m}\left(\frac{1}{2}\right)^{2m} & \text{, if } n = 2m \text{ is even} \end{cases}$$

Now Pitman notes we can tell recurrence or transience by looking at the fact that the total number of visits to 0 follows a geometric distribution with parameter $(1 - \rho_0)$

$$\mathbb{E}_0(\text{total \# visits to 0}) = \sum_{n=1}^{\infty} P^n(0,0)$$

But we know that $\binom{2m}{m}(\frac{1}{2})^{2m}$ is the same as the probability of m heads and m tails in $2m$ tosses. Increasing tosses gives a very "flat" normal curve because the mean of $\mathbb{E}_0 S_{2m} = 0$ and the variance tends to infinity, because the variance of each summed term is 1, the mean square is

$$\mathbb{E}_0 S_{2m}^2 = \underbrace{1 + 1 + \cdots + 1}_{2m} = 2m$$

We call this *diffusion*, in that on average the center of our distribution goes no where, but the distribution spreads out and flattens. Using Stirling's formula[1]

$$n! \sim \left(\frac{n}{e}\right)^n \sqrt{2\pi n}$$

and applying this to our earlier expression to show that

$$P^{2m}(0,0) \sim \frac{C}{\sqrt{m}}$$

where C is some constant and $\sim$ means the ratio tends to 1 as $n \to \infty$.

5.4.1 Recurrence versus Transience

To see recurrence versus transience, we look at (from earlier)

$$\sum_{n=1}^{\infty} P^n(0,0) = \sum_{m=1}^{\infty} P^{2m}(0,0) \sim \sum_{n=1}^{\infty} \frac{C}{\sqrt{m}} = \infty$$

(A rather paradoxical fact) This implies that the expected return time to 0 is infinite

$$\mathbb{E}_0 T_0 = \infty$$

although we are sure to eventually return with probability 1. Recall the definition of recurrent gives

$$\mathbb{P}_x(T_x < \infty) = 1 \iff \mathbb{P}_x(T_x \geq n) \to 0, \text{ as } n \to \infty.$$

Also, we should know that positive recurrence implies recurrence, but the converse is not necessarily true. Pitman summarizes that on our homework, we can quote the result: If we have a stationary probability measure, then the chain is *positive recurrent*.

5.5 Notion of x-Blocks of a Markov chain

Start at x (for simplicity) or wait until we hit x. Then look at the successive return times $T_x^{(i)}$ which is the i^{th} copy of T_x. Now recall this has the Strong Markov Property, which gives us two things:

[1] or Normal Approximation

(i) Every $T_x^{(i)}$ has the same distribution as T_x.

(ii) Further, they are independent copies. That is, $T_x^{(1)}, T_x^{(2)}, \ldots$ are independent.

Now Pitman mentions a variation on this theme of x-blocks, which explains many things.

5.5.1 Example: x-Blocks

Let $N_{xy}^{(i)} :=$ the # of visits to y in the i^{th} x-block of length T_x. In our previous in-class example, this gives a sequence

$$2, 0, 6, 0, 4, 2, \ldots$$

Now for some book keeping, consider what happens if we sum over all states y. Of course, this just gives the length of $T_x^{(i)}$ by "Accounting 101."

$$\sum_{y \in \mathcal{S}} N_{xy}^{(i)} = T_x^{(i)}$$

Note we must agree that $N_{xx}^{(i)} = 1$ for this to work. Now this implies that there is a formula involving expectations. Taking expectation starting at x

$$\sum_{y \in \mathcal{S}} \mathbb{E}_x N_{xy}^{(i)} = \mathbb{E}_x T_x^{(i)}$$

where this is really the same equation for all i by the Strong Markov Property. Fix x, y and look at $N_{xy}^{(1)}, N_{xy}^{(2)}, \ldots$, each of which

(i) $N_{xy}^{(i)}$ has the same distribution as $N_{xy} := N_{xy}^{(1)}$.

(ii) Further, the $N_{xy}^{(i)}$ are independent and identically distributed.

Pitman reminds us that as we return to x, via the Strong Markov Property, nothing of the past changes our expectations or distributions going forward.

5.6 Positive Recurrent Chains (P irreducible)

Notice that if $\mathbb{E}_x T_x < \infty$, and we define N_{xy} as we have earlier, then we can let

$$\begin{aligned} \mu(x, y) &:= \mathbb{E}_x(N_{xy}) \\ \mu(x) &:= \mathbb{E}_x T_x = \text{ mean length of } x\text{-block} \end{aligned}$$

Correspondingly to our Accounting 101, we write

$$\sum_{y \in \mathcal{S}} \mu(x, y) = \mu(x) < \infty$$

Further, we can show (see text for details) that if we sum

$$\sum_y \mu(x,y)P(y,z) = \mu(x,z)$$

or in other words, $\mu(x,\cdot)$ is a stationary measure, **not** a stationary probability, as it is an unnormalized measure). That is

$$\mu(x,\cdot)P = \mu(x,\cdot)$$

This is important because it gives us a simple explicit construction of a stationary measure $\mu(x,\cdot)$ for every state x in state space $\mathcal{S}$ of a positive recurrent (PR) irreducible chain with matrix P. Notice that this is not just any measure. By convention, we say that the number of times we visit x in the duration of T_x is 1 (this is necessary to satisfy our constructions today). That is, we must not count a visit twice, and we must set

$$\mu(x,x) := 1$$

in order to get

$$\sum_{y\in\mathcal{S}} \mu(x,y) = \mu(x) < \infty$$

Now to get a stationary probability measure, we take

$$\pi(y) = \frac{\mu(x,y)}{\sum_z \mu(x,z)} = \frac{\mu(x,y)}{\mu(x)}$$

and this does **not** depend on x. We can take any reference state and we get the same thing when we look at these ratios.

5.6.1 Explanation of the Key Formula

We may ask why we have

$$\sum_y \mu(x,y)P(y,z) = \mu(x,z)$$

Recall that $\mu(x,y)$ is the expected number of hits on y before T_x. That is,

$$\mu(x,y) = \mathbb{E}_x(\#\text{ of hits on } y \text{ before } T_x)$$

Now, every time we hit y, then $P(y,z)$ is the probability that the next step is to state z. Therefore, at least intuitively, $\mu(x,y)P(y,z)$ has a particular meaning. That is

$$\mu(x,y)P(y,z) = \mathbb{E}_x(\ \#\text{ of transitions } y\to z \text{ before } (\leq) T_x\ )$$

The distribution of a single x-block gives the following formulas for the invariant probability measure π

$$\pi(x) = \frac{1}{\mathbb{E}_x T_x}, \quad \frac{\pi(y)}{\pi(x)} = \mu(x,y)$$

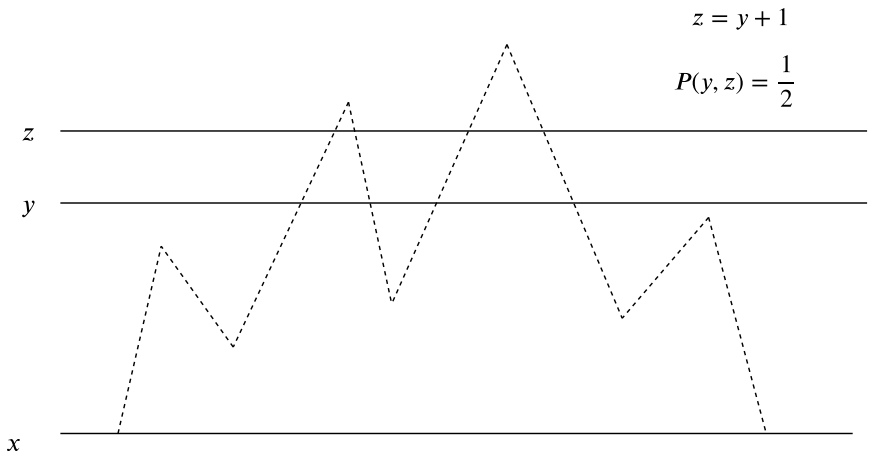

5.7 Limit Theorem

If we let $N_n(y) := \sum_{k=1}^{n} \mathbb{1}(X_k = y) = \#$ of hits on y in first n steps, then

$$\begin{aligned}\mathbb{E}_x \frac{N_n(y)}{n} &= \text{ mean } \# \text{ hits on } y \text{ per unit time up to } n \\ &= \frac{1}{n} \sum_{k=1}^{n} P^k(x, y) \to \pi(y)\end{aligned}$$

We have this Cesáro mean convergence always for irreducible positive recurrent chains, these themselves do not necessarily converge, but their average converges. Now if we additionally impose *aperiodicity*, we have

$$P^n(x, y) \to \pi(y)$$

always for irreducible and positive recurrent and aperiodic transition matrix P. See §1.8 in the text for the proof.

5.8 Review and Audience Questions

A null recurrent chain has a stationary measure with reference state x assigned as measure 1. If we do this on a simple symmetric random walk, we find that the expected time spent in every state on an x-block is 1, which explains why we expect to spend so much time to return back to x. If we have a stationary probability measure

$$\mathbb{E}_\pi \frac{1}{n} \sum_{k=1}^{n} P^k(x, y) = \pi(y)$$

we can argue that the stationary measure must be approached in the limit (at least in the Césaro sense) and is therefore unique as the limit must be unique.

LECTURE 6

First Step Analysis and Harmonic Equations

Pitman opens to questions regarding irreducible, aperiodic, recurrent (both positive and null), or transient. For a nice transition probability matrix, there exists a stationary probability π so that

$$\lim_{n\to\infty} P^n = \begin{pmatrix} \pi \\ \pi \\ \vdots \\ \pi \end{pmatrix} \tag{6.1}$$

Pitman asks us to recall the single most important formula regarding recurrent chain and its expected time for returning to x starting from x, $\mathbb{E}_x T_x$. For a nice (irreducible, positive recurrent),

$$\mathbb{E}_x T_x = \frac{1}{\pi(x)} \tag{6.2}$$

where $\pi(x)$ is the long run average (fraction of) time spent in state x. Recall that we defined that we hit x exactly once per x-cycle on average, which is equal to once per $\mathbb{E}$ cycle. This makes sense intuitively, where expecting to take a long time before returning to state x corresponds to not being in state x as often.

6.1 Hitting Places

Recall that we used the notation

$$T_A := \min\{n \geq 1 : X_n \in A\}$$
$$T_x := \min\{n \geq 1 : X_n = x\}$$

This is not trivial for X with $X_0 = x$. For analysis of hitting places (and time), it's often easier to have our discrete-time sequence start at 0. Hence we define

$$V_A := \min\{n \geq 0 : X_n \in A\} \tag{6.3}$$

Pitman notes that this is not a universal notation and we might see T, V, τ used for this definition, but for this text and course, we will use V_A for this purpose.

Thoerem 1.28 (Durett p. 55)

Consider a Markov chain with state space S. Take two non empty, (necessarily disjoint) $A, B \subseteq S$. Let $C := S - (A \cup B)$ and assume C is finite

Assumptions Suppose we have $h : S \to \mathbb{R}$ such that

$$h(a) = 1, \ \forall\, a \in A \tag{6.4}$$

$$h(b) = 0, \ \forall\, b \in B \tag{6.5}$$

$$h(x) = \sum_y P(x, y)h(y), \ \forall\, x \in C \tag{6.6}$$

Suppose also that

$$\mathbb{P}_x(V_{A\cup B} < \infty) > 0, \ \forall\, x \in C$$

Then

$$h(x) = \mathbb{P}_x(V_A < V_B) \tag{6.7}$$

The point of the theorem is that for a typical Markov chain X, and disjoint sets of states A and B, the chance that X hits A before B can be found, as a function of the initial state x, by solving the system of linear equations (6.6) subject to the obvious boundary conditions (6.4) and (6.5). Very commonly, we'll write the equation (6.6) in matrix notation, where h is a *column vector* as $h(x) = (Ph)(x)$ for $x \in C$. It is very convenient to assume (with no loss of generality) that both A and B are absorbing sets of states. Then (6.6) reduces to simply $h = Ph$, an equality of row vectors indexed by $x \in S$. This is because $h(x) = Ph(x)$ holds trivially for any absorbing state x (meaning $P(x, x) = 1$). Note that because C is assumed finite, a variant of Durrett's Lemma 1.3 shows that (6.1) is equivalent to

$$\mathbb{P}_x(V_{A\cup B} < \infty) = 1, \ \forall\, x \in C \tag{6.8}$$

For a chain with infinite state space S and C infinite, this condition is adequate as a replacement for (6.1), provided it is assumed that h is a bounded or non-negative function, so there are no problems with the definition of Ph.

Intuitively, regard A and B as sets of boundary states. Graphically, it is convenient to place A, the set of target states, at the top of a 2 (or higher dimensional) lattice, and place B as all three remaining boundaries of the lattice (left, right, bottom edges).

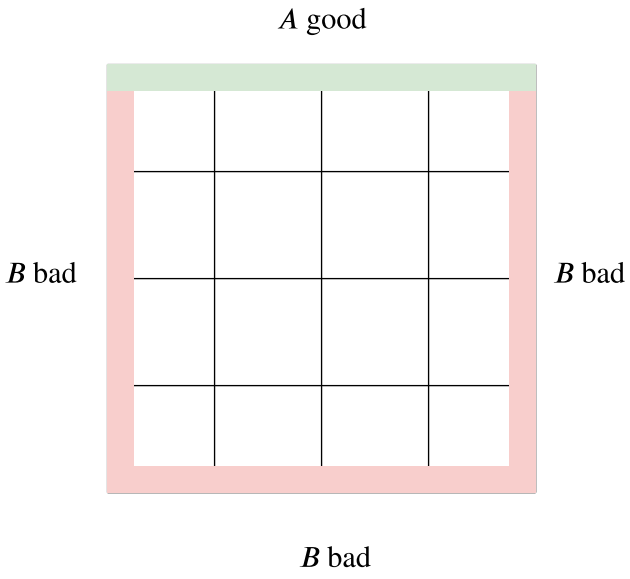

6.2 Method of Solution (Durret p.54)

Pitman notes that the method is more important than the solution here. From the text, "Let $h(x)$ be the probability of hitting A before B, starting from X." We call this technique **first step analysis**. "By considering what happens at the first step." That is, we assert that we start at $X_0 := x$, and we condition on time X_1, the value of the chain at time 1. Generalizing, let Y be any nonnegative (for simplicity) random variable, which is a function of $X_0, X_1, X_2, \ldots$ some Markov chain with transition matrix P. Consider $\mathbb{E}_x Y$ as a function of x, and notice we can write, by summing out all states $z \in S$, using $\sum_{z \in S} \mathbb{1}(X_1 = z) = 1$

$$\begin{aligned}
\mathbb{E}_x Y &= \mathbb{E}_x \sum_{z \in S} \mathbb{1}(X_1 = z) Y \\
&= \sum_z \mathbb{E}_x \big[\mathbb{1}(X_1 = z) Y \big] \\
&= \sum_z \mathbb{P}_x(X_1 = z) \mathbb{E}_x(Y \mid X_1 = z) \\
&= \sum_z P(x, z) \mathbb{E}_x(Y \mid X_1 = z)
\end{aligned}$$

which is simply computing the $\mathbb{P}_x$ expectation of Y by conditioning on X_1 Commonly, Y can be written as a function of $X_1, X_2, \ldots$ and this can be further simplified. We can do this for instance when Y is the indicato $Y = \mathbb{1}(V_A < V_B)$ in the setting of the above theorem. Then

$$\boxed{\mathbb{E}_x Y = \mathbb{P}_x(V_A < V_B)} \tag{6.9}$$

Something else is true as well via first step analysis. Take $x \notin A \cup B$. Look at the probability that V_A happens before V_B, provided that we know $X_1 = z$. Now if z is one of the boundary cases, this is trivial. So we treat in cases, using the Markov property

$$\mathbb{P}_x(V_A < V_B \,|\, X_1 = z) = \begin{cases} 1 & , z \in A \\ 0 & , z \in B \\ \mathbb{P}_z(V_A < V_B) & , \text{else} \end{cases}$$

as you should convince yourself.

> Does this probability of hitting A before B have anything to do with $P(c, \cdot)$ for $c \in A \cup B$?

We agree on the edge cases, for starting in A or B. Now we make this key observation, which is not mentioned in the text. Because of our definitions, namely the possibility of being there at time zero, the answer is NO!

With this in mind, we modify the problem at hand to make the entire set of states $A \cup B$ absorbing. That is, $P(c, c) := 1, \forall\, c \in A \cup B$. That is to say when we arrive, we stick there, and we solve the problem under these circumstances.
Notice that we agreed by conditioning on X_1 that

$$h(x) := \mathbb{P}_x(V_A < V_B), \text{ for} x \notin A \cup B$$

solves the *Harmonic equation*

$$\boxed{h(x) = \sum_y P(x, y) h(y)} \tag{6.10}$$

Notice that if we make $A \cup B$ absorbing, then this harmonic equation above is true for ALL $x \in A \cup B$. Now we arrive at a reformulation of the theorem.

Pitman's Version of Durett's Theorem

Assume that

- P has $A \cup B$ as absorbing states and
- $\mathbb{P}_x(\text{hit } A \cup B \text{ eventually}) = 1, \forall\, x \in S$

then

$$h(x) := \mathbb{P}_x(\text{hit } A \text{ before } B)$$

is the *unique* bounded or non-negative solution of $h = Ph$, subject to the *boundary condition* that $h = \mathbb{1}_A$ (the indicator of A) on $A \cup B$.

This is fundamentally the same as Durett's theorem, but with some tinkering, we have a more elegant statement as here. Notice that $h = Ph$ is a very special equation,

whose as solutions solve various problems. In order to understand this equation, it is important to understand what is Pf for a function (column vector) f (assume either nonnegative or bounded so that we can make sense of the summations). Then the action of the transition matrix P on a column vector f gives us:

$$\big(Pf\big)(x) = \sum_{y \in S} P(x,y) f(y),$$

summing over all y in the state space. $P(x,y)$ gives the probability distribution over values y, depending on the initial state x and $f(y)$ simply gives the return from state y. Hence directly by our notation, we have:

$$\big(Pf\big)(x) = \mathbb{E}_x f(X_1).$$

Hence

$$\big(Ph\big)(x) = \mathbb{E}_x h(X_1)$$

as the meaning of $(Ph)(x)$. Another way to say this is by looking at the conditional expectation (knowing X_0)

$$\mathbb{E}\big[h(X_1) \,|\, X_0\big] = \big(Ph\big)(X_0)$$

Pitman makes the following claim: If $h = Ph$ (that is, h solves the harmonic equation), then the expectation (starting at x) of h of any variable (X_n) is

$$\mathbb{E}_x\big[h(X_n)\big] = h(x)$$

which is true by $n = 1$ by $(Ph)(x) = \mathbb{E}_x h(X_1)$ from above (that is, $h = Ph$). Now, this is true for $n = 1, 2, 3, \ldots$ by induction and the Markov property. If we trust this for now (we may revisit this later), we may want to assume that

$$h = \begin{cases} 1, & \text{on } A \\ 0, & \text{on } B, \end{cases}$$

then we can write

$$h(x) = \mathbb{E}_x h(X_n) = \sum_{y \in S} P^n(x,y) h(y),$$

as our familiar notation for a Markov chain. Then we can equivalently write this as a summation over the three state cases

$$h(x) = \sum_{y \in A} P^n(x,y) h(y) + \sum_{y \in B} P^n(x,y) h(y) + \sum_{y \in S - A - B} P^n(x,y) h(y)$$

Recall that we've set $A \cup B$ to be absorbing, so the first two terms are simply

$$\sum_{y \in A} P^n(x,y) h(y) = \mathbb{P}(V_A \leq n)$$

$$\sum_{y \in B} P^n(x,y) h(y) = 0$$

Hence

$$h(x) = \mathbb{P}_x(V_A \leq n) + 0 + \sum_{y \in S-A-B} P^n(x,y)h(y)$$

Now if we take $n \to \infty$, then

$$\lim_{n\to\infty} h(x) = \lim_{n\to\infty} \mathbb{E}_x h(X_n) = P_x(V_A < \infty) + \underbrace{\lim_{n\to\infty} \sum_{y \in S-A-B} P^n(x,y)h(y)}_{=0}$$

$$= \mathbb{P}_x(V_A < \infty)$$

because $\mathbb{P}_x(\text{hit } A \cup B \text{ eventually}) = 1$ via our assumption, and the sum which tends to 0 is bounded above by the maximum absolute value of $h(y)$ over $y \in S - A - B$ times $\mathbb{P}_x(V_{A\cup B} > n)$ which tends to 0 by the assumption that $\mathbb{P}_x(V_{A\cup B} < \infty) = 1$.

6.3 Canonical Example: Gambler's Ruin for a Fair Coin

The state space is $S := \{0, 1, 2, \ldots, N\}$, and the goal state is $A = \{N\}$, and the bad state is $B = \{0\}$. The transition matrix is then

$$P = \begin{bmatrix} 1 & 0 & 0 & \cdots & \\ \frac{1}{2} & \frac{1}{2} & 0 & \cdots & \\ 0 & \frac{1}{2} & \frac{1}{2} & & 0 \cdots \\ 0 & 0 & 0 & \cdots & 1 \end{bmatrix}$$

Now let (X_n) be the simple random walk with absorbing states $\{0, N\}$. Then

$$\mathbb{P}_x(\text{hit } N \text{ before } 0) = h(x)$$

is desired. The equation $h = Ph$ becomes

$$h(x) = \frac{1}{2}h(x+1) + \frac{1}{2}h(x-1) \text{ for } 0 < x < N$$

and we set the boundary conditions

$$h(N) := 1 \qquad h(0) := 0$$

Now the harmonic equation $h(x) = \frac{1}{2}h(x+1) + \frac{1}{2}h(x-1)$ says that the graph of $h(x)$ is a straight line, over integer values x, passing through 0 and 1. Hence $h(x) = \frac{x}{N}$ is the unique solution to this system of equations. Here it is easy that we are certain to eventually hit the boundary states. Hence for the simple symmetric random walk started at $0 \leq x \leq N$

$$\mathbb{P}_x(\text{hit } N \text{ before } 0) = \frac{x}{N}$$

which is a famous result as due to Abraham de Moivre around 1730.

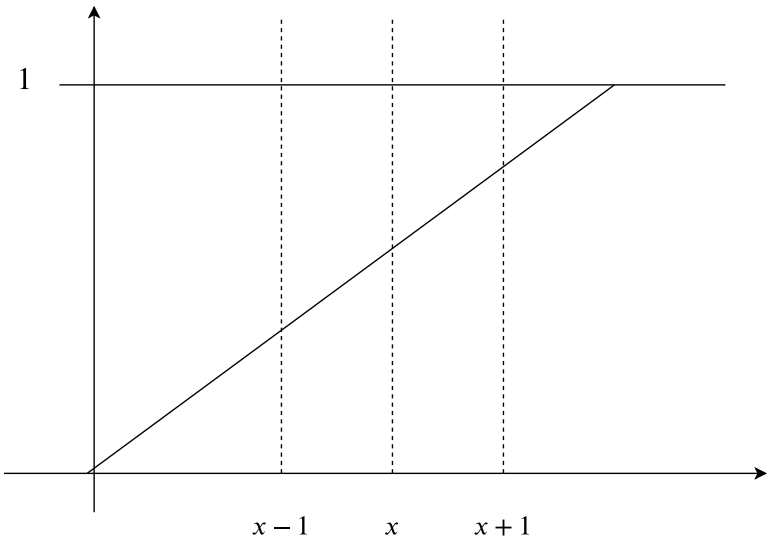

6.3.1 Gambler's Ruin with a Biased Coin

What are the harmonic equations? We reason that this results in the same equations, with slight modifications

$$h(x) = ph(x+1) + qh(x-1) \quad (0 < x < N)$$

which we may solve via algebra as done in Durett (p. 58). Pitman shows us a more clever way, related to the idea of a *martingale*. There is a discussion of this problem in the context of martingales at the end of the text, as aspects of a hitting-time problem (we will revisit this at the end of the course). Observe that $h(x) = x$ is no longer harmonic when $q \neq p$ (biased coin). Now by some good guesswork you can discover that

$$h(x) = \left(\frac{q}{p}\right)^x$$

is a harmonic function for the $p - q$ walk. We check this

$$\begin{aligned} Ph(x) &= p\left(\frac{q}{p}\right)^{x+1} + q\left(\frac{q}{p}\right)^{x-1} \\ &= \left(\frac{q}{p}\right)^x = h(x) \end{aligned}$$

This is a bit clever, but it is not a bad idea to try a solution of the form $h(x) = r^x$ of the harmonic equations, and if you do that you will get a quadratic which forces $r = 1$ (boring) or $r = q/p$ (very useful) as above. As soon as we have found this

$h(x)$, we can argue as before: from $h = Ph$ get $h = P^n h$ and so for each $n \geq 0$

$$h(x) = \mathbb{E}_x \left(\frac{q}{p}\right)^{X_n}$$

$$= \left(\frac{q}{p}\right)^{N} \mathbb{P}_x(\text{hit } N \text{ before } n) + \left(\frac{q}{p}\right)^{0} \mathbb{P}_x(\text{hit } 0 \text{ before } n) + \sum_{y \notin \{0,N\}} \cdots$$

Now taking $n \to \infty$, this final term goes to zero. Hence in the limit
and additionally

$$\mathbb{P}_x(\text{hit } N) + \mathbb{P}_x(\text{hit } 0) = 1$$

Now we have two equations and two unknowns. Solve these, and you get the solution found by Durrett on p.58.

LECTURE 7

First Step Analysis Continued

7.1 First Step Analysis: Continued

The simple idea here is to derive equations by conditioning on step 1. We can find all sorts of things about Markov chains by doing exactly this. Pitman notes that the text keeps doing this technique without explicitly pointing it out. Recall that first step analysis for a Markov chain $(X_0, X_1, X_2, \dots)$, we consider some random variable

$$Y = Y(X_0, X_1, X_2, \dots)$$

If we know $\mathbb{E}_x Y$ for all states x and we want to compute the expectation of Y for a chain with X_0 assigned a probability distribution $\lambda = \lambda(x) \;\; x \in S$, denoted $\mathbb{E}_\lambda Y$, we would take

$$\mathbb{E}_\lambda Y = \sum_{x \in S} \lambda(x) \mathbb{E}_x Y$$

Put simply, the expectation of a random variable Y is the expectation of the expectation of Y conditioned on X_0. That is,

$$\mathbb{E}(Y) = \mathbb{E}\big[\mathbb{E}(Y \mid X_0)\big].$$

We may want to condition on X_1 as well, which is how we derived the harmonic equations from the previous lecture. Let's look at an example where we can do this again.

7.1.1 Example: Mean Hitting Times

Suppose we have a set of states A (we can make them absorbing as a matter of technique, it makes no difference to the answer), and consider

$$V_A := \min\{n \geq 0 : X_n \in A\}.$$

We want to find $\mathbb{E}_x V_A$ for any initial state x. If $x \in A$, then we trivially have $\mathbb{P}_x(V_A = 0) = 1$ and hence $\mathbb{E}_x V_A = 0$. Now, for any state x define a function for the mean

$$m(x) = m_A(x) := \mathbb{E}_x V_A$$

where we drop the subscript A as it is understood from context. We want equations for $m(x)$.

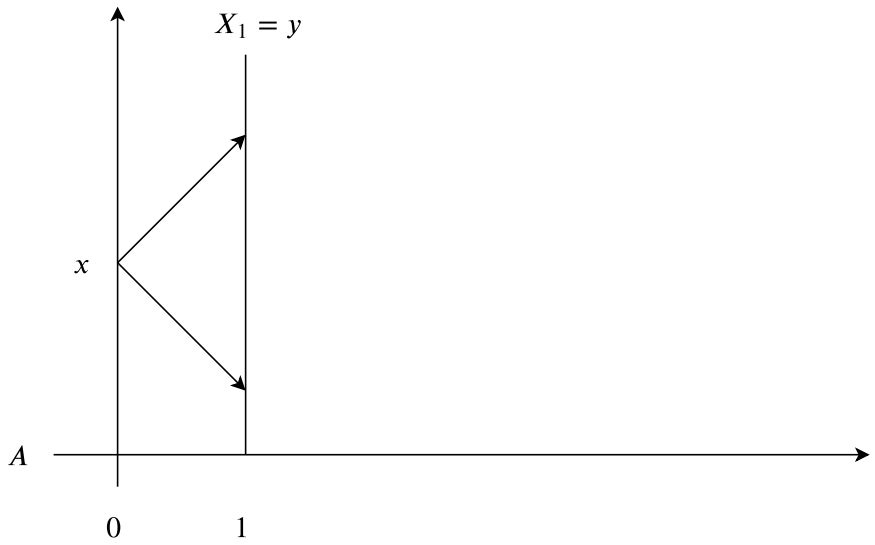

From x, we hit $X_1 = y$ with probability $P(x, y)$. Now given $X_0 = x$, $X_1 = y$, for $x \notin A$ we have

$$\mathbb{E}(V_A \,|\, X_0 = x, X_1 = y) = 1 + \mathbb{E}_y(V_A)$$

Notice that this is correct if $y \in A$. If we happen to hit A =at time 1, then $V_A = 1$ and the second term $\mathbb{E}_y(V_A)$ is zero. Additionally, this is correct if $y \notin A$, that is, because $x \notin A$ we are certain to take at least 1 step, with $\mathbb{E}_x V_A \geq 1$. This means that we can write down a system of equations, relating to the mean times

$$m(x) = 1 + \sum_{y \in S} P(x, y) m(y) \qquad (x \notin A)$$

This system should be solved together with the *boundary condition*

$$m(x) = 0 \qquad (x \in A)$$

If we have only a finite number of non-absorbing states, then we have a finite number of linear equations and this number of unknowns.

In the text, Theorem 1.29 on page 62 states that as long as we can reach the boundary from the any state in the interior (in some number of steps) with positive probability, provided there are only a finite number of interior states, this system of equations will have a unique solution. In practice, in examples, you just write down the system of linear equations and solve them by standard methods or software.

7.1.2 Application: Duration of a Fair Game

The usual Gambler's Ruin for a fair coin. Text Example 1.52 on page 66, We start with $\$x$ and play for $\pm\$1$ gains with equal probability until we hit either \$0 or some

$\$N$. Last lecture, we showed

$$\mathbb{P}_x(\text{reach } N \text{ before } 0) = \frac{x}{N}$$

Set $A := \{0, N\}$ as our absorbing states and V_A the duration of the game, where

$$V_A := \min\{n \geq 0 : X_n \in A\}$$

Recall that there remains the scenario of never hitting the boundary A, but we have already found before that the probability assigned to this uncountable infinite number of never-ending paths is zero.

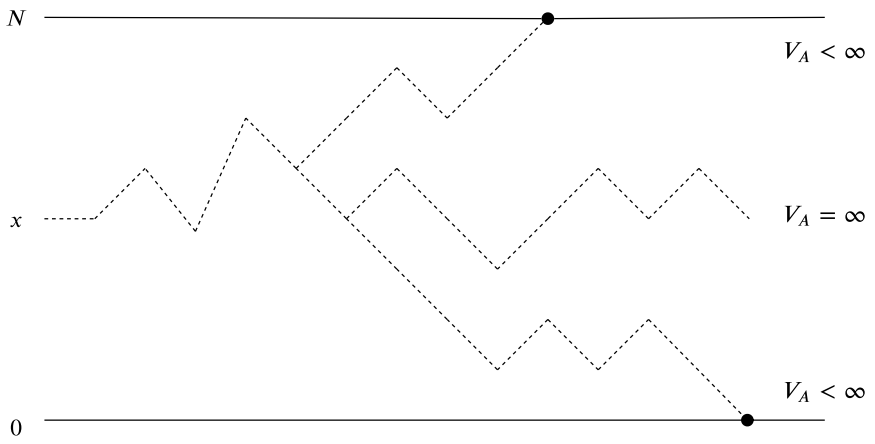

To see this, notice that for any 'block' of N steps, there is a strictly positive probability that we hit a boundary state. That is $\mathbb{E}_x V_A < \infty$ with probability 1 and for all $x \in S$.

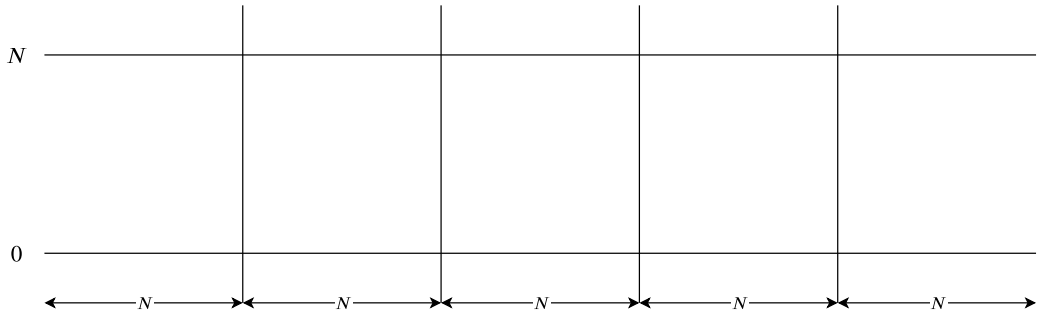

We use this argument to form the geometric bound as we have before in a previous lecture. To find $m(x) := \mathbb{E}_x V_A$ we first write out the boundary conditions. That is,

$$m(0) = m(N) = 0$$

Now the nontrivial cases, we again break into two parts

$$m(x) = 1 + \frac{1}{2}m(x+1) + \frac{1}{2}m(x-1) \quad 0 < x < N$$

Then we solve for this system of equations. Recall that previously we considered the simpler system $h(x) = \frac{1}{2}h(x+1) + \frac{1}{2}h(x-1)$, which implies that $h(x)$ is linear, more pedantically *affine*. For constants a, b, we have

$$h(x) = ax + b$$

Now consider the addition of a quadratic term:

$$m(x) = cx^2 + ax + b$$

Then we observe

$$\frac{1}{2}c(x+1)^2 + \frac{1}{2}c(x-1)^2 = cx^2 + \underbrace{\frac{1}{2}c(2x) + \frac{1}{2}c(-2x)}_{=0} + c$$

Now from this sort of consideration, $m(x)$ as above solves the equation

$$\frac{1}{2}m(x+1) + \frac{1}{2}m(x-1) = c + m(x)$$

Hence, we conclude that our system of equations is solved by a quadratic function of the form $m(x) = cx^2 + ax + b$, where $c = -1$.

7.1.3 Summarizing our Findings

Observe

$$g_1(x) = ax + b \implies \frac{1}{2}g_1(x+1) + \frac{1}{2}g_1(x-1) = g_1(x)$$

$$g_2(x) = cx^2 \implies \frac{1}{2}g_2(x+1) + \frac{1}{2}g_2(x-1) = g_2(x) + c$$

These together imply

$$g(x) = cx^2 + ax + b = (g_1 + g_2)(x) \implies \frac{1}{2}g(x+1) + \frac{1}{2}g(x-1) = g(x) + c$$

Hence we have that

$$m(x) := cx^2 + bx + a$$

solves our equations from earlier if and only if $c = -1$. Then plugging this in, we have

$$m(x) = -x^2 + bx + a$$

and additionally recall that $m(0) = m(N) = 0$. There's only one quadratic that satisfies these, namely

$$m(x) = -x(x - N) = \boxed{x(N - x)}$$

In summary, with the idea to try a quadratic (which Pitman notes is not too different from noticing before that a harmonic function for the fair gambler's ruin chain must be linear), finding the exact solution is not hard. See text for solution of the mean duration of an unfair game, and many further examples.

7.2 Conditioning on other variables

Commonly in the analysis of Markov chains it is effective to condition on X_0 or on X_1. Also common to condition on X_n (example in homework). Now we would like to consider that these may not be the only variables on which we would like to condition. There may be more clever techniques, where we employ our imagination to find a more apt conditioning variable, often a suitable random time. Also, exploiting the addition rule for expectation, after breaking a random variable into a sum of two variables, should be kept in mind.

7.2.1 Runs in Independent Bernoulli(p) Trials

We want to find the mean time until we see N successes in a row. Let τ_N be the random number of trials required. e.g. for $N = 3$ if the outcome of the trials is

$$(X_1, X_2, \ldots) = (0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, \ldots)$$

then $\tau_N = 11$. Note that in treating Bernoulli trials, and more generally i.i.d. sequences, it is customary to start indexing by 1, whereas for general discussion of Markov chains it is customary to start indexing by 0. Python programmers should like the Markovian convention.

The exact distribution of τ_N is tricky. You can find its generating function if you like, but we only want the expectation here, which is relatively simple. Of course, because $\mathbb{P}(\tau_N \geq N) = 1$

$$\begin{aligned}\mathbb{E}\tau_N &= \sum_{k=N}^{\infty} k\ \mathbb{P}(\tau_N = k) \\ &= \sum_{k=N}^{\infty} \mathbb{P}(\tau_N \geq k)\end{aligned}$$

but neither the point probabilities in the first equality nor the tail probabilities needed for the second tail sum formula sum have a simple formula. Hence we ask, "What should we condition on?" As a student suggests, try τ_{N-1}. That is, having a string of $N-1$ ones in a row.

$$\tau_N = \tau_{N-1} + \Delta_N, \quad \text{where } \Delta_N = \begin{cases} 1 \text{ with probability } p \\ 1 + \text{a copy of } \tau_N \text{ otherwise} \end{cases}$$

If you are eager to see the run of N ones, you are disappointed if the trial following trial τ_{N-1} is a 0, as this means you must start over again. However, this *regeneration* of the problem is exactly what is needed to help find the sequence of means μ_N

Define $\mu_N := \mathbb{E}\tau_N$, then the above observation gives

$$\mu_N = \mu_{N-1} + 1 + q\mu_N$$

where rearranging gives

$$\mu_N = \frac{\mu_{N-1} + 1}{p}$$

We test this

$$\mu_1 = \frac{1}{p}$$

by the mean of geometric. Similarly,

$$\mu_2 = \frac{\left(\frac{1}{p} + 1\right)}{p} = \frac{1+p}{p^2}$$

and repeating this gives

$$\mu_N = \frac{1 + p + p^2 + \cdots + p^{N-1}}{p^N}$$

In summary, we solved this problem by noticing that to get to N in a row, we needed to first get to $N-1$ in a row, and then condition on the next trial. Here is another approach:

7.2.2 Conditioning on the First Zero

Define G_0 as the first $n \geq 1$ such that $X_n = 0$ (that is, wait for the first 0). In other words, G_0 is one plus the length of the first run of 1s. Then $G_0 \sim \textbf{Geometric}(q)$, where q is the failure probability. It seems reasonable to try to find $\mathbb{E}\tau_N$ by conditioning on G_0, as G_0 is closely related to τ_N, and we know the distribution of G_0. If $G_0 > N$, then $\tau_N = N$. On the other hand, if $G_0 = g \leq N$, then the problem starts over: there is the equality in distribution

$$(\hat{\tau}_N - g \,|\, G_0 = g) \stackrel{d}{=} \tau_N \qquad (0 < g < N)$$

meaning that conditional given $G_0 = g$) the remaining time $\tau_N - g$ has the same distribution as τ_N. This is by a rather obvious form of the Strong Markov Property for Bernoulli trials.

Therefore, by conditioning on G_0, we have

$$\mathbb{E}\tau_N = \left[\sum_{g=1}^{N} \mathbb{P}(G_0 = g)(g + \mathbb{E}\tau_N)\right] + \mathbb{P}(G_0 > N)N$$

Now let $\mu_N := \mathbb{E}\tau_N$, so that the earlier equation gives

$$\mu_N = \sum_{g=1}^{N} p^{g-1} q(g + \mu_N) + p^N N$$

We look at a simple $N = 2$ case. Here in this solution, we have:

$$\mu_2 = p^0 q(1 + \mu_2) + pq(2 + \mu_2) + p^2 \cdot 2$$
$$\mu_2(1 - q - pq) = q + 2pq + 2p^2$$

hence easily $\mu_2 = (1 - p)/p^2$ as before. You can easily check this method gives the same conclusion as before for general N.

LECTURE 8

Infinite State Spaces and Probability Generating Functions

8.1 Infinite State Spaces

§1.11 is starred in the text, but is not optional for our course. We'll discuss techniques for both finite and infinite state spaces, in particular

- probability generating functions
- potential kernel (AKA) Green matrix

Pitman gives a list of additional resources with nice problems worth trying. See Bibliography for further details.

[3] Grimmett, Geoffrey R. and Stirzaker, David R. *Probability and Random Processes*

[4] Asmussen, S{}o ren *Applied Probability and Queues*

[5] Norris, J. R. *Cambridge Series in Statistical and Probabilistic Mathematics*

[6] Feller, William. *An Introduction to Probability Theory and its Applications*

8.2 Review of Mathematics: Power Series

Know the following by heart, because it'll be on the midterm.

8.2.1 Binomial Theorem

The most important case of the binomial expansion

$$(1+x)^n = \sum_{k=0}^{n} \binom{n}{k} x^k$$

where we should observe

$$\binom{n}{k} = \frac{n!}{(n-k)!k!} = \frac{n(n-1)\cdots(n-k+1)}{k(k-1)\cdots 1}$$

and it is an important insight that the numerator is a polynomial in n. Pitman comments that no one realized why this is important until about 1670. The reason is that this form can be extended to other powers, namely $n := -1, \frac{1}{2}, \frac{-1}{2}$, or any real number $n \to r \in \mathbb{R}$. Now look at

$$(1+x)^r = \sum_{k=0}^{\infty} \binom{r}{k} x^k \qquad (|x| < 1)$$

which is valid for all real r and all real or complex x with $|x| < 1$. Notice that the combinatorial meaning of r 'choose' k makes sense only for $r = n$ a positive integer and k a non-negative integer. But the meaning of $\binom{r}{k}$ is extended to all real numbers r and all non-negative integers k by treating $\binom{n}{k}$ as a polynomial of degree k in n, then substituting r in place of n in this polynomial.
This is the instance with $f(x) = x^r$ of the *Taylor expansion* of a function f about the point 1

$$f(1+x) = f(1) + f'(1)x + \frac{f''(1)}{2!}x^2 + \cdots$$

which for suitable f is valid for $|x| < R$, where R is the radius of convergence. Usually for our purposes, $R \geq 1$. As another Taylor expansion (around 0 instead of 1)

$$e^x = \sum_{k=0}^{\infty} \frac{x^k}{k!}$$

We get exponentials arising as limit of binomial probabilities (e.g. the Poisson distribution). Also, recall that the geometric distribution converges to the exponential distribution with suitable scaling.

8.3 Probability Generating Functions

Suppose we have a non-negative integer-valued random variable X, which for simplicity will have non-negative integer values $X \in \{0, 1, 2, \dots\}$.

Probability Generating Function (PGF)

The *probability generating function* for a discrete $X \in \{0, 1, 2, \dots\}$ is

$$\phi_X(z) := \mathbb{E}z^X$$

We usually take $0 \leq z \leq 1$. When discussing PGFs, we may push z to $|z| \leq 1$, but in this course we will work entirely with PGFs defined as a function of an argument $z \in [0, 1]$.

Then $\phi_X(z) \in [0,1]$ too, and there are many contexts in which $\phi_X(z)$ acquires meaning as the probability of something. Now we can write the above as a power series. Recall that

$$\mathbb{E}g(X) = \sum_{n=0}^{\infty} \mathbb{P}(X=n)g(n) \tag{8.1}$$

so

$$\phi_X(z) := \mathbb{E}Z^X = \sum_{n=0}^{\infty} \mathbb{P}(X=n)z^n = \sum_{n=0}^{\infty} P_n z^n$$

where $p_n := \mathbb{P}(X=n)$. We worked with PGFs very briefly in a previous lecture, for dice probabilities, namely taking X uniform on $\{1,2,3,4,5,6\}$, and we looked at

$$\phi_X(z) = \frac{1}{6}(z + \ldots + z^6)$$

Recall this is where Pitman asked us to look at powers of this expansion in Wolfram Alpha. Notice that by convention, $0^0 = 1$, so $\phi_X(0) = \mathbb{P}(X=0)$. Now for any PGF, we have

$$\begin{aligned}
\frac{d}{dx}\phi_X(z) &= \frac{d}{dz}\sum_n \mathbb{P}(X=n)z^n \\
&= \sum_n \mathbb{P}(X=n)\frac{d}{dz}z^n \\
&= \sum_n \mathbb{P}(X=n)nz^{n-1}
\end{aligned}$$

and so we see that

$$\mathbb{E}X = \frac{d}{dz}\phi_X(z)\Big|_{z=1^-}$$

where we must approach $z=1$ from the left if the radius of convergence R is exactly $R=1$, but typically $R>1$ and you can just evaluate the derivative at $z=1$.
Perhaps we'd like to compute the variance. We ask, what happens if we differentiate twice?

$$\left(\frac{d}{dz}\right)^2 \phi_X(z) = \sum_{n=0}^{\infty} \mathbb{P}(X=n)n(n-1)z^{n-2}$$

Again we'd like the z factor to go away, so we set $z := 1$ and we have

$$\begin{aligned}
\mathbb{E}\left[X(X-1)\right] &= \sum_{n=0}^{\infty} \mathbb{P}(X=n)n(n-1) \\
&= \left(\frac{d}{dz}\right)^2 \phi_X(z)\Big|_{z=1^-}
\end{aligned}$$

Recall that $X_\lambda \sim$ **Poisson** (λ) if and only if:

$$\mathbb{P}(X_\lambda = n) = \frac{e^{-\lambda}\lambda^n}{n!},$$

which via the generating function implies

$$\phi_{X_\lambda}(z) = \sum_{n=0}^{\infty} \frac{e^{-\lambda}\lambda^n z^n}{n!} = e^{-\lambda}e^{\lambda z} = e^{\lambda(z-1)}$$

Easily from the above analysis by d/dz, or otherwise,

$$\begin{aligned} \mathbb{E}X_\lambda &= \lambda \\ \mathbb{V}\mathrm{ar}(X_\lambda) &= \lambda \end{aligned}$$

A (good) question arises whether $\phi_X(z)$ is a probability. The answer is yes, because after all the range of values is between 0 and 1, and any such function can be interpreted as a probability. Notably, we have

$$\phi_X(z) = \mathbb{P}(X \leq G_{1-z})$$

where G_p for $0 \leq p \leq 1$ denotes a random variable independent of X with the geometric (p) distribution on $\{0, 1, \ldots\}$: for $n \geq 0$
Then

$$\mathbb{P}(G_p = n) = (1-p)^n p, \text{ and } \mathbb{P}(G_p \geq n) = (1-p)^n.$$

In summary, we can think of a probability generating function as a probability, and we only need that G_{1-z} is independent of X.
Now if X, Y are independent, then

$$\begin{aligned} \mathbb{E}z^{X+Y} &= \mathbb{E}\left[z^X z^Y\right] \\ &= \left[\mathbb{E}z^X\right]\left[\mathbb{E}z^Y\right] \\ &= \phi_X(z)\,\phi_Y(z) \end{aligned}$$

Hence the PGF of a sum of independent variables is the product of their PGFs.

Example: Let $G_p \sim$ **Geometric**(p) on $\{0, 1, 2, \ldots\}$. Then

$$\mathbb{P}(G_p = n) = (1-p)^n p, \text{ for } n = 0, 1, 2, \ldots$$

Now if we want to look at the probability generating function, we have

$$\mathbb{E}(z^{G_p}) = \sum_{n=0}^{\infty} q^n p z^n = \frac{p}{1-qz}$$

for $p + q = 1$ and $|z| < 1$. Now we look at

$$T_r := G_1 + G_2 + \cdots + G_r$$

where $r = 1, 2, 3, \ldots,$ and G_i are all independent geometrically distributed with the same parameter p. The interpretation is to see G_p as the number of failures before the first success. That is, the number of 0s before the first 1 in independent **Bernoulli**()) 0,1 trials. Then similarly,

$$T_r = T_{r,p} = \text{ number of 0s before } r^{\text{th}} \text{ 1 in indep. } \textbf{Bernoulli}(p) \text{ 0,1 trials}$$

Looking at i.i.d. copies of G_p we use generating functions

$$\begin{aligned}
\mathbb{E}z^{T_r} &= \left(\frac{p}{1-qz}\right)^r = p^r(1-qz)^{-r} \\
&= p^r\left(1+(-qz)\right)^{-r} \\
&= \sum_{n=0}^{\infty}\binom{-r}{n}(-qz)^n, \\
&= p^r\sum_{n=0}^{\infty}\frac{(r)_{n\uparrow}}{n!}q^n z^n
\end{aligned}$$

where we simply plug into Newton's binomial formula. Notice this is

$$\mathbb{E}z^{T_r} = p^r\sum_{n=0}^{\infty}\frac{(r)_{n\uparrow}}{n!}q^n z^n$$

where

$$(r)_{n\uparrow} := r(r+1)\cdots(r+n-1)$$

$$\frac{(r)_{n!}}{n!} = \binom{r+n-1}{n}$$

From 134, we know this to be the negative binomial distribution. The above formula can be derived directly by counting: $\binom{r+n-1}{n}$ is the number of ways to place the n failures in the first $r+n-1$ trials, and the last $(n+r)^{\text{th}}$ trial must be a 1. But the generating function technique used above is instructive, and can be applied to more difficult problems.

8.4 Probability Generating Functions and Random Sums

Suppose we have $Y_1, Y_2, \ldots$ i.i.d. non-negative integer random variables, with probability generating function

$$\phi_Y(z) = \mathbb{E}z^{Y_k} = \sum_{n=0}^{\infty}\mathbb{P}(Y_k = n)z^n$$

the same generating function for all Y_k. Now consider another random variable, $X \geq 0$, integer valued, assumed independent of the sequence of $Y's$, and look at:

$S_X = Y_1 + Y_2 + \ldots + Y_X$, the sum of X independent copies of Y. Then

$$S_n = Y_1 + \cdots + Y_n$$
$$S_X = Y_1 + \cdots + Y_X$$

Now if $X = 0$ with 0 copies of Y, then our convention is to set the empty sum to give 0. We wish to find the PGF of S_X. The random index X is annoying, so try conditioning on it

$$\begin{aligned}
\mathbb{E}z^{S_x} &= \sum_{n=0}^{\infty} \mathbb{P}(X = n)\mathbb{E}\left(z^{S_n}\right) \\
&= \sum_{n=0}^{\infty} \mathbb{P}(X = n)\left[\phi_Y(z)\right]^n \\
&= \phi_X\left[\phi_Y(z)\right]
\end{aligned}$$

which is a composition of generating functions. In the middle line, recognize this is a generating function, just evaluated at a different location. Notice that for this to hold, we needed to assume that X is independent of $Y_1, Y_2, \ldots$. Also, there are some easy consequences for moments which you can derive directly or by generating functions, especially $\mathbb{E}S_X = (\mathbb{E}Y)\mathbb{E}X$ and you can get a formula for $\mathbb{E}S_X^2$ and hence the variance of S_X.

8.5 Application: Galton-Watson Branching Process

Assume that we're given some probability distribution (offspring distribution)

$$p_0, p_1, p_2, \ldots$$

Start with some fixed number k of individuals in generation 0, where each of these k individuals has offspring with distribution according to X. Our common notation is

$$Z_n := \#\text{ of individuals in generation } n$$

and so we have the following equality in distribution

$$(Z_1 \mid Z_0 = k) \stackrel{d}{=} X_1 + X_2 + \cdots + X_k$$

where the X_i are i.i.d. $\sim p$. Continuing the problem, given $Z_0, Z_1, \ldots, Z_n$ with $Z_n = k$, then $Z_{n+1} \sim X_1 + \cdots + X_k$. It's intuitive to draw this as a tree, where individuals of generation 0 have some number of offspring and some have none. We create a branching tree from one stage to the next. Clearly,(Z_n) is a Markov chain on $\{0, 1, 2, \ldots\}$. Note that state $k = 0$ is absorbing, which fits with the convention of empty sums i.e. summing 0 copies of the offspring variable gives 0.

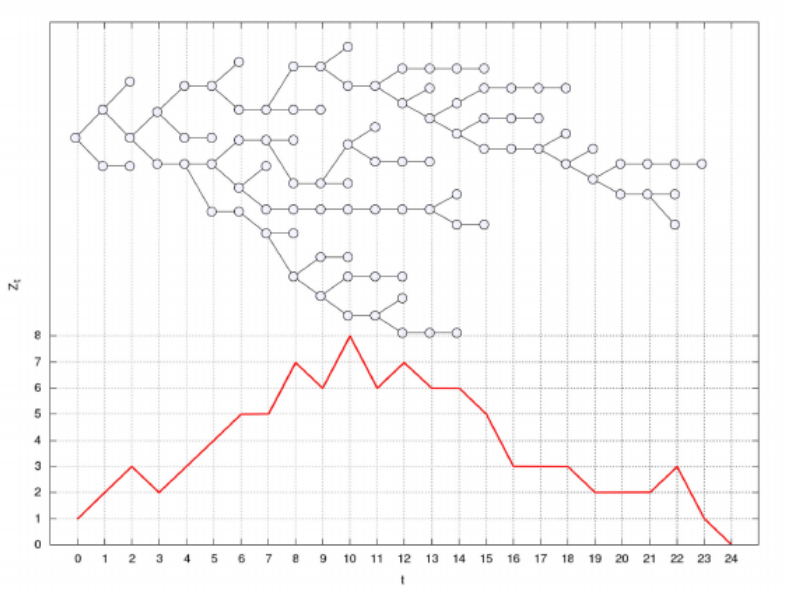

Figure 8.1: See [7] A realization of the Galton-Watson process. At the top, the tree associated to the process is shown, starting from the left ($Z_0 = 1$). At the bottom, the evolution of the number of elements originated in each generation t are displayed.

Here's a visualization of the branching process.

Now, we should expect that generating functions should be helpful, as we are iterating random sums. We'll iterate the composition of generating functions. For simplicity, start with $z_0 = 1$. Let $\phi_n(s) = \mathbb{E}\left(s^{Z_n}\right)$ for $0 \leq s \leq 1$. We see that

$$Z_{n+1} = \text{ sum of } Z_n \text{ copies of } X$$

Hence

$$\phi_1(s) = \sum_{n=0}^{\infty} p_n s^n = \mathbb{E}s^X$$

which we define as the **offspring generating function**. To find ϕ_2, we look at $\phi_1(\phi_1(s))$. That is,

$$\begin{aligned} \phi_2(s) &= \text{ PGF of sum of } Z_1 \text{ copies of } X \\ &= \phi_1\left[\phi_1(s)\right] \end{aligned}$$

Continuing, we similarly have

$$\begin{aligned} \phi_3(s) &= \text{ PGF of sum of } Z_2 \text{ copies of } X \\ &= \phi_1(\phi_1(\phi_1(s))) \end{aligned}$$

and so on. Now Pitman presents the famous problem of finding the probability of

extinction

$$\mathbb{P}_1(\text{extinction}) = \mathbb{P}_1(Z_n = 0 \text{ for large } n)$$

$$= \lim_{n\to\infty} \mathbb{P}_1(Z_n = 0)$$

Now we ask, how do we find $Z_n = 0$? We basically have a formula for this. What is the probability that $Z_1 = 0$? This is simply

$$\mathbb{P}_1(Z_1 = 0) = p_0$$

Then

$$\mathbb{P}_1(Z_2 = 0) = \phi(\phi(0)) = \phi(p_0)$$

and similarly,

$$\mathbb{P}_1(Z_3 = 0) = \phi(\phi(\phi(0))) = \phi(\phi(p_0))$$

and so on. See figure 8.2. This gives the exact formula in general that

$$\mathbb{P}_1(Z_n = 0) = \phi_{n-1}(0)$$

where ϕ_{n-1} is the n th iterate of the offspring generating function ϕ. Because ϕ is continuous, it follows that the extinction probability

$$s_0 := \lim_{n\to\infty} \phi_n(s)$$

is a root of the equation

$$s = \phi(s)$$

In general s_0 is the least root s of this equation with $0 \leq s \leq 1$. Note that $s = 1$ is always a root. Even if you aren't a fan of generating functions, you should note that they are inescapable in the solution to the branching extinction problem.
By analysis of the graph of ϕ, which is convex with $\phi(0) = p_0$ and derivative $\phi'(1-) = \mu$, there are three cases:

- *supercritical* ($\mu > 1$): then there is a unique root s_0 with $0 \leq s_0 < 1$ and $\phi(s_0) = s_0$. This is the extinction probability.
- *subcritical* ($\mu < 1$): then the unique root is $s_0 = 1$: extinction is certain;
- *critical and non-degenerate* ($\mu = 1$) and $p_0 > 0$: then $s_0 = 1$. See figure 8.3. See figure 8.3. Because the generating function ϕ is convex, the only root returned from fixed point iteration is precisely at 1. This implies that if $\mu = 1$ and $p_0 > 0$ the probability of extinction $\mathbb{P}_1(\text{extinction}) = 1$. The fluctuations of Z_n in this case lead with probability one to extinction.

There is a very annoying case for branching processes that we should not forget, which is the *degenerate case* where $p_1 := \mathbb{P}(X = 1) = 1$, which just makes the population stay at 1, $\mathbb{P}_1(Z_n = 1) = 1$ for all n, and the extinction probability is 0. There is no random fluctuation in it. The book presents this conclusion in different ways.

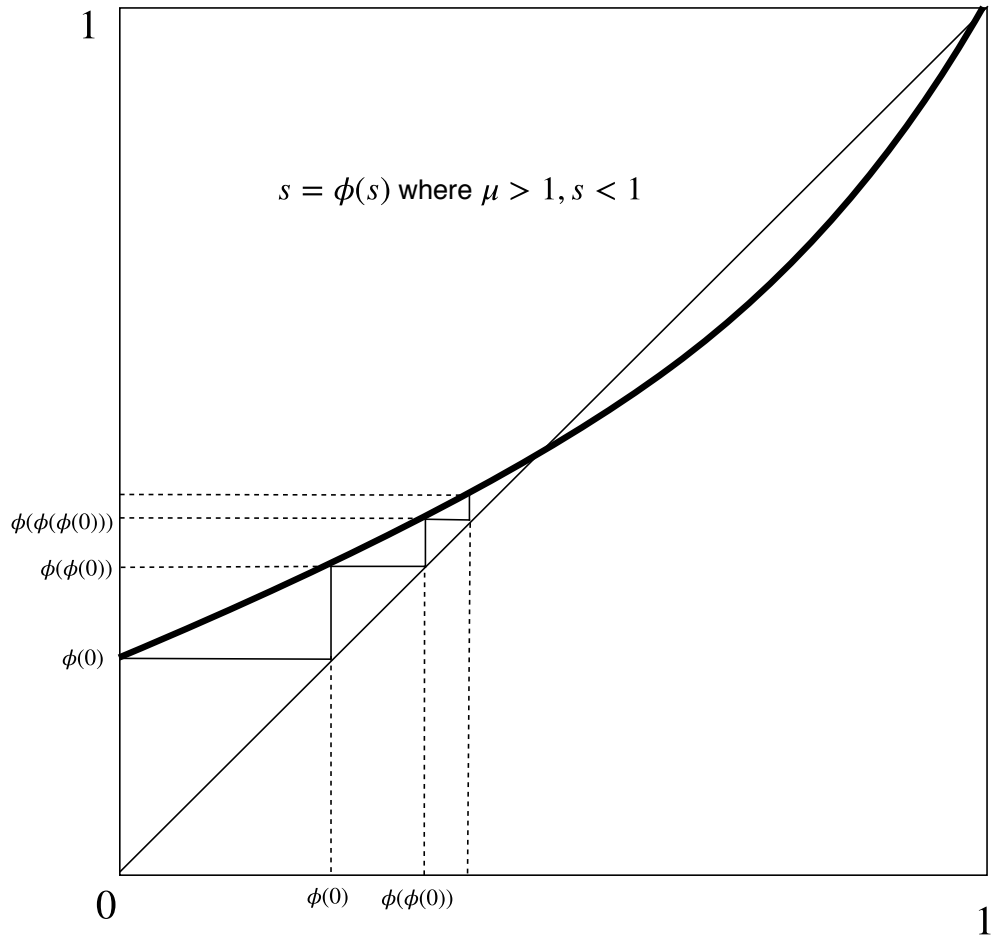

Figure 8.2: Supercritical. As an example, we sketch ($\phi(s)$ with respect to s for the generating function of **Poisson** $(3/2)$. This gives a fixed point iteration returning the unique root s of $s = \phi(s)$ with $s < 1$. Here the mean is larger than 1 ($\phi'(1) = \mu > 1$).

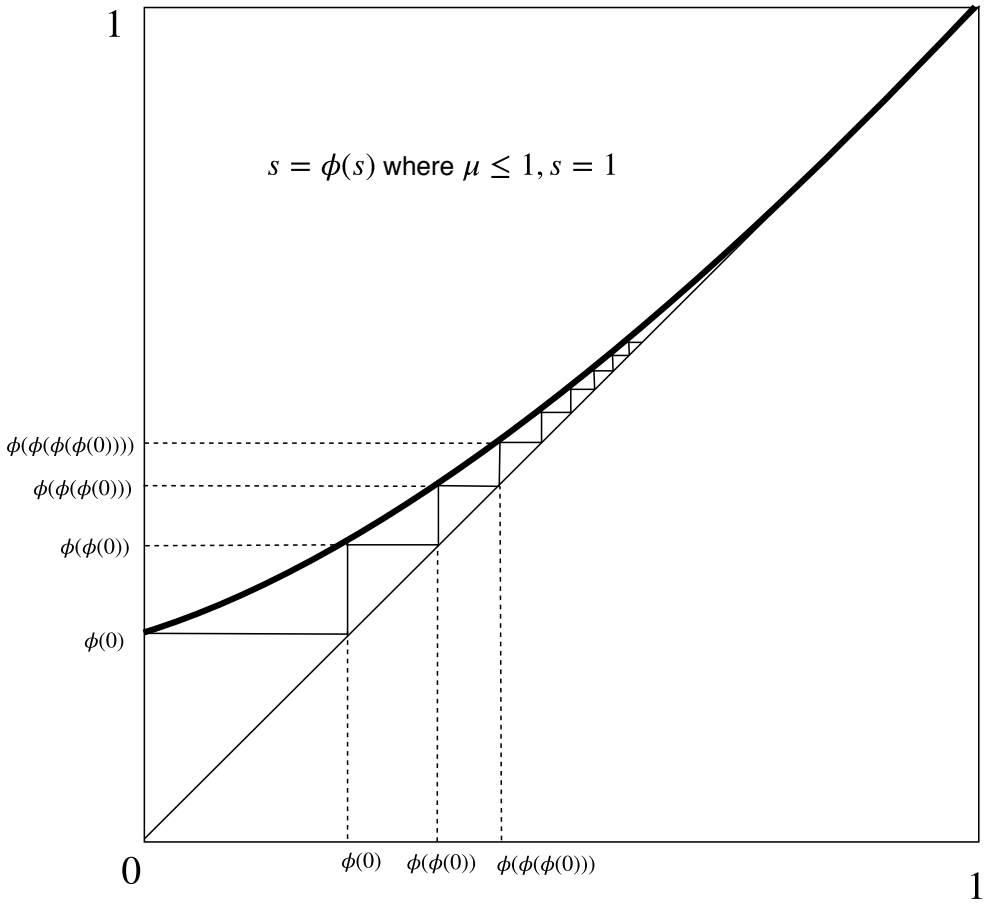

Figure 8.3: Subcritical or critical non-degenerate. We see that if $p_0 > 0$ and $\mu := \sum_n np_n \leq 1$, then the probability generating function is a convex curve with slope ≤ 1 at 1 and value $p_0 > 0$ at 0. So the curve cannot ever cross the diagonal s. There is a no root of $\phi(s) = s$ with $s < 1$, and hence $\phi(\phi(\phi(\cdots(0)))) \to 1$ as $n \to \infty$

LECTURE 9

Potential Theory (Green Matrices)

9.1 Potential Theory (Green Matrices)

Green was an English mathematician who mainly worked on differential equations. This is a concept borrowed from differential equations, applied to our present context. Let P be a transition matrix on a countable state space S. Define

$$\begin{aligned} G(x,y) &:= \sum_{n=0}^{\infty} P^n(x,y) \\ &= \mathbb{E}_x \sum_{n=0}^{\infty} \mathbb{1}(X_n = y) \\ &= \mathbb{E}_x N_y \end{aligned}$$

As a bit of book-keeping, we define

$$N_y := \text{ total \# visits to } y \text{ including a visit at time } n = 0$$

We can tell if states are transient or recurrent by looking at the Green matrix. Recall that we saw before that

$$G(x,x) < \infty \iff x \text{ is transient}$$

$$G(x,x) = \infty \iff x \text{ is recurrent}$$

Remark: If S is finite, then it is obvious that $G(x,x) = \infty$ for some x. Particularly, some state must be recurrent. Let us start with the case where we have *infinite state space* and a *transient chain*. Recall that we defined

$$\begin{aligned} T_y &:= \min\{n \geq 1 : X_1 = y\} \\ V_y &:= \min\{n \geq 0 : X_n = y\} \end{aligned}$$

The following is generally true with no assumptions. There is a relation between $G(x, y)$ and $G(y, y)$. Always, $G(x, y) \leq G(y, y)$, and the ratio is a hitting probability

Key Fact

For all x and y (including $x = y$),

$$G(x, y) = \mathbb{P}_x(V_y < \infty)G(y, y)$$

Why is this formula true? Simply, $\mathbb{E}_x N_y$ is computed by conditioning on the event $(N_y > 0)$ which is identical to the event $(V_y < \infty)$. On this event, once we get to y, the chain starts over as if from y. To be more formal, we would cite the Strong Markov Property. Now let's look at a key example.

9.1.1 Example

Consider a simple random walk on the integers $\mathbb{Z} := \{\dots, -1, 0, 1, \dots\}$, where from any integer, we go left one with probability q and right one with probability p. Pitman notes that we have something similar to this on homework 4, where we are moving on a circle. However, we could very easily unwrap the circle into the integer line. Let's compute the potential kernel. First of all, we ask if we really need the first parameter x, where $x, y \in \mathbb{Z}$. The definition of subtraction (translation invariance of the transition matrix) gives us:

$$G(x, y) = G(0, y - x)$$

So it's enough to discuss $G(0, y)$. Because of the key fact above, most of the action is looking at $G(y, y) = G(0, 0)$, and hence

$$G(x, y) = \underbrace{G(0, 0)}\,\underbrace{\mathbb{P}_x(\text{event of hitting } y)}$$

These two are our key ingredients. Let's first compute $G(0, 0)$. Notice that this random walk can only come back on even n (this is a periodic walk).

$$\begin{aligned} G(0, 0) &:= \sum_{n=0}^{\infty} P^n(0, 0) \\ &= \sum_{m=0}^{\infty} P^{2m}(0, 0) \\ &= \sum_{m=0}^{\infty} \binom{2m}{m} p^m q^m \end{aligned}$$

and comparing this against the case where $p = q = \frac{1}{2}$ and adjusting, we have:

$$G(0, 0) = \sum_{m=0}^{\infty} \binom{2m}{m} 2^{-2m} (4pq)^m$$

Now Pitman states the following fact and requires that we perform this tedious computation once in our life

$$\binom{2m}{m}2^{-2m} = \frac{\left(\frac{1}{2}\right)m\uparrow}{m!} = \frac{\left(\frac{1}{2}\right)\left(\frac{1}{2}+1\right)\cdots\left(\frac{1}{2}+m-1\right)}{m(m-1)\cdots 1}$$

and we know (from the previous lecture) that for $|x|<1$

$$(1+x)^r = \sum_{m=0}^{\infty}\binom{r}{m}x^r$$

$$\implies (1-x)^{-r} = \sum_{m=0}^{\infty}\binom{-r}{m}(-x)^r$$

$$= \boxed{\sum_{m=0}^{\infty}\frac{(r)_{m\uparrow}}{m!}x^m}$$

where we call this the negative binomial expansion. Bringing this back to the problem at hand (recognizing the negative binomial coefficient), we have

$$G(0,0) = \sum_{m=0}^{\infty}\frac{\left(\frac{1}{2}\right)_{m\uparrow}}{m!}(4pq)^m$$

by negative binomial expansion with $r := \frac{1}{2}$ and $x := 4pq$. Then this gives

$$G(0,0) = (1-4pq)^{-\frac{1}{2}}$$

Pitman reminds that in using this expansion, we should always be cautious for convergence in ensuring $|x|<1$. Hence in this problem, provided $4pq<1$ (equivalent to $p\neq\frac{1}{2}$)
Notice that if $p=\frac{1}{2}$, then in our formula we have $(1-1)=0$ to a negative power, which gives us ∞. We can easily check

$$\binom{2m}{m}\left(\frac{1}{2}\right)^{2m} \sim \frac{c}{\sqrt{m}} \qquad (m\to\infty)$$

This precisely gives $G(0,0)=\infty$ in the case $p=\frac{1}{2}$. Hence

$$G(0,0) = \begin{cases}\infty, & p=\frac{1}{2}\\ (1-4pq)^{-\frac{1}{2}}, & p\neq\frac{1}{2}\end{cases}$$

Pitman says we can be a bit cuter about this. Notice

$$1-4pq = 1-4p+4p^2 = (2p-1)^2$$

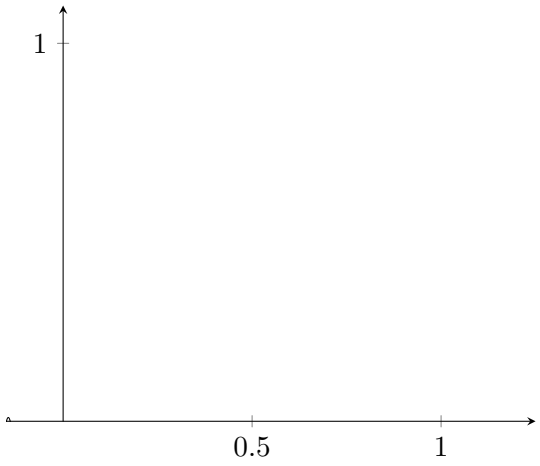

Figure 9.1: Graph of $4p(1-p)$

Therefore,

$$\begin{aligned} G(0,0) &= \left[(2p-1)^2\right]^{-\frac{1}{2}} \\ &= \frac{1}{|2p-1|} \\ &= \frac{1}{2\left|p-\frac{1}{2}\right|} \end{aligned}$$

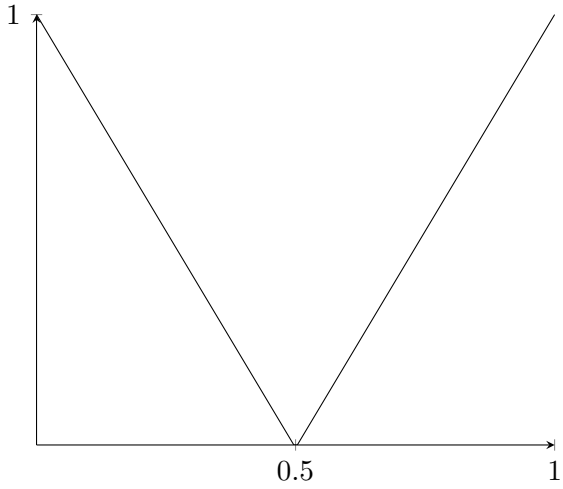

Figure 9.2: Graph of $|2x-1|$. Similar graph present in homework, with the exception of a smoother valley towards $1/2$, as compared to the sharp corner present here.

9.2 Escape Probability

We'll continue the p, q walk. Consider the probability, starting at 0, that we never come back, which we write as $\mathbb{P}_0(T_0 = \infty)$, and set this equal to w.
Now we ask, what is w in terms of $G(0,0)$? Recall that $G(0,0) = \mathbb{E}_0 N_0$, the expected number of hits on 0. Again, our convention is to count the starting position at time 0. Then

$$\mathbb{P}_0(N_0 = 1) = \mathbb{P}_0(T_0 = \infty) = w$$

which is familiar from a past discussion. That is, under P_0, starting at 0, N_0 has a very friendly distribution. *Brief Recap: Geometric Distribution* Recall that for $N \sim \textbf{Geometric}(p)$ on $\{1, 2, 3, \dots\}$ with probability of success is p. Then

$$\mathbb{P}(N = n) = (1-p)^{n-1} p$$

Hence for our problem

$$\mathbb{P}_0(N_0 = n) = (1-w)^{n-1} w$$

so

$$N_0 \sim \textbf{Geometric}(w)$$

Hence we can do away with our placeholder w, so that $w = \frac{1}{G(0,0)}$, so that

$$\mathbb{P}_0(T_0 = \infty) = \frac{1}{G(0,0)} = 2|p - 1/2|$$

The homework problem graph is very similar to this, but with a curve instead of a sharp 'valley' at $1/2$

9.3 More Formulas for Simple Random Walks (SRW)

We consider the same context as our example. Let $S_n = \Delta_1 + \cdots + \Delta_n$, where each Δ_i is either $+1$ or -1 with probability p or q, respectively. What is the probability that S_n tends to $+\infty$? This depends on p. In the recurrent case, this probability will be zero. Now if $p > q$, then we have a "drift" for our expectation that carries us off to ∞. That tells us that

$$\mathbb{P}(S_n \to +\infty) = \begin{cases} 0, & p < q \\ 1, & p > 1 \end{cases}$$

or more neatly, using indicator notation

$$\mathbb{P}(S_n \to +\infty) = \mathbb{1}\left(p > \frac{1}{2}\right)$$

$$\mathbb{P}(S_n \to -\infty) = \mathbb{1}\left(p < \frac{1}{2}\right)$$

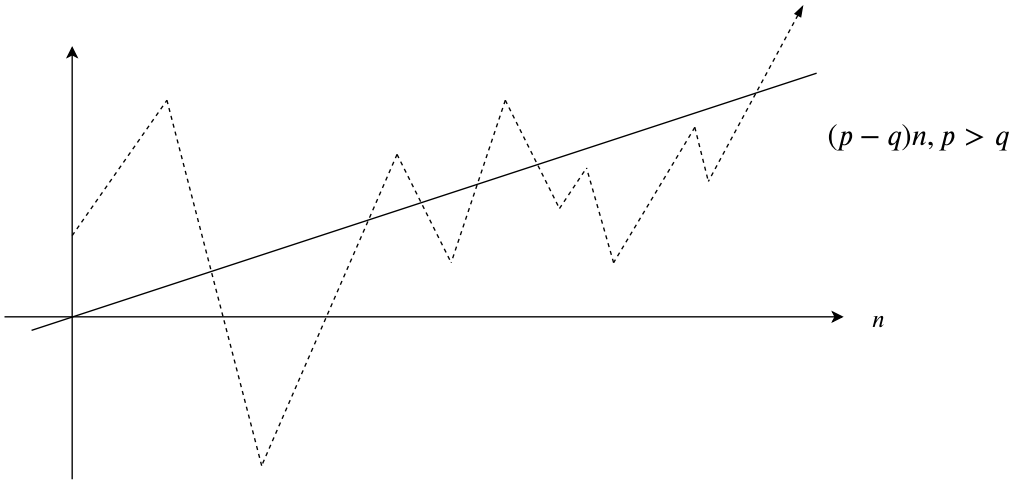

Now assume $p > q$. Then for $x \geq 1$, we have

$$\mathbb{P}_x(T_0 < \infty) = \left(\frac{q}{p}\right)^x$$

by Gambler's ruin probability (in the limit) from a previous lecture. Now $\frac{q}{p} < 1$, so

$$\mathbb{P}_x(T_0 = \infty) = 1 - \left(\frac{q}{p}\right)^x$$

$$\mathbb{P}_{-x}(T_0 < \infty) = 1$$

because the drift up ($\uparrow$) takes us to $+\infty$ with probability 1 and must hit 0 along the way.

9.4 Green's Matrix for Finite State Space S

The entries $G(x, y)$ are only interesting if y is a transient state. Take the example where the set A is an absorbing set of states, and let $S - A$ (or $S \setminus A$) be interior states, and

$$\mathbb{P}_x(T_A < \infty) = 1 \qquad \forall x \in S$$

In the Gambler's Ruin example, take equal probability ($\frac{1}{2}\uparrow, \frac{1}{2}\downarrow$). Let $S := \{0, 1, \ldots, N\}$ and $A = \{0, N\}$. In the matrix (see next page), Q is a $(S-A)\times(S-A)$ matrix. We claim that for $x, y \in S - A$

$$G(x, y) = \sum_{n=0}^{\infty} P^n(x, y) \quad \text{(from earlier)}$$

$$= \sum_{n=0}^{\infty} Q^n(x, y)$$

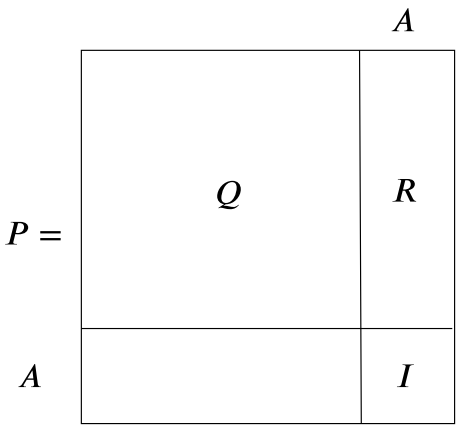

because $P^n(x,y)$ is a sum of products along paths through interior states only, and $P(w,z) = Q(w,z))$ for all transitions (w,z) contributing to such products. Notice that Q is **not** stochastic; in fact, we say that Q is "sub-stochastic". We see

$$(Q\mathbb{1})(x) = \mathbb{P}_x(X_1 \notin A)$$

which will sometimes be less than 1. In any case, it's ≤ 1. We want to focus only on the non-degenerate (interesting) part of G. So, assuming $\mathbb{P}_x(V_a < \infty) > 0$ (there is some positive probability), then $G(x,a) = \infty$ for every a. This follows from

$$G(x,a) = \mathbb{P}_x(V_a < \infty) \underbrace{G(a,a)}_{=\sum_{n=0}^{\infty} 1 = \infty}$$

Remark Now we'll throw away all the absorbing states for our discussion, so that all our matrices are indexed by $S - A$. We're shrinking our matrix to focus on the interesting portion of our potential kernel. Hence

$$G = \sum_{n=0}^{\infty} Q^n, \text{ on } S - A$$

$$= I + Q + Q^2 + Q^2 + \cdots$$

and compare against

$$QG = Q + Q^2 + Q^3 + \cdots$$

and subtracting these gives

$$G - QG = G - GQ = G(I - Q) = I$$

which implies

$$\boxed{G = (I - Q)^{-1}} \qquad (\star\star\star)$$

Computationally, this boils down to simply using our computers to crunch the inverse. Of course, for large matrix powers, we may run into underflow, overflow, or computationally singular matrices. However, there are ways to treat this issue within numerical linear algebra.

9.4.1 Return to Gambler's Ruin

We'd like to find $G(x, \cdot)$, which is the row x of the Green matrix for Gambler's Ruin. Take

$$G - GQ = I \implies G = I + GQ$$

Now the row $G(x, \cdot)$ is determined in general by

$$G(x, y) = \mathbb{1}(x = y) + \sum_z G(x, z)Q(z, y)$$

which for the Gambler's Ruin chain becomes

$$G(x, y) = \mathbb{1}(x = y) + \frac{1}{2}G(x, y - 1) + \frac{1}{2}G(x, y + 1)$$

with the boundary conditions that

- if $y = 1$ then $G(x, 0) = 0$
- if $y = N - 1$ then $G(x, N) = 0$

Now we'd like to graph $y \mapsto G(x, y)$.

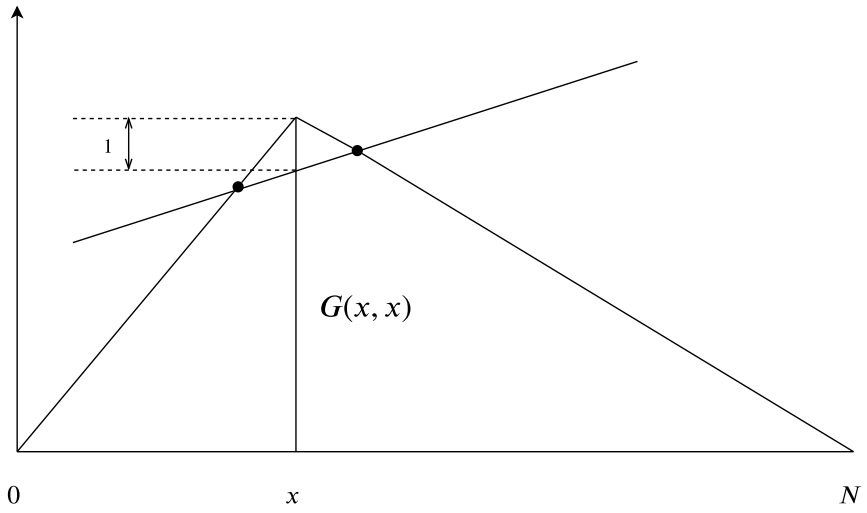

If we ignore the indicator term, the right two terms gives a straight line from where we start at x, to the boundary, where the function must vanish. So the graph is a tent with its peak at x, and the equation for x says this peak value $G(x, x)$ is 1 greater than the average of values to its right and left. That indicates how high the peak is, so the equations determine $G(x, \cdot)$ completely. From this consideration

$$G(x, y) = \begin{cases} \dfrac{y}{x}G(x, x), & 0 \le y \le x \\ \dfrac{N - y}{N - x}G(x, x), & x \le y \le N \end{cases}$$

Also

$$G(x,x) - 1 = \frac{1}{2}\left(\frac{x-1}{x} + \frac{N-(x+1)}{N-x}\right) G(x,x)$$

which is easily solved to give

$$G(x,x) = \frac{2x(N-x)}{N}$$

We have two important checks on this calculation. We know

$$\mathbb{P}_x(\text{ hit 0 before } N) = 1 - \frac{x}{N} = \frac{N-x}{N}$$

But we can compute this by conditioning on T_0: for $0 < x < N$

$$\begin{aligned}
\mathbb{P}_x(\text{ hit 0 before } N) &= \sum_{n=0}^{\infty} \mathbb{P}_x(T_0 = n+1) \\
&= \sum_{n=0}^{\infty} \mathbb{P}_x(X_n = 1, X_{n+1} = 0) \\
&= \sum_{n=0}^{\infty} \mathbb{P}_x(X_n = 1)\,\frac{1}{2} \\
&= G(x,1)\,\frac{1}{2} = \frac{1}{x}\frac{2x(N-x)}{N}\frac{1}{2} = \frac{N-x}{N}
\end{aligned}$$

Similarly, by the same method, get

$$\mathbb{P}_x(\text{hit } N \text{ before } 0) = \frac{x}{N}$$

as before. Also from before, we found $\mathbb{E}_x V_{0N} = x(N-x)$ which we can now check using the Green matrix

$$\mathbb{E}_x V_{0N} = \mathbb{E}_x \sum_{y=1}^{N-1} N_y \tag{9.1}$$

$$= \sum_{y=1}^{N-1} G(x,y) \tag{9.2}$$

$$= \sum_{y=1}^{x-1} G(x,y) + G(x,x) + \sum_{y=x+1}^{N-1} G(x,y) \tag{9.3}$$

$$= \cdots = x(N-x) \tag{9.4}$$

where you can easily fill in the "$\cdots$."

9.4.2 Conclusion

Whenever it is possible to evaluate the Green matrix G you have immediate access to both hitting probabilities and mean hitting times as in the example above. See the text around (1.26) on page 62 for other applications of the same method.

LECTURE 10

The Fundamental Matrix of a Positive Recurrent Chain

10.1 Comments on Homework 5

This week's homework is posted as a worksheet on bCourses, and there is one correction as posted on Piazza. The first two problems are quite easy. The first is on branching processes.

For the second problem, we'll be frustrated if we don't know about Poisson thinning. This is Poisson thinning in disguise within a Markov chain.
In words, this means that if we have a Poisson (μ) number of independent Bernoulli(p) trials, the count of successes is Poisson (λp). We should use the 'obvious' generating function to show this.

Now, a hint for the Kac identities, which are about a stationary process of 0s and 1s. With an informal notation, the claim is that in a stationary sequence of 0s and 1s $\mathbb{P}(1\underbrace{000}_{n}) = \mathbb{P}(\underbrace{000}_{n}1)$. This is much weaker than assuming reversibility (this is only for one such pattern). As a hint:

$$\mathbb{P}(1000) = \mathbb{P}(*000) - \mathbb{P}(0000),$$

where $*$ acts as a wild-card and can be a 0 or 1. Once you have worked with this idea, it takes about 4 lines to solve the problem.

The exercise on tail generating functions is just routine for us to get to get practice with generating functions.

The renewal generating function problem is relatively easy from the result of #4. We'll discuss a bit of renewal theory today in-class.

10.2 Renewal Generating Functions

For our discussions today, we assume P is irreducible.

We look at the summation

$$\sum_{n=0}^{\infty}\left(u_n - \frac{1}{\mu}\right) = \sum_{n=0}^{\infty}\left[P^n(0,0) - \pi(0)\right],$$

where there is an underlying Markov chain, P is our transition matrix with stationary probability π, and we assume irreducible and aperiodic.
We take the state 0 to be a special state. Then in the language of Renewal theory (by definition, returning to our initial state),

$$u_n = P^n(0,0) = \mathbb{P}_0(\text{return to 0 at time } n) = \mathbb{P}(\text{renewal at } n).$$

and the mean inter-renewal time is $\mu = 1/\pi(0)$. This formula arises from looking at a natural generalization of the Potential kernel (Green matrix):

$$G := \sum_{n=0}^{\infty} P^n = (I - P)^{-1},$$

for a Markov matrix P. This matrix G is very useful for transient chains, when all entries of G are finite. We may ask, what is the Green matrix if P is recurrent? In general

$$G(x,y) = \mathbb{E}_x \sum_{n=0}^{\infty} \mathbb{1}(X_n = y) = \sum_{n=0}^{\infty} P^n(x,y).$$

In words, this is $\mathbb{E}_x(\#$ of hits on y with infinite time horizon). We know

$$G(x,y) = \infty, \forall_{x,y\in S} \iff P \text{ is recurrent}$$

Finite state irreducible P are very interesting. They have a stationary distribution π, with $\pi P = \pi$. However, $G(x,y) = \infty$ for all x and y, which is of no interest. However,, there is another matrix associated with a positive recurrent chain which is nearly as informative as $G = (1-P)^{-1}$ for a transient chain.
Assume for simplicity that S is finite, and that P is aperiodic. Then we know a lot about P^n. We know

$$P^n(x,y) \to \pi(y) \text{ as } n \to \infty.$$

Informally, the process loses track of its starting state x, and no matter what the value x of X_0 the distribution of X_n given $X_0 = x$ approaches the limit distribution π as $n \to \infty$.
As an aside, we invent the notation $\mathbb{1}$ for a column vector of all 1s: so $\mathbb{1}(x) = 1$, for all $x \in S$. So $\pi P = \pi$, and $\pi\mathbb{1} = 1$.

A bit 'cuter': we use matrix notation to write simply:

$$P^n \to \Pi := \mathbb{1}\pi \text{ as } n \to \infty$$

where the limit matrix Π is the matrix with all rows equal to π. Notice that $P^n \to \Pi$ rapidly as $n \to \infty$. If we're careful about this,

$$\left|P^n(x,y) - \pi(y)\right| \le c\rho^n \text{ for some } c < \infty \text{ and } 0 < \rho < 1.$$

This implies that

$$\sum_{n=0}^{\infty} \left|P^n(x,y) - \pi(y)\right| < \infty,$$

which then implies that

$$\sum_{n=0}^{\infty} (P^n - \Pi)$$

exists, entrywise as a limit matrix.
Look what happens when we square $(P - \Pi)$. Recall that matrix multiplication is not commutative; however, we can still perform the expansion:

$$\begin{aligned}(P-\Pi)^2 &= (P-\Pi)(P-\Pi)\\ &= P^2 - P\Pi - \Pi P + \Pi^2\\ &= P^2 - \Pi - \Pi + \Pi\\ &= P^2 - \Pi\end{aligned}$$

as you can easily show. In general, the product of any number of factors P and Π is Π, so long as there is at least one Π. Hence,

$$(P-\Pi)^n = P^n - \Pi \text{ for } n = 1, 2, \ldots$$

where we need only use the binomial theorem to evaluate the coefficient of Π. Beware that

$$(P-\Pi)^0 = I \neq I - \Pi = P^0 - \Pi$$

which means that the $n = 0$ term must be treated separately in calculations such as the following:

$$\begin{aligned}\sum_{n=0}^{\infty} (P^n - \Pi) &= I - \Pi + \sum_{n=1}^{\infty} (P^n - \Pi)\\ &= I - \Pi + \sum_{n=1}^{\infty} (P-\Pi)^n\\ &= \sum_{n=0}^{\infty} (P-\Pi)^n - \Pi\end{aligned}$$

Recall that

$$\sum_{n=0}^{\infty} K^n = (I-K)^{-1},$$

for suitable K (like a sub-stochastic matrix). Now this is the case for $K := P - \Pi$, and we can define:

$$\boxed{Z := \sum_{n=0}^{\infty} (P-\Pi)^n = (I - P + \Pi)^{-1}}$$

which is called the *fundamental matrix* of the irreducible, recurrent Markov chain. Our homework is to look at

$$\sum_{n=0}^{\infty} \left(u_n - \frac{1}{\mu} \right),$$

for a renewal sequence u_n. We can always write this, for a suitable Markov matrix P with invariant probability π , $\Pi = \mathbb{1}\pi$ and fundamental matrix Z as above, (see "renewal chain" early in the text), as

$$\sum_{n=0}^{\infty} \left(u_n - \frac{1}{\mu} \right) = \sum_{n=0}^{\infty} (P^n(0,0) - \pi(0)) = Z(0,0) - \pi(0) \tag{10.1}$$

To summarize this discussion:

- for an aperiodic recurrent P with finite state space S and $\pi P = \pi$, let $\Pi := \mathbb{1}\pi$. Then the matrix $I - P + \Pi$ has an inverse Z as above.

This result is true also for aperiodic P, except that the usual partial sums of the series diverge, and so the series must be evaluated using an Abel sum:

$$Z := (I - P + \Pi)^{-1} = \lim_{s \uparrow 1} \sum_{n=0}^{\infty} (1-P)^n s^n$$

as discussed further in the homework for the diagonal entries of Z corresponding to (10.1).

Example: Consider the period 2 transition probability matrix

$$P = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix} \implies \Pi = \begin{bmatrix} 1/2 & 1/2 \\ 1/2 & 1/2 \end{bmatrix} \implies I - P + \Pi = \begin{bmatrix} 3/2 & -1/2 \\ -1/2 & 3/2 \end{bmatrix}$$

which is easily inverted to give

$$Z = (I - P + \Pi)^{-1} = \begin{bmatrix} 3/4 & 1/4 \\ 1/4 & 3/4 \end{bmatrix} \tag{10.2}$$

In this example $P^n(0,0)$ for $n = 0, 1, 2, \dots$ gives the sequence $(1, 0, 1, 0, 1, 0, 1, 0, \dots)$, with $\pi(0) = \frac{1}{2}$ and $\mu = 1/\pi(0) = 2$. Notice that

$$(P^n(0,0) - \pi(0), n = 0, 1, 2, \dots) = \left(\frac{1}{2}, -\frac{1}{2}, \frac{1}{2}, -\frac{1}{2}, \cdots\right),$$

which gives an oscillating sum if left untreated. Using an Abel sum, we can check that

$$\lim_{s\uparrow 1} \left(P^n(0,0) - \pi(0)\right) s^n = \lim_{s\uparrow 1} \frac{1}{2} \sum_{n=0}^{\infty} (-s)^n = \lim_{s\uparrow 1} \frac{1}{2} \frac{1}{(1+s)} = \frac{1}{4} \tag{10.3}$$

Whereas according to (10.1) and (10.2) the same limit is evaluated as

$$Z(0,0) - \pi(0) = \frac{3}{4} - \frac{1}{2} = \frac{1}{4}. \tag{10.4}$$

As detailed in later sections there are lots of formulas for features of a recurrent Markov chain in terms of entries of the fundamental matrix Z. For transient matrices, we express things in terms of entries of the Green matrix. For recurrent chains, we express things in terms of entries of Z.

10.3 Variance of Sums Over a Markov chain

A first indication of the importance of the matrix Z comes from computation of variances in the Central Limit Theorem (CLT) for Markov chains.
A key special case arises when $P = \Pi$. This means that under $\mathbb{P}_\lambda$ with $X_0 \sim \lambda$ all the following variables are iid with $X_1 \sim \pi$, $X_2 \sim \pi$, and so on.
In this (iid) case with $P = \Pi$, and more generally, we look at:

$$S_n(f) := \sum_{k=1}^{n} f(X_k),$$

which we may regard as the reward from n steps of the chain if we are paid $\$f(x)$ for each value x. In the iid case, and more generally under $\mathbb{P} = \mathbb{P}_\pi$ with $X_0 \sim \pi$, so the chain is stationary,

$$\begin{aligned} \mathbb{E}S_n(f) &= n\mathbb{E}f(X_1) \\ &= n\pi f, \end{aligned}$$

where $\pi f = \sum_x \pi(x) f(x)$ is a real number. Notice that the states X_n of the chain can be abstract, but we assume f to take on numerical values, so that we may discuss the expectation and variance of sums like $S_n(f)$. Continuing in the iid case $P = \Pi$, we have

$$\mathbb{V}\mathrm{ar}(S_n(f)) = n\mathbb{V}\mathrm{ar}(f(X_1)),$$

and $\mathbb{V}\mathrm{ar}(X) = \mathbb{E}(X^2) - \mathbb{E}(X)^2$ gives, in matrix notation:

$$\sigma^2(f) := \mathbb{V}\mathrm{ar}(X_1(f)) = \pi f^2 - (\pi f)^2 .$$

Then by the Central Limit Theorem for sums of iid random variables, provided $\pi f^2 < \infty$,

$$\mathbb{P}\left(\frac{S_n(f) - n\pi f}{\sigma(f)\sqrt{n}} \le z\right) \to \Phi(z),$$

the standard normal CDF. So what about for a Markov chain?

10.3.1 The Mean

Assuming now that (X_n) is an irreducible finite state Markov chain with $X_0 \sim \lambda$

$$\begin{aligned}
\mathbb{E}_\lambda \sum_{k=1}^{n} f(X_k) &= \lambda \left(\sum_{k=1}^{n} P^k\right) f \\
&= n\lambda \left(\frac{1}{n}\sum_{k=1}^{n} P^k\right) f \\
&\sim n \underbrace{\lambda \Pi}_{=\pi} f = n\pi f, \text{ as } n \to \infty
\end{aligned}$$

Notice that the mean per step πf is the same as in the iid case. If $\lambda = \pi$ the formula for the mean of $S_n(f)$ is still exactly $n\pi f$, but for $\lambda \neq \pi$ this formula only holds asympotically in the limit as $n \to \infty$ instead of exactly.

10.3.2 The Variance

At least in the aperiodic case, because of the convergence in distribution of X_n to π, it is intuitively obvious that the behavior of $\mathbb{V}\mathrm{ar}S_n(f)$ for large n can't depend much on the initial distribution λ. We certainly see this for the mean, and it holds here for the variance as well, even for periodic chains. So the most important case for variance computations is the stationary case with $\lambda := \pi$. Then we can compute

$$\begin{aligned}
\mathbb{V}\mathrm{ar}_\pi\left(S_n(f)\right) &= \mathbb{V}\mathrm{ar}_\pi\left(f(X_1) + f(X_2) + \cdots + f(X_n)\right) \\
&= \sum_{k=1}^{n} \mathbb{V}\mathrm{ar}_\pi f(X_k) + 2 \sum_{1 \le j < k \le n} \mathrm{Cov}\left[f(X_j), f(X_k)\right] \\
&= n\, \mathbb{V}\mathrm{ar}_\pi f(X_1) + 2\sum_{l=1}^{n-1} (n-l)\mathrm{Cov}\left[f(X_0), f(X_l)\right]
\end{aligned}$$

because stationarity of the process implies $\mathrm{Cov}\left[f(X_j), f(X_k)\right] = \mathrm{Cov}\left[f(X_0), f(X_{k-j})\right]$ and for $1 \le l \le n-1$ there are $n-l$ pairs $1 \le j < k \le n$ with $k - j = l$.

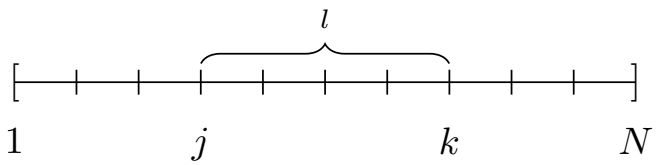

Now we are interested in what happens for large n. We see that

$$\frac{\mathbb{V}\mathrm{ar}_\pi(S_n(f))}{n} = \mathbb{V}\mathrm{ar}_\pi f(X_1) + 2\sum_{l=1}^{n-1}\left(1-\frac{l}{n}\right)\mathrm{Cov}\left[f(X_0), f(X_l)\right]$$

To simplify, assume that $\pi f = 0$ which is the same as $\mathbb{E}_\pi f(X_k) = 0$ because $X_k \sim \pi$ under $\mathbb{P}_\pi$, the probability with the stationary measure. Then

$$\frac{\mathbb{V}\mathrm{ar}_\pi(S_n(f))}{n} = \pi f^2 + 2\sum_{l=1}^{n-1}\left(1-\frac{l}{n}\right)\mathbb{E}_\pi f(X_0)f(X_l),$$

where we get rid of the subtraction terms by our assumption. Now we ask how to compute $\mathbb{E}_\pi\left[f(X_0)f(X_l)\right]$, where we must respect the joint distribution between X_0 and X_l, which involves P^l. Use Markov chain properties and condition on X_0:

$$\begin{aligned}
\mathbb{E}_\pi\left[f(X_0)f(X_l)\right] &= \mathbb{E}_\pi\left[f(X_0)(\overbrace{\mathbb{E}_\pi f(X_l)|X_0}^{=(P^l f)(X_0)})\right] \\
&= \mathbb{E}_\pi f(X_0)(P^l f)(X_0) \\
&= \mathbb{E}_\pi\left(f\cdot P^l f\right)(X_0) \\
&= \pi(f\cdot P^l f),
\end{aligned}$$

where $(f\cdot g)(x) = f(x)g(x)$; in other words, not matrix multiplication of column vector times column vector, which does not make sense.

10.3.3 The Central Limit Theorem

For f with $\pi f = 0$ we have derived the following exact formula

$$\frac{\mathbb{V}\mathrm{ar}_\pi\left[S_n(f)\right]}{n} = \pi f^2 + 2\sum_{l=1}^{n-1}\left(1-\frac{l}{n}\right)\pi\left(f\cdot P^l f\right).$$

Now we would like to see what happens when n is large (as in the Central Limit Theorem). Let $n\to\infty$. Assume P is aperiodic, so that $P^l \to \Pi$, and use $1-\frac{l}{n}\uparrow 1$. Then

$$\frac{\mathbb{V}\mathrm{ar}_\pi\left[S_n(f)\right]}{n} \xrightarrow{n\to\infty} \sigma^2(f) \tag{10.5}$$

where the asymptotic variance $\sigma^2(f)$ per step of the chain can be evaluated from the limit of the exact formula above as

$$\begin{aligned}
\sigma^2(f) &= \pi f^2 + 2\sum_{l=1}^{\infty} \pi\left(f \cdot P^l f\right) \\
&= \pi f^2 + 2\pi f\left(\sum_{l=1}^{\infty} P^l f\right) \\
&= \pi f^2 + 2\pi f \sum_{l=1}^{\infty}\left(P^l - \Pi\right) f \text{ because } \pi f = 0 \\
&= \pi f^2 + 2\pi f\,(Z - I)\, f \text{ where } Z := \sum_{l=0}^{\infty}(P - \Pi)^l \\
&= \boxed{2\pi f \cdot Zf - \pi f^2}.
\end{aligned}$$

To summarize, the fundamental matrix Z arises naturally in the formula for the limiting variance per unit time in the sum $S_n(f)$ of values of a function f of a stationary Markov chain.

Central Limit Theorem for Markov Chains

The CLT works for any irreducible and positive recurrent Markov chain and any f with $\pi f^2 < \infty$ with this evaluation of the asymptotic mean and variance.

See Asmussen's text APQ [4] for a more careful statement and proof.

10.4 Further Applications of the Fundamental Matrix

The following material was not covered in lecture, but serves to review many of the basic properties of Markov chains, and provides further applications of the fundamental matrix of a recurrent Markov chain.

10.4.1 Stopping times

First recall from the text, page 13, the definition of a stopping time T for a discrete time stochastic process $(W_n) = (W_n, n = 0, 1, 2\ldots)$ with countable state space. That is, T is a random variable with values in $\{0, 1, 2, \ldots, \infty\}$, such that for each $n = 0, 1, 2, \ldots$ the event $(T = n)$ is determined by the values of $W_0, \ldots, W_n$. Formally, the indicator function $\mathbb{1}(T = n)$ is a function of $(W_0, \ldots, W_n)$:

$$\mathbb{1}(T = n) = f_n(W_0, \ldots, W_n) \qquad (n = 0, 1, 2, \ldots) \tag{10.6}$$

for some $\{0, 1\}$-valued function f_n of $n+1$ variables $W_k \in S$ for $0 \le k \le n$. Because

$$1 - \mathbb{1}(T > n) = \mathbb{1}(T \le n) = \sum_{k=0}^{n} \mathbb{1}(T = k) \qquad (n = 0, 1, 2, \ldots)$$

equivalent condition are

$$\mathbb{1}(T \le n) = g_n(W_0, \ldots, W_n) \qquad (n = 0, 1, 2, \ldots) \tag{10.7}$$

for some $\{0, 1\}$-valued function g_n of the same variables, and

$$\mathbb{1}(T > n) = h_n(W_0, \ldots, W_n) \qquad (n = 0, 1, 2, \ldots) \tag{10.8}$$

for some $\{0, 1\}$-valued function h_n of the same variables. Intuitively, for each n you can tell whether or not any of the events $(T = n)$, $(T > n)$ or $(T \le n)$ has occurred just by looking at the variables $W_0, \ldots, W_n$. It is only ever necessary to check one of the above conditions (10.6) (10.7) (10.8) for each $n = 0, 1, 2, \ldots$, because each of these conditions implies both the others by simple manipulation of indicator variables. Examples of stopping times T are the now familiar first hitting times $T = V_A$, allowing a hit at time 0, with

$$\mathbb{1}(V_A > n) = \prod_{k=0}^{n} \mathbb{1}(W_k \notin A) \qquad (n \ge 0) \tag{10.9}$$

$$\mathbb{1}(V_A = n) = \mathbb{1}(W_n \in A) \prod_{k=1}^{n} \mathbb{1}(W_k \notin A) \qquad (n \ge 0). \tag{10.10}$$

with the convention for $n = 0$ in the product $\prod_{k=1}^{n}$ that the empty product $\prod_{k=1}^{0} = 1$ (usual convention that an empty product equals 1, which corresponds by taking logs to the convention that an empty sum equals 0). Similarly, the first passage times $T_A := \min\{n \ge 1 : W_n \in A\}$ are stopping times, with

$$\mathbb{1}(T_A > 0) = 1 \tag{10.11}$$

$$\mathbb{1}(T_A > n) = \prod_{k=1}^{n} \mathbb{1}(W_k \notin A) \qquad (n \ge 1) \tag{10.12}$$

$$\mathbb{1}(T_A = 0) = 0 \tag{10.13}$$

$$\mathbb{1}(T_A = n) = \mathbb{1}(W_n \in A) \prod_{k=1}^{n-1} \mathbb{1}(W_k \notin A) \qquad (n \ge 1). \tag{10.14}$$

Note that

- for both $T = V_A$ and $T = T_A$ the logical description of $\mathbb{1}(T > n)$ as a product of indicators, corresponding to an intersection of events, is slightly simpler that the corresponding description of $\mathbb{1}(T = n) = \mathbb{1}(T > n - 1) - \mathbb{1}(T > n)$.
- the definition of a stopping time T relative to (W_n) does not involve any assumptions about the distribution of the process (W_n).

Now let (W_n) be some process defined on a probability space with underlying probability measure $\mathbb{P}$, and let (X_n) be another process derived as $X_n = x_n(W_0, \ldots, W_n)$ for some function x_n of $W_0, \ldots, W_n$. Common examples are

- sums and products, $X_n = W_0 + \cdots + W_n$ and $X_n = W_0 \cdots W_n$.
- expansions of the state space to allow extra randomization: $W_n = (X_n, Y_n)$ for some process (Y_n).

Let P be a transition probability matrix on the countable state space of X. Say that $(X_0, X_1, \ldots)$ is *Markov with transition matrix P relative to the history of* (W_n), abbreviated *Markov* (P) *relative to* (W_n), if

- $X_n = x_n(W_0, \ldots, W_n)$ for some function x_n of $W_0, \ldots, W_n$
- for every $n = 0, 1, 2, \ldots$, every choice of states x and y of the process X, and every event A_n with $\mathbb{1}_{A_n} = f_n(W_0, \ldots, W_n)$ for some function f_n of $W_0, \ldots, W_n$,

$$\mathbb{P}(A_n \text{ and } X_n = x \text{ and } X_{n+1} = y) = \mathbb{P}(A_n \text{ and } X_n = x)P(x, y) \qquad (10.15)$$

which is equivalent by $\mathbb{P}(B \,|\, A) = \mathbb{P}(AB)/\mathbb{P}(A)$ to

$$\mathbb{P}(X_{n+1} = y \,|\, X_n = x \text{ and } A_n) = P(x, y) \qquad (10.16)$$

for all choices of x and A_n with $\mathbb{P}(X_n = x \text{ and } A_n) > 0$.

Typically, (10.15) is used for computations. Taking $\mathbb{1}_{A_n}$ to be a function of $X_0, \ldots, X_n$ shows that if (X_n) is Markov (P) relative to the history of (W_n), then (X_n) is Markov (P) relative to its own history. This is just the usual time-homogeneous Markov property of (X_n). If (X_n) is Markov relative to a richer history (W_n) than just its own history, it means that given $X_n = x$ any additional information in $(W_0, \ldots, W_n)$, beyond what is already encoded in $(X_0, \ldots, X_n)$, is of no use in predicting the next value X_{n+1}: the distribution of this variable given $(X_n = x)$ is $P(x, \cdot)$, no matter what is known about $(W_0, \ldots, W_n)$ besides that the event $(X_n = x)$ has occurred.

Suppose now that T is a stopping time relative to the history (W_n) and that the process X is Markov (P) relative to (W_n), under a probability measure $\mathbb{P} = \mathbb{P}_\lambda$ which assigns X_0 some arbitrary initial distribution λ. Then (10.15) for $A_n = (T = n)$ reads

$$\mathbb{P}_\lambda(T = n \text{ and } X_n = x \text{ and } X_{n+1} = y) = \mathbb{P}_\lambda(T = n \text{ and } X_n = x)P(x, y) \qquad (10.17)$$

so summing over $n = 0, 1, 2, \ldots$ gives

$$\mathbb{P}_\lambda(T < \infty \text{ and } X_T = x \text{ and } X_{T+1} = y) = \mathbb{P}_\lambda(T < \infty \text{ and } X_T = x)P(x, y) \qquad (10.18)$$

That is, for any initial distribution λ of X_0

- under $\mathbb{P}_\lambda$ given $T < \infty$ and $X_T = x$ the distribution of X_{T+1} is $P(x, \cdot)$.

This iterates easily to give the *Strong Markov Property*:

- under $\mathbb{P}_\lambda$ given $T < \infty$ and $X_T = x$ the process $(X_T, X_{T+1}, \cdots)$ has the same distribution as the original chain $(X_0, X_1, \ldots)$ under $\mathbb{P}_x(\cdot) := \mathbb{P}_\lambda(\cdot \mid X_0 = x)$.

This formulation of the strong Markov property is exactly as in Durrett's Theorem 1.2 on page 13, except that Durrett derives the result only for stopping times T of (X_n) itself. The above argument shows that the Strong Markov Property holds also for stopping times T of any process (W_n) such that (X_n) is Markov (P) relative to the history of (W_n). Such stopping times can involve additional randomization beyond the chain (X_n), and are sometimes called *randomized stopping times of* (X_n). Examples of such stopping times involving extra randomization are

- T that is independent of (X_n), taking $W_n = (T, X_n)$;
- $T := \min\{n \geq 0 : X_n = Y\}$ for a state Y chosen independently of (X_n), taking $W_n = (X_n, Y)$;
- $T := \min\{n \geq 0 : p(X_n) \leq U_n\}$ for some function $p : S \to [0,1]$ and (U_n) i.i.d. uniform $[0,1]$ variables independent of (X_n). Here $W_n := (X_n, U_n)$. So given you have not stopped before time n, after observing values of both X_k and U_k for $0 \leq k < n$, and $X_n = x$ you stop at time n with probability $p(x)$, according to whether or not the current uniform variable $U_n \leq x$.

So the Strong Markov Property holds in examples such as these involving extra randomization.

10.4.2 Occupation Measures for Markov chains

Similarly to (10.19), the general formula (10.15) for $A_n = (T > n)$ reads

$$\mathbb{P}_\lambda(T > n \text{ and } X_n = y \text{ and } X_{n+1} = z) = \mathbb{P}_\lambda(T > n \text{ and } X_n = y)P(y, z). \quad (10.19)$$

Summing this over $n = 0, 1, 2, \ldots$ and all states x gives

$$\mathbb{E}_\lambda \sum_{n=0}^{\infty} \mathbb{1}(T > n, X_{n+1} = z) = \sum_y \mathbb{E}_\lambda \left(\sum_{n=0}^{\infty} \mathbb{1}(T > n, X_n = y) \right) P(y, z). \quad (10.20)$$

For any initial probability distribution λ, and any stopping time T of some history (W_n) relative to which (X_n) is Markov with transition matrix P, define measures λG_T and λP_T on the state space of the chain as follows: for $y \in S$

$$\lambda G_T(y) := \mathbb{E}_\lambda \sum_{n=0}^{\infty} \mathbb{1}(T > n, X_n = y) = \mathbb{E}_\lambda \sum_{n=0}^{T-1} 1(X_n = y) \quad (10.21)$$

$$\lambda P_T(y) := \mathbb{E}_\lambda \sum_{n=0}^{\infty} \mathbb{1}(T = n, X_n = y) = \mathbb{P}_\lambda(T < \infty, X_n = y). \quad (10.22)$$

Observe that

- the *pre-T occupation measure* $\lambda G_T(\cdot)$ describes the expected numbers of hits of various states y counting only times n with $0 \le n < T$.

- $\lambda P_T(\cdot)$ is the distribution of X_T on the event $T < \infty$; so $\lambda P_T(\cdot)$ is a sub-probability measure with total mass $\lambda P_T \mathbb{1} = \mathbb{P}_\lambda(T < \infty) \in [0, 1]$.

For purposes of matrix operations, each of these measures on S should be treated as a *row vector*.

- for any non-negative function f with $\lambda G_T |f| < \infty$

$$\lambda G_T f = \sum_{y \in S} \lambda G_T(y) f(y) = \mathbb{E}_\lambda \sum_{n=0}^{T-1} f(X_n) \in [0, \infty] \tag{10.23}$$

- In particular, for $f = \mathbb{1}$, the function with constant value 1, the total mass of $\lambda G_T(\cdot)$ is

$$\lambda G_T \mathbb{1} = \sum_{y \in S} \lambda G_T(y) = \mathbb{E}_\lambda T \in [0, \infty] \tag{10.24}$$

- If this expectation $\mathbb{E}_\lambda T < \infty$, then (10.23) holds with a finite expectation $\lambda G_T f$ for every bounded f.

Let δ_x be the row vector $\delta_x(y) = \mathbb{1}(x = y)$ with mass 1 at x and mass 0 elsewhere. Let $G_T(x, \cdot) := \delta_x G_T(\cdot)$ be the pre-T occupation measure and $P_T(x, \cdot) := \delta_x P_T(x, \cdot)$ the distribution of X_T on $(T < \infty)$ for a chain started in state x. Let G_T and P_T denote the $S \times S$ matrices with these rows. As the notation suggests, and is justified by conditioning on X_0,

$$\lambda G_T(\cdot) = \sum_x \lambda(x) G_T(x, \cdot) \text{ and } \lambda P_T(\cdot) = \sum_x \lambda(x) P_T(x, \cdot). \tag{10.25}$$

Then (10.20) gives the following general *occupation measure identity* for any stopping time T of a Markov chain (X_n), including randomized stopping times, as discussed above:

$$\lambda G_{T+1} = \lambda G_T + \lambda P_T = \lambda + \lambda G_T P. \tag{10.26}$$

This identity is just two different ways of evaluating the pre-$(T + 1)$ occupation measure λG_{T+1}:

- firstly by peeling off the contribution of the last term in (10.21)for $n = T$ on the event $(T < \infty)$, and

- secondly by peeling off the first term for $n = 0$, and evaluating the remaining terms by (10.20).

For every stopping time T of a Markov chain with transition probability matrix P, this identity relates the initial distribution λ of X_0 and the pre-T occupation

measure λG_T to the distribution λP_T of X_T. The occupation measure identity can also be written more compactly as an identity of matrices:

$$G_{T+1} = G_T + P_T = I + G_T P \tag{10.27}$$

where for each $x \in S$ the identity of row x in (10.27) is the identity (10.26) for $\lambda = \delta_x$. As a simple example, if $T = N$ has constant value $N \geq 0$, then (10.27) reduces to

$$\sum_{n=0}^{N} P^n = \left(\sum_{n=0}^{N-1} P^n\right) + P^N = I + \left(\sum_{n=0}^{N-1} P^n\right) P \tag{10.28}$$

which is obviously true for any matrix P with well defined powers, not just for transition matrices.

Call a measure μ on S *locally finite* if $\mu(z)$ is finite for every $z \in S$. Provided the initial distribution λ and the stopping time T are such that the occupation measure λG_T is locally finite, which is typically obvious by geometric bounds, the occupation measure identity (10.26) can be rearranged as

$$\lambda G_T(I - P) = \lambda - \lambda P_T. \tag{10.29}$$

You easily can check that this rearrangement of the occupation measure identity is justified in each of the following cases. These cases include several instances of (10.29) which have been discussed in previous lectures and in the text.

(a) The state space S is infinite, every state $z \in S$ is transient, without any restriction on the stopping time $T \in [0, \infty]$. In particular $T = \infty$ is allowed for such a chain. Then $P_T = P_\infty$ is the zero matrix, and $G_T = G_\infty$ with

$$G_\infty = \sum_{n=0}^{\infty} P^n = (I - P)^{-1}$$

the Green matrix associated with P, as discussed in Lecture 9.

(b) $T = V_A$, the first hitting time of A, counting a hit at time $n = 0$, for any subset A of S such that $\mathbb{P}_x(V_A < \infty) > 0$ for all $x \notin A$. Let $Q_A(x, y) := P(x, y)\mathbb{1}(x \notin A, y \notin A)$ be the restriction of P to $(S - A) \times (S - A)$, and write simply I for the similarly restricted identity matrix. Then, as discussed in Lecture 9,

$$G_{V_A}(x, y) = \left(\sum_{n=0}^{\infty} Q_A^n\right)(x, y) = (I - Q_A)^{-1}(x, y) \text{ for } x \notin A \text{ and } y \notin A \tag{10.30}$$

and $G_{V_A}(x, y) = 0$ else. The evaluation of rows of the pre-V_A occupation matrix G_{V_A} was detailed in Lecture 9 for the fair Gambler's Ruin chain on $S = \{0, 1 \ldots, N\}$ with $A = \{0, N\}$, using a the occupation measure identity (10.29) to argue that in this case the graph of the function $y \mapsto G_{V_A}(x, y)$ is linear on each of the intervals $[0, x]$ and $[x, N]$, vanishing at 0 and N, with peak value $G_{V_A}(x, x)$ which is 1 greater than the average of neighboring values $\frac{1}{2}G_{V_A}(x, x - 1) + \frac{1}{2}G_{V_A}(x, x + 1)$.

(c) λ and T are such that λG_T is locally finite, with $\mathbb{P}_\lambda(T < \infty) = 1$ and $X_T \stackrel{d}{=} X_0$. Then (10.29) becomes $\lambda G_T(I - P) = 0$, so λG_T is a P-invariant measure.

(d) In particular, if P is irreducible and recurrent, $\lambda = \delta_x$ and $T = T_x$ for any fixed state x, then and $X_0 \stackrel{d}{=} X_{T_x}$, so $\mu_x(\cdot) := G_{T_x}(x, \cdot)$, the occupation measure of a single x-block of the chain, gives a strictly positive, locally finite P-invariant measure:

$$\mu_x(\cdot) = \mu_x(\cdot)P \text{ with } 0 < \mu_x(y) < \infty \qquad (y \in S). \tag{10.31}$$

This is Durrett Theorem 1.24 on page 48. The above derivation of this result follows essentially the same steps as Durrett's proof.

10.4.3 Positive recurrent chains: the ergodic theorem

Focusing now on the case when P is irreducible and positive recurrent, formula (10.31) gives a different invariant measure $\mu_x(\cdot)$ for each $x \in S$. Normalizing $\mu_x(\cdot)$ by its total mass

$$m_{xx} := \mu_x(\cdot)\mathbb{1} = \mathbb{E}_x T_x$$

gives a stationary probability measure π for P. It appears at first as if this stationary probability measure $\pi_x(\cdot) = \mu_x(\cdot)/m_{xx}$ for P might depend on the choice of reference state x. But it does not, for it can be shown in many ways that there can be at most one invariant probabiity measure for an irreducible transition matrix P. Hence the basic formula

$$\pi(x) = \frac{1}{m_{xx}} \qquad (x \in S) \tag{10.32}$$

for the unique stationary probability measure π for an irreducible and positive recurrent chain. In particular, this uniqueness of π and (10.32) are consequences of the following *ergodic theorem* for an irreducible chain (Durrett's Theorem 1.22): If π is any stationary probability measure for an irreducible P, then for every initial distribution λ of X_0, and every function f with $\pi|f| < \infty$, there is the convergence

$$\frac{1}{n}\sum_{k=1}^{n} f(X_k) \to \pi f := \sum_y \pi(y) f(y) \text{ as } n \to \infty. \tag{10.33}$$

In particular, for $f(y)$ the indicator function $f_x(y) = \mathbb{1}(y = x)$, this states that the long run limiting relative frequency of times the chain hits state x is $\pi f_x = \pi(x)$. Technically, the convergence (10.33) holds *almost surely*, meaning with $\mathbb{P}_\lambda$ probability one, no matter what the intial distribution λ. A complete formulation and proof of this result is beyond the scope of this course. But this notion of *almost sure convergence* implies what is called *convergence in probability*: for every initial distribution λ

$$\mathbb{P}_\lambda\left(\left|\frac{1}{n}\sum_{k=1}^{n} f(X_k) - \pi f\right| > \epsilon\right) \to 0 \qquad (\forall \epsilon > 0). \tag{10.34}$$

If S is finite, and more generally if $\pi f^2 < \infty$, this can be proved, as in the case of i.i.d. X_i, when $P = \Pi := \mathbb{1}\pi$ is the transition matrix with all rows equal to π, by

bounding the variance, as discussed earlier, and using Chebychev's inequality, From either kind of convergence of averages of $f(X_k)$ over the path of the Markov chain, granted that the limiting average value is a constant πf, this constant is obviously unique. Hence the uniqueness of π. To summarize: for an irreducible transition matrix P with countable state space S, the following conditions are equivalent:

- $m_{xx} < \infty$ for some (hence all) $x \in S$;
- P has an invariant measure μ with $\mu(x) \geq 0$ for all $x \in S$ and $0 < \mu\mathbb{1} < \infty$;
- P has unique invariant probability measure $\pi(x) = 1/m_{xx}$.

Then the transition matrix P, or the associated Markov chain, is called *positive recurrent*, (10.33) and (10.34) imply

$$\frac{1}{n}\sum_{k=1}^{n} P^k \to \Pi := \mathbb{1}\pi \text{ as } n \to \infty. \tag{10.35}$$

If P is aperiodic, then (10.35) can be strengthened (Durrett Theorem 1.23 on page 52) to

$$P^n \to \Pi \text{ as } n \to \infty \tag{10.36}$$

In both (10.35) and (10.36), and other formulas involving limits of matrices in this course, the meaning of convergence is that each entry of the matrix on the left converges to the corresponding entry of the matrix in the right. For transition matrices on both sides, as in (10.35) and (10.36), the rows on both sides are probability measures, meaning the entries are non-negative with row sums 1. Convergence of entries of a sequence of transition matrices then implies convergence of each row of probability measures in the metric of *total variation distance* between probability measures. For instance, in (10.36), for each initial state x,

$$\sum_y |P^n(x,y) - \pi(y)| \to 0 \text{ as } n \to \infty \tag{10.37}$$

as in the last estimate in Durrett's proof of (10.36) on page 53. This implies things like

$$\mathbb{E}_x f(X_n) = (P^n f)(x) = \sum_y P^n(x,y) f(y) \to \pi f \text{ as } n \to \infty \tag{10.38}$$

for every bounded function f, and that this convergence holds uniformly over all f with $|f(x)| \leq b$ for any finite bound b.

10.4.4 Occupation measures for recurrent chains

Suppose now that the transition matrix P is irreducible and recurrent. For any non-empty set of states $A \subseteq S$ in the state space S let $T_A := \min\{n \geq 1 : X_n \in A\}$ be the first passage time to A. Recurrence of the chain implies $\mathbb{P}_x(T_A < \infty) = 1$ for all $x \in S$. Let $G_A(x, \cdot)$ denote the pre-T_A occupation measure starting from state x. A fairly complete analysis of the process $(X_n, 0 \leq n \leq T_A)$, with initial fixed value $x \notin A$ and final random value $X_{T_A} \in A$ is obtained by consideration of the pre-T_A occupation matrix G_A with rows $G_A(x, \cdot)$. Observe that

- if $x \notin A$ then $G_A(x,y) = 0$ for $y \in A$, because the pre-T_A occupation measure only counts hits on A strictly before time T_A.
- these rows $G_A(x,\cdot)$ for $x \notin A$ are therefore completely determined by the restriction of the matrix G_A to $(S-A)\times(S-A)$, as displayed in (10.30), which is just the inverse of the restriction $I-Q_A$ of $I-P$ to $(S-A)\times(S-A)$.
- if $a \in A$ then $G_A(a,y) = \mathbb{1}(a=y)$ for $y \in A$, since the only possible hit of y strictly before T_A is a hit at time 0, and conditioning on X_1 gives

$$G_A(a,y) = \mathbb{1}(a=y) + \mathbb{1}(y \notin A)\sum_{x \notin A} P(a,x)G_A(x,y) \qquad (a \in A) \tag{10.39}$$

  So the rows $G_A(a,\cdot)$ for $a \in A$ are easily determined from the rows $G_A(x,\cdot)$ for $x \notin A$.

Consequently, to compute the entire pre-T_A occupation matrix G_A for any non-empty subset A of states of a recurrent Markov chain, the main task is to compute the restriction of G_A to $(S-A)\times(S-A)$ by inverting the restriction of $I-P$ to $(S-A)\times(S-A)$ matrix $(I-Q_A)$.
Observe that if

$$V_A := \min\{n \ge 0 : X_n \in A\} = T_A\mathbb{1}(X_0 \notin A)$$

then

$$\mathbb{P}_x(T_A = V_A) = 1 \qquad (x \notin A).$$

So for initial states $x \notin A$ the the pre-T occupation measures $G_T(x,\cdot)$ are identical for $T = T_A$ and $T = V_A$:

$$G_A(x,\cdot) := G_{T_A}(x,\cdot) = G_{V_A}(x,\cdot) \qquad (x \notin A).$$

However, for $x = a \in A$, instead of (10.39) there is the trivial evaluation $G_{V_A}(a,\cdot) = 0$, which is sometimes convenient as a boundary condition. For instance, this condition fits well with the equations satisfied by the function

$$m_A(x) := \mathbb{E}_x V_A = \begin{cases} 0, & x \in A \\ \mathbb{E}_x T_A = \sum_{y\in S} G_A(x,y) = G_A(x,\cdot)\mathbb{1}, & x \notin A. \end{cases}$$

If P is positive recurrent, then $m_A(x) < \infty$ for all x. By conditioning on X_1, the function $m(x) = m_A(x)$, regarded as a column vector indexed by states $x \in S$, solves the system of equations

$$m(x) = 1 + \sum_{y\in S} P(x,y)m(y) \text{ with } m(a) = 0 \text{ for } a \in A. \tag{10.40}$$

According to Theorem 1.29 on page 62 of the text, if $S-A$ is finite, this system of equations has unique solution $m(x) = m_A(x)$. With the present assumption that P is irreducible and positive recurrent, this uniqueness can also be shown if $S-A$ is

infinite. Once this mean function $m_A(x)$ is found for starting states $x \notin A$, for a starting state $a \in A$ the mean first return time to A is found, either by summing (10.39) over $y \in S$, or by conditioning on X_1:

$$\mathbb{E}_a T_A = 1 + \sum_{x \notin A} P(a,x) m_A(x) \qquad (a \in A). \tag{10.41}$$

Starting from any state $X_0 = x$, as well as the mean first passage times $\mathbb{E}_x T_A$, the distributions and moments of many other functions of the path $(X_n, 0 \le n \le T_A)$ are easily expressed in terms of the pre-T_A occupation matrix G_A, using consequences of the Strong Markov Property. Consider for instance, the random count of hits on y before T_A

$$N_{yA} := \sum_{n=0}^{T_A - 1} \mathbb{1}(X_n = y)$$

whose mean given $X_0 = x$ is $\mathbb{E}_x N_{yA} = G_A(x,y)$. If $y \in A$ this count is just the constant random variable $\mathbb{1}(x = y)$, either 1 or 0 depending on the starting state x. For $y \notin A$ let

$$p_{xyA} := \mathbb{P}_x(N_{yA} > 0) = \frac{G_A(x,y)}{G_A(y,y)}$$

be the probability starting from x of reaching y before T_A, where the second equality is due to the Strong Markov Property. Then N_{yA} has the modified geometric distribution

$$\mathbb{P}_x(N_{yA} = 0) = 1 - p_{xyA} \tag{10.42}$$

$$\mathbb{P}_x(N_{yA} = n) = p_{xyA}(1 - e_{yA})^{n-1} e_{yA} \qquad (n = 1, 2, \ldots). \tag{10.43}$$

where the escape probability

$$e_{yA} = \mathbb{P}_y(N_{yA} = 1) = \mathbb{P}_y(T_A < T_y) = \frac{1}{G_A(y,y)}$$

is determined by the mean $G_A(y,y)$ of the $\mathbb{P}_y$ distribution of N_{yA}, which is geometric (e_{yA}) on $\{1, 2, \ldots\}$. Thus for $y \notin A$, due to the Strong Markov Property

- for each starting state x with $p_{xyA} > 0$ the $\mathbb{P}_x$ distribution of N_{yA} given $(N_{yA} > 0)$ does not depend on x, and is identical to the $\mathbb{P}_y$ distribution of N_{yA} which is geometric (e_{yA}) on $\{1, 2, \ldots\}$.

These formulas for the distribution of N_{yA} starting from varous states x are just variations of the results presented in Durrett's text for $T = \infty$ instead of $T = T_A$ in formula (1.5) on page 18 and Lemma 1.11 on page 19. These formulas show that the path of a recurrent Markov chain stopped when it first hits some set of states A is full of counting variables N_{yA} with modified geometric distributions, whose parameters can be read from the pre-T_A occupation matrix G_A.
This analysis is of particular interest when $A = \{a\}$ for a single state a. Then the pre-T_a occupation matrix

$$G_a(x,y) = \mathbb{E}_x N_{ya} = \mathbb{E}_x \sum_{k=0}^{T_a - 1} \mathbb{1}(X_n = y) \tag{10.44}$$

gives the expected number of hits on y before T_a, starting from any state x. However, if there are for instance a finite number N of states, for each target state a a different $(N-1)\times(N-1)$ submatrix of $I-P$ must be inverted to obtain the matrix G_a. Or if it is desired to obtain the full matrix of mean first passage times $(m_{xa}, x, a \in S)$, a corresponding system of $N-1$ linear equations (10.40) must be solved to obtain the m_{xa} for $x \neq a$, and thence $m_{aa} = 1 + \sum_{x\neq a} P(a,x)m_{xa}$. But to do this for all states $a \in S$ involves repeating this computational process N times over.

10.4.5 The fundamental matrix of a positive recurrent chain

There is a much more efficient way to evaluate all the Green matrices $G_a(x,y)$ and first passage moments $m_{xa} = G_a(x,\cdot)\mathbb{1}$ of a positive recurrent Markov chain with transition matrix P. This involves the following three step process:

- Step 1. Find the invariant probability vector π. This can be done more or less quickly, depending on the structure of P. Before getting into linear algebra, it is best to first try the shortcuts for this:
  - check the detailed balance equations to see if there is a reversible equilibrium, as if there is it will be easy to find.
  - check if P is doubly stochastic, in which case the uniform distribution is invariant.
  - check if there is some family of initial distributions λ for which it is particularly easy to find λP, and look for π amongst these initial distributions (as in Problem 2 of Worksheet 5).

  Worst case, if none of these shortcuts works,
  - either solve the system of equations $\pi P = \pi$ and $\pi\mathbb{1} = 1$ by linear algebra,
  - or compute the pre-T_a occupation kernel G_a for some fixed state a by linear algebra, and evaluate $\pi(x) = G_a(a,x)/G_a(a,\cdot)\mathbb{1}$.
- Step 2. Set $\Pi := \mathbb{1}\pi$, and invert $I - P + \Pi$ to obtain the *fundamental matrix*

$$Z := (I - P + \Pi)^{-1} \tag{10.45}$$

- Step 3. Read the Green matrices and first passage moments from the simple formulas for all these quantities in terms of entries of Z, as detailed below.

The method of Steps 2 amd 3 will now be motivated by studying the relation between of the matrix $I - P + \Pi$ and the Green matrix G_a for some arbitrary fixed target state a, assuming for simplicity that the state space S is finite. But it is known that all the formulas obtained in this case are valid for any irreducible and positive recurrent chain, just with finite sums replaced by infinite sums, with some care to ensure the infinite sums are convergent. Further motivation for the study of the fundamental matrix Z derived from an irreducible positive recurrent Markov matrix P is provided by the fact that Z is involved in evaluating the variance of the sum

$\sum_{k=1}^{n} f(X_k)$ appearing in the ergodic theorem (10.33), which is needed to provide normal approximations to the distribution of this sum for large n. Here is a more formal statement of Step 2.

Theorem 1 (Existence of the fundamental matrix Z)**.** *Let P be an irreducible transition matrix with finite state space S. Let I be the $S \times S$ identity matrix, let π with*

$$\pi P = \pi \text{ and } \pi \mathbb{1} = 1$$

be the unique invariant probability measure for P, and let $\Pi := \mathbb{1}\pi$ be the $S \times S$ transition matrix with all rows equal to Π. Then

1. *The matrix $I - P + \Pi$ is invertible, with inverse Z that is uniquely determined by either one of the two matrix equations*

$$(I - P + \Pi)Z = I = Z(I - P + \Pi) \tag{10.46}$$

2. *This matrix Z shares with P the basic properties*

$$\pi Z = \pi, \qquad Z\mathbb{1} = \mathbb{1}, \qquad \text{hence} \qquad \Pi Z = Z\Pi = \Pi \tag{10.47}$$

3. *There is the identity $(I - P)^n = P^n - I$ for $n \geq 1$, but not for $n = 0$. If P is aperiodic, then*

$$Z = \sum_{n=0}^{\infty} (P - I)^n = I + \sum_{n=1}^{\infty} (P^n - \Pi) \tag{10.48}$$

   meaning that if $Z_N := \sum_{n=0}^{N} (I - P)^n$ is the matrix of partial sums up to N, then $Z(x, y) = \lim_{N \to \infty} Z_N(x, y)$ for all $x, y \in S$.

4. *If P is periodic, the parital sums Z_N of the series* (10.48) *series are not convergent. But whatever the period of P, the sum can be evaluated by Abel's method:*

$$Z = \lim_{s \uparrow 1} \sum_{n=0}^{\infty} (I - P)^n s^n = I + \lim_{s \uparrow 1} \sum_{n=1}^{\infty} (P^n - \Pi) s^n \tag{10.49}$$

Proof. By linear algebra, a square matrix M is invertible iff there is no non-trivial linear relation among its column vectors: that is, the equation $Mg = 0$ for an unknown column vector g has unique solution $g = 0$. So to establish the existence of Z it is enough to show that the equation

$$(I - P + \Pi)g = 0$$

for an unknown column vector g implies $g = 0$. Pre-multiply this equation by the row vector π to see it implies $\pi g = 0$, hence $\Pi g = 0$, and $(I - P)g = 0$. So g is a P-harmonic column vector with

$$g = Pg = P^n g \qquad (n = 1, 2, \ldots)$$

But for an irreducible P with finite state space, every P-harmonic g is constant: $g = m\mathbb{1}$ for some scalar multiplier m, and $\pi g = 0$ forces $m = 0$, hence $g = 0$. To confirm that $g = m\mathbb{1}$ for some m, let $m := \min_{y \in S} g(y)$ and let x be a state with $g(x) = m$. Then

$$0 = g(x) - m = (P^n[g - m\mathbb{1}])(x) = \sum_y P^n(x,y)[g(y) - m] \geq 0.$$

Hence $P^n(x,y)[g(y) - m] = 0$ for every state y. By irreducibility of P, there is an n such that $P^n(x,y) > 0$, hence $g(y) = m$. So $g = m\mathbb{1}$. This proves the first assertion that $I - P + \Pi$ is invertible. The remaining assertions are then easily checked, using $\Pi P = P\Pi = \Pi$ to show that any finite product of factors P and Π containing at least one Π reduces to Π, and hence $(P - I)^n = P^n - \Pi$ for $n \geq 1$ by using the binomial expansion of $0 = (1 - 1)^n$ to identify the coefficient of Π in the expansion of $(P - I)^n$. In the aperiodic case, the difference $|(P^n - \Pi)(x,y)|$ is bounded by some constant time ρ^n with $0 < \rho < 1$ by the estimate in Durrett's proof of (10.36) on page 53. This implies the series in (10.48) is absolutely convergent, hence that Z defined by (10.48) satisfies (10.46). The discussion of the Abel sum (10.49) in the periodic case just requires a bit more analysis. □

Corollary 1.1 (Solution of Poisson equations)**.** *For a finite state irreducible transition probability matrix P, with $\Pi = \pi\mathbb{1}$ and $Z = (I - P + \Pi)^{-1}$ as above,*

(a) [Poisson equation for a column vector] *For each given column vector f*

$$\text{there exists a solution } g \text{ of } (I - P)g = f \qquad \Longleftrightarrow \qquad \pi f = 0 \tag{10.50}$$

in which case, for each scalar m

$$\text{the unique solution } g \text{ of } (I - P)g = f \text{ and } \pi g = m \text{ is } g = Zf + m\mathbb{1}. \tag{10.51}$$

(b) [Poisson equation for a row vector] *For each given row vector δ*

$$\text{there exists a solution } \mu \text{ of } \mu(I - P) = \delta \qquad \Longleftrightarrow \qquad \delta\mathbb{1} = 0 \tag{10.52}$$

in which case, for each scalar m

$$\text{the unique solution } \mu \text{ of } \mu(I - P) = \delta \text{ and } \mu\mathbb{1} = m \text{ is } \mu = \delta Z + m\pi. \tag{10.53}$$

Proof. Dealing with the case of row vectors, suppose that μ is a solution of the Poisson equation $\mu(I - P) = \delta$ for some given vector δ. Post-multiply the Poisson equation by $\mathbb{1}$ to see that $\delta\mathbb{1} = 0$ is necessary for existence of any solution μ. Continuing to assume that $\mu(I - P) = \delta$, let $m := \mu\mathbb{1}$, so $(\mu - m\pi)\Pi = 0$ and the Poisson equation gives $(\mu - m\pi)(I - P) = \delta$. Add these two equations to get

$$(\mu - m\pi)(I - P + \Pi) = \delta \qquad \text{hence} \qquad \mu - m\pi = \delta Z$$

by post-multiplication by $(I-P+\Pi)^{-1} = Z$. Finally, these steps are easily reversed to show for every δ with $\delta\mathbb{1} = 0$ and every scalar m that $\mu := \delta Z + m\pi$ solves $\mu(I-P) = \delta$ and $\mu\mathbb{1} = m$. The corresponding result for column vectors is obtained by a dual argument, using pre-multiplication of the Poisson equation $(I-P)g = f$ by π to obtain the necessary condition $\pi f = 0$. □

Note the striking duality exposed by the above analysis of the two Poisson equations

$$(1-P)g = f \qquad \text{and } \mu(I-P) = \delta$$

according to whether the transition matrix P is regarded as an operator on the left of column vectors g or on the right of row vectors μ. For each action of P, the Poisson equation with 0 on the right side is solved by any P-invariant vector:

- the only P-invariant column vectors are the constant function $g = m\mathbb{1}$ for some scalar m;
- the only P-invariant row vectors are scalar multiples $\mu = m\pi$ of the invariant probability π.

In each case, the solution of the Poisson equation is clearly unique only up to addition of P-invariant vector, that is some multiple of $\mathbb{1}$ or of π as the case may be. The key points are that for each given column vector f or row vector δ on the right side of the Poisson equation, subject to the obvious necessary condition for existence of a solution (that is $\pi f = 0$ or $\delta\mathbb{1} = 0$),

- a specific solution of the Poisson equation is obtained as either Zf or δZ;
- the most general solution of the Poisson equation is then either $Zf + m\mathbb{1}$ or $\delta Z + m\pi$, as the case may be.

Corollary 1.2 (Occupation measures derived from the fundamental matrix)**.** *For an irreducible Markov chain with transition matrix P with invariant probability π and fundamental matrix $Z = (I-P+\Pi)^{-1}$, let T be any randomized stopping time of the chain, and for a given initial distribution λ let λG_T be the pre-T occupation measure*

$$\lambda G_T(z) := \mathbb{E}_\lambda \sum_{n=0}^{T-1} 1(X_n = z) \qquad (z \in S) \tag{10.54}$$

regarded as a row vector, with total mass $\lambda G_T \mathbb{1} = \mathbb{E}_\lambda T$. Assuming $\mathbb{E}_\lambda T < \infty$, the occupation measure λG_T is related to the initial distribution λ and the final distribution λ_T of X_T by the Poisson equation

$$\lambda G_T(I-P) = \lambda - \lambda_T \tag{10.55}$$

whose solution is

$$\lambda G_T = (\lambda - \lambda_T)Z + (\mathbb{E}_\lambda T)\pi \tag{10.56}$$

Proof. The Poisson equation (10.55) for λG_T is read from the general occupation measure identity (10.26), using $\mathbb{E}_\lambda T < \infty$ to justify finiteness of the measure λG_T. Then (10.56) is read from the general form (10.53) of the solution of the Poisson equation. □

For the stopping time $T = T_y := \min\{n \geq 1 : X_n = y\}$ this formula (10.56) establishes a close connection between the fundamental matrix Z and the pre-T_y occupation matrix G_y with (x, z) entry

$$G_y(x, z) := \mathbb{E}_x N_{zy} \text{ where } N_{zy} := \sum_{n=0}^{T_y - 1} \mathbb{1}(X_n = z)$$

is the number of hits on z strictly before T_y, counting a hit at time $n = 0$ if $x = z$. Abbreviate

$$m_{xy} := \mathbb{E}_x T_y = \sum_z G_y(x, z) = G_y(x, \cdot)\mathbb{1}$$

for the mean first passage time from x to y, and

$$m_{\lambda y} := \mathbb{E}_\lambda T_y = \sum_x \lambda(x) m_{xy} = \sum_z (\lambda G_y)(z) = \lambda G_y \mathbb{1}$$

for the mean first passage time to state y starting from $X_0 \sim \lambda$. Formula (10.56) in this case reads

$$\lambda G_y(\cdot) = \lambda Z(\cdot) - Z(y, \cdot) + m_{\lambda y}\pi(\cdot) \tag{10.57}$$

The choice of the stationary initial distribution $\lambda = \pi$ is of special interest. The feature of Z noted in (10.47) that $\pi Z = \pi$ makes (10.57) for $\lambda = \pi$ simplify to

$$\pi G_y(\cdot) = -Z(y, \cdot) + (1 + m_{\pi y})\pi(\cdot) \tag{10.58}$$

which rearranges as

$$Z(y, \cdot) = (1 + \pi G_y \mathbb{1})\pi(\cdot) - \pi G_y(\cdot). \tag{10.59}$$

Thus the fundamental matrix $Z = (I - P + \Pi)^{-1}$ can be described in probabilistic terms as follows:

- row y of Z is a scalar multiple $m\pi$ of the stationary probability π, minus the pre-T_y occupation measure for the chain started with $X_0 \sim \pi$, where $m = 1 + \mathbb{E}_\pi T_y = 1 + \pi G_y \mathbb{1}$ is determined by the row sum $Z(y, \cdot)\mathbb{1} = 1$.

In particular, evaluating the row vector (10.59) at y gives the diagonal entries of the fundamental matrix Z in terms of π and mean first moments:

$$Z(y, y) = (\mathbb{E}_\pi T_y)\pi(y) = \sum_x \pi(x) m_{xy} \pi(y) = \frac{\mathbb{E}_y T_y^2 + m_{yy}}{2m_{yy}^2} \tag{10.60}$$

where $m_{yy} = \mathbb{E}_y T_y$ and the last equality is due to Kac's identity (Problem 3 of Worksheet 5)

$$\mathbb{P}_\pi(T_y = n) = \mathbb{P}_y(T_y \geq n)/m_{yy} \qquad (n = 1, 2, \ldots) \tag{10.61}$$

which shows how the distribution of T_y for $X_0 \sim \pi$ determines the distribution of T_y given $X_0 = y$ and vice versa. As a checks on these formulas: summing (10.61) over n gives 1 on the left side by the assumed recurrence of the chain, and 1 on the right side by the tail sum formula for $\mathbb{E}_y T_y$. Also, (10.60) can be checked by renewal theory, using the infinite sum (10.48) and probability generating functions, as indicated in Problem 5 of Worksheet 5.
Similarly, taking $\lambda = \delta_x$ in (10.57) gives

$$Z(x,y) - Z(y,y) = \mathbb{1}(x = y) - m_{xy}\pi(y). \tag{10.62}$$

Hence, by adding (10.60) and (10.63), there is a general formula for entries of Z in terms of mean first passage times:

$$Z(x,y) = \mathbb{1}(x = y) + (m_{\pi y} - m_{xy})\pi(y). \tag{10.63}$$

The case $x = y$ of (10.63) reduces to (10.60) by $\pi(x) = 1/m_{xx}$, which allows the diagonal mean first passage times m_{xx} to be determined from π. Subtract (10.63) from (10.60) to see that similarly, all the off diagonal mean first passage times can be read from the fundamental matrix Z and π:

$$m_{xy} = \frac{Z(y,y) - Z(x,y)}{\pi(y)} \qquad (x \neq y) \tag{10.64}$$

As a check on these formulas, the simple special case that P governs a sequence of i.i.d. random variables with distribution π corresponds to

$$P = \Pi \iff Z = I \iff m_{xy} = m_{yy} \text{ for all } x. \tag{10.65}$$

10.4.6 Exercises.

1. Deduce that

$$G_y(x,z) = (m_{xy} - m_{xz}1(x \neq z) + m_{yz}1(y \neq z))\pi(z). \tag{10.66}$$

This can be rearranged as

$$m_{xy} + m_{yz}1(y \neq z) = G_y(x,z)m_{zz} + m_{xz}1(x \neq z) \tag{10.67}$$

which is just two different ways of calculating $\mathbb{E}_x T_{yz}$ where T_{yz} is the first visit to z at or after time T_y. See Chapter 2 of the Aldous-Fill book `https://www.stat.berkeley.edu/users/aldous/RWG/book.html` for further discussion.

2. Check directly as follows, without discussion of the Poisson equation, that Z defined by (10.63) is the inverse of $I - P + \Pi$. Recall that by a first step analysis, for all pairs of states x and z, including $x = z$,

$$m_{xz} = 1 + \sum_{y \neq z} P(x,y)m_{yz} \tag{10.68}$$

For Z defined by (10.63), use (10.68) to evaluate the $(PZ)(x,z)$, and show that $PZ = Z - I + \Pi$. Check also that $\Pi Z = \Pi$, and hence that $Z(I - P + \Pi) = I$.

3. Evaluate all entries of Z explicitly for a two state chain.

4. Evaluate one row of Z explicitly for a random walk on three states in a circle, with probabilities p and q for clockwise and anticlockwise steps, with $p+q=1$. How can the other rows of Z be derived from this row with no further calculation?

5. (Asmussen Proposition 7.1 [4]) Show that for each f with $\pi f = 0$ the unique solution g of the Poisson equation $(I-P)g = f$ with $g(y) = 0$ is

$$g(x) = \delta_x G_y f = \mathbb{E}_x \sum_{n=0}^{T_y - 1} f(X_n)$$

6. Is there a corresponding description of the solution of $\mu(I-P) = \delta$ with $\delta \mathbb{1} = 1$ and $\delta(y) = 0$?

10.5 References

The material on pre-T occupation measures is based on my *Occupation measures for Markov chains* Adv. Appl. Probab. 9, 69-86 (1977). See Asmussen's *Applied probability and queues* (§I.7[4]) and Chapter 2 of the Aldous-Fill book `https://www.stat.berkeley.edu/users/aldous/RWG/book.html` for further discussion of the fundamental matrix Z. Beware that Aldous-Fill use the notation Z for $Z - \Pi$, and T_x^+ for our T_x and T_x for our V_x.

LECTURE 11

Poisson Processes, Part 1

11.1 Introduction: Poisson Processes

We'll do a quick review of basic properties of the Poisson distribution. Recall that we have the **Binomial**(n, p)distribution of

$$S_n(p) := X_1(p) + X_2(p) + \cdots + X_n(p),$$

where the $X_i(p)$ are independent **Bernoulli**(p) variables. We have a nice construction of all these variables at once. Take $U_1, U_2, \ldots$ iid over interval $[0, 1]$ and let

$$X_n(p) := \mathbb{1}\{U_n \leq p\}$$

Observe

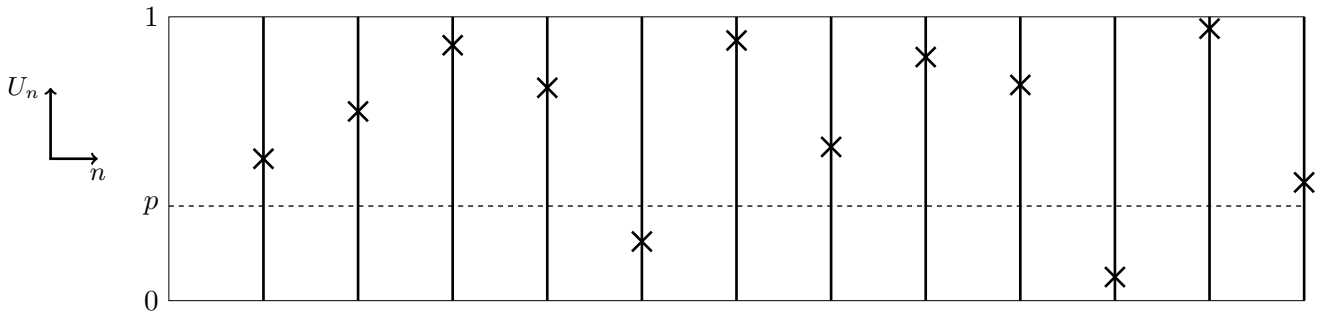

n	1	2	3	4	5	6	7	8	9	10	11	12
$X_n(p)$	0	0	0	0	1	0	0	0	0	1	0	0

Now we look at when $n \to \infty$, with $p \downarrow 0$, so that $np \equiv \mu$ is fixed. In terms of the above construction, we are simply lowering the dotted line at level p, and pushing

$n = \mu/p$ to infinity (through integer values, or round μ/p to the nearest integer, it makes no difference in the limit), to keep the expected number of points $\times$ below the bar in n trials to equal μ, either exactly for simplicity as in the following calculations, or approximately enough so that $np \to \mu$ as both $n \to \infty$ and $p \downarrow 0$. Working with $np \equiv \mu$ exactly

$$\mathbb{E}S_n(p) = np \equiv \mu$$

so that

$$\begin{aligned}
\mathbb{P}(S_n(p) = k) &= \binom{n}{k} p^k (1-p)^{n-k} \\
&= \frac{(n)_{k\downarrow}}{k!} \left(\frac{\mu}{n}\right)^k \left(1 - \frac{\mu}{n}\right)^{n-k} \\
&\to e^{-\mu} \frac{\mu^k}{k!} \text{ as } n \to \infty \text{ and } p \downarrow 0
\end{aligned}$$

This is how the Poisson distribution arises as the limit of binomial distributions with large n and small p. A more careful argument (see text Theorem 2.9 on page 104) shows that for any collection of independent indicator variables $X_1, \ldots, X_n$ with $p_k = \mathbb{P}(X_k = 1)$ with sum S_n so $\mathbb{E}S_n = p_1 + \cdots + p_n\mu$, the **Poisson** (μ) distribution will be a good approximation to the distribution of S_n whenever $\mu(\max_k p_k)$ is small. In fact the total variation distance between this distribution of S_n (depending on $p_1, \ldots, p_n$) and **Poisson** (μ) is at most $\mu(\max_k p_k)$, which for any fixed μ is small provided all the p_k are sufficiently small. Now when $\mu \geq 0$, let N_μ or $N(\mu)$ denote a random variable with this **Poisson** (μ) limit law

$$\mathbb{P}(N_\mu = k) = e^{-\mu} \frac{\mu^k}{k!} \quad k = 0, 1, 2, \ldots$$

Some basics

$$\mathbb{E}N_\mu = \mu \quad \text{and} \quad \mathbb{V}\text{ar}(N_\mu) = \mu$$

Notice that in the **Binomial**(n, p) setup

$$\mathbb{V}\text{ar}(S_n(p)) = np(1-p) = \mu(1-p) \to \mu$$

as $p \downarrow 0$. So the variance for Poisson is the limit of binomial variances, as should be expected. You should check these moment formulas by summation. Additionally, by probability generating functions, we have

$$G_{N(\mu)}(z) := \mathbb{E}z^{N(\mu)} = \sum_{n=0}^{\infty} z^n e^{-\mu} \frac{\mu^n}{n!} = e^{-\mu} e^{\mu z} = e^{-\mu(1-z)}$$

Take the derivatives $\frac{d}{dz}, \frac{d^2}{dz^2}$ at $z = 1$. Gives us $\mathbb{E}N_\mu$ and $\mathbb{E}N_\mu(N_\mu - 1)$, hence the formula for variance.

Exercise There is a pretty formula for $\mathbb{E}\binom{N_\mu}{k}$. Find it and prove it.

11.2 Sum of Independent Poissons

Take $N_1, N_2, \ldots, N_m$ independent Poissons with means $\mu_1, \mu_2, \ldots, \mu_m$. Then

$$N_1 + \cdots + N_m \sim \textbf{Poisson}\,(\mu_1 + \cdots + \mu_m)$$

This is intuitively clear as a limit of the corresponding result for sums of independent binomials with parameters (n_i, p) with $n_i p \equiv \mu_i$ and $p \downarrow 0$. Alternatively, we can easily derive it with probability generating functions. Or by induction on m, from the case $m = 2$, using the convolution formula and the binomial theorem to evaluate

$$\mathbb{P}(N_1 + N_2 = n) = \sum_{k=0}^{n} \mathbb{P}(N_1 = k)\mathbb{P}(N_2 = n - k)$$

11.3 Poissonization of the Multinomial

What is the distribution of the following random vector?

$$\Big((N_1, N_2, \ldots, N_m)\,|\,N_1 + N_2 + \cdots + N_m = n\Big)$$

for independent Poisson(μ_i) variables N_i? For $m = 2$, the last computation suggested above shows as a byproduct that

$$(N_1 \mid N_2 = n) \stackrel{d}{=} \textbf{Binomial}(n, p) \text{ for } p = \mu_1/(\mu_1 + \mu_2).$$

For general m, we can compute the above conditional distribution using Bayes rule. We should perform this computation once in our lives, as this is the basis of the entire theory of Poisson processes. It is only of interest to consider $n_1, \ldots, n_k$ with $n_1 + \cdots + n_m = n$, the given value for the sum. So consider

$$\begin{aligned}
&\mathbb{P}(N_1 = n_1, N_2 = n_2, \ldots, N_m = n_m \mid N_1 + N_2 + \cdots + N_m = n) \\
&\qquad = \frac{\mathbb{P}(N_1 = n_1, N_2 = n_2, \ldots, N_m = n_m)}{\mathbb{P}(N_1 + \cdots + N_m = n)} \\
&\qquad = \frac{\dfrac{e^{-\mu_1}\mu_1^{n_1}}{n_1!} \cdots \dfrac{e^{-\mu_m}\mu_m^{n_m}}{n_m!}}{\dfrac{e^{-(\mu_1+\cdots+\mu_m)}(\mu_1 + \cdots + \mu_m)^{n_1+\cdots+n_m}}{(n_1 + \cdots + n_m)!}} \\
&\qquad = \frac{n!}{n_1! \cdots n_m!}\, p_1^{n_1} p_2^{n_2} \cdots p_m^{n_m}
\end{aligned}$$

Notice how the exponentials cancel across the numerator and denominator, with $\mu = \mu_1 + \cdots + \mu_m$ and $p_k = \frac{\mu_k}{\mu}$. We recognize this as the familiar *multinomial* distribution. Hence for $N = N_1 + \cdots + N_m$ as before, we have

$$(N_1, \ldots, N_m \,|\, N = n) \stackrel{d}{=} \textbf{Multinomial}(n, p_1, p_2, \ldots, p_m)$$

In words, any vector of independent Poisson counts $N_1, \ldots, N_m$ conditioned on the total count equal to n is distributed like the counts of values with probabilities $(p_1, \ldots, p_m)$ in n Multinomial trials. To state the conclusion more formally:

Poissonization of the Multinomial

The following two conditions on a random vector of counts $N_1, \ldots, N_m$ are equivalent:

(1) the $N_1, \ldots, N_m$ are independent Poissons with means $\mu_1, \ldots, \mu_m$.

(2) $N_1 + \cdots + N_m$ is **Poisson** $(\mu_1 + \cdots + \mu_m)$ and given $N_1 + \cdots + N_m = n$, the vector $(N_1, \ldots, N_m)$ is multinomial$(n, p_1, \ldots, p_m)$ with $p_k = \frac{\mu_k}{\mu_1 + \cdots + \mu_m}$.

This statements is both important and easy to check once formulated. It is not obvious at first, but it becomes very familiar and quite intuitive as you work with Poisson processes.

Corollary

Suppose we have $Y_1, Y_2, \ldots$ iid with probability distribution

$$\mathbb{P}(Y_i = k) = p_k,$$

for some probability distribution $(p_1, \ldots, p_m)$ on $\{1, \ldots, m\}$. Assume that N is independent of $Y_1, Y_2, \ldots$ and $N \sim$ **Poisson** (μ). Define

$$N_k := \sum_{i=1}^{N} \mathbb{1}(Y_i = k),$$

which is the number of Y values equal to k in the N trials (hence equal to 0 if $N = 0$).
Then $N_1, N_2, \ldots, N_m$ are independent **Poisson** $(\mu_1, \ldots, \mu_m)$.

Proof. By constuction, $N_1 + N_2 + \cdots + N_m = N$, and given $N = n$, the $(N_1, \ldots, N_m)$ are multinomial$(n, p_1, \ldots, p_k)$. Now plug into the theorem with $\mu + k = p_k \mu$ so that $\mu_1 + \cdots \mu_m = \mu$. □

Remark It is not really necessary to have an infinite iid sequence Y_i to construct counts N_k as above. All that is needed is a Poisson variable N, and for each $n \geq 1$,

conditionally given $N = n$ a sequence $Y_1, \ldots, Y_n$ of iid variables with the required distribution $(p_1, \ldots, p_m)$. Then the same conclusion holds, for the same reasons.

To summarize, if in multinomial trials, we randomize the number of trials N with a Poisson distribution, we make the counts independent. This is highly non obvious at first, but this is a key fact which will be exploited heavily for Poisson Processes. See the text Theorem 2.1 on page 108, Theorem 2.12 on page 110 and Theorem 2.15 on page 115 for several slight variations on this theme, expressed in terms of Poisson processes. The underlying distributional relation that makes all these results work is the Poissonization of the multinomial.

11.4 Poisson Point Processes (PPP)

We'll presume that we've seen the Poisson processes on a line. It is more interesting and intuitive to start looking at a PPP in a strip of the plane. The setting is as follows

- Let $\Delta_i \sim \textbf{Poisson}\,(\lambda)$ be the Poisson count for box i with intensity $\lambda > 0$, some fixed rate per unit area.

- If N_t is the number of points in $[0, t] \times [0, 1]$, then it follows

- $N_t = \Delta_1 + \Delta_2 + \ldots + \Delta_t \sim \textbf{Poisson}\,(\lambda t)$ for $t \in \mathbb{Z}^+$.

We can obtain insight for the likely number of counts for $\Delta_i \sim \textbf{Poisson}\,(2)$, by looking at its histogram. Take note of the the fact that for a random variable with distribution $\textbf{Poisson}\,(m)$, the values with highest probability occur at m and $m - 1$.

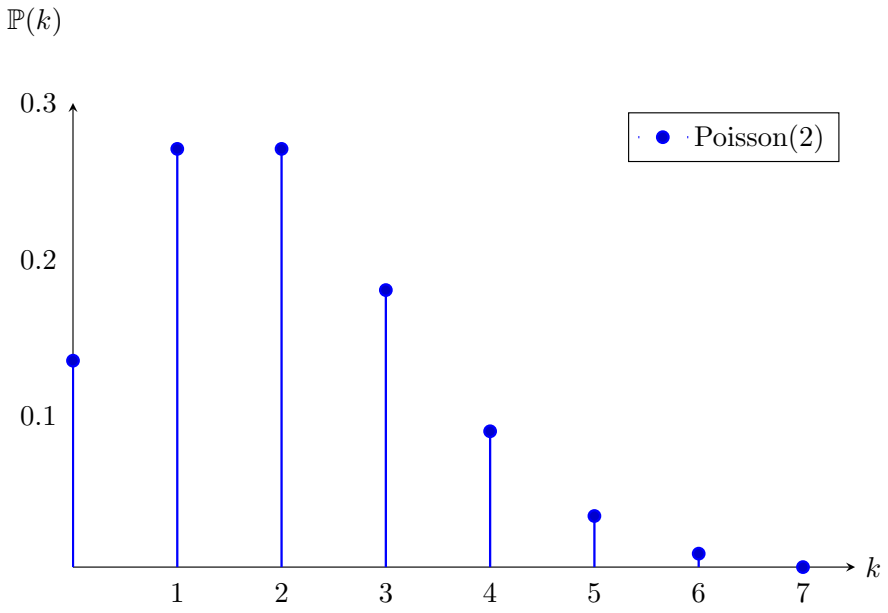

With this in mind, given $\Delta_i = n$, we throw down n iid points with uniform probability in the i^{th} square $[i, i+1] \times [0, 1]$. Study the following diagrams.

11.4.1 PPP Strips

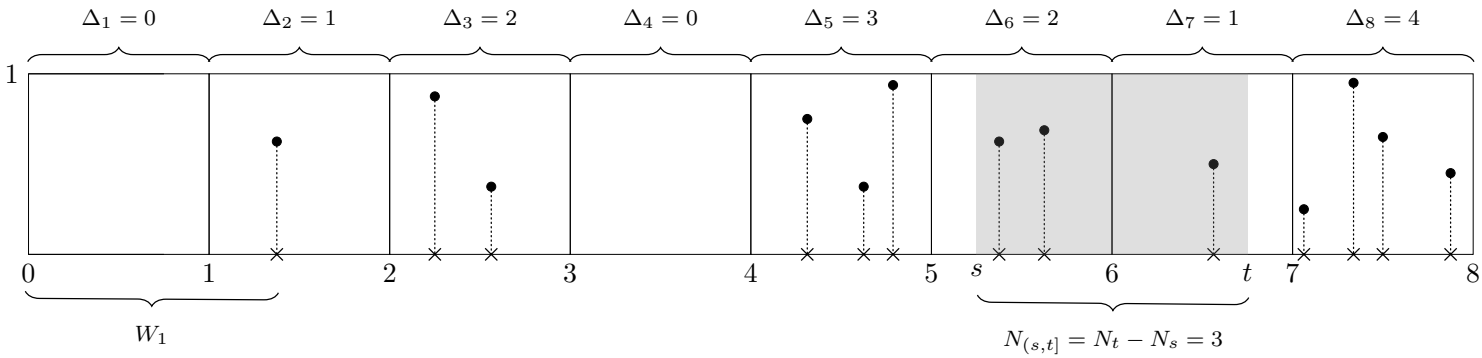

Figure 11.1: We project each "dot" onto the horizontal axis. This exploits two things: (1) the time between arrivals, the first shown below the strip, denoted W_1 and (2) the counts of arrivals between an interval say $(s, t]$. Recall that N_t denoted the counts before and including time t. Hence for $N_{(s,t]}$ we simply take the difference between N_t and N_s.

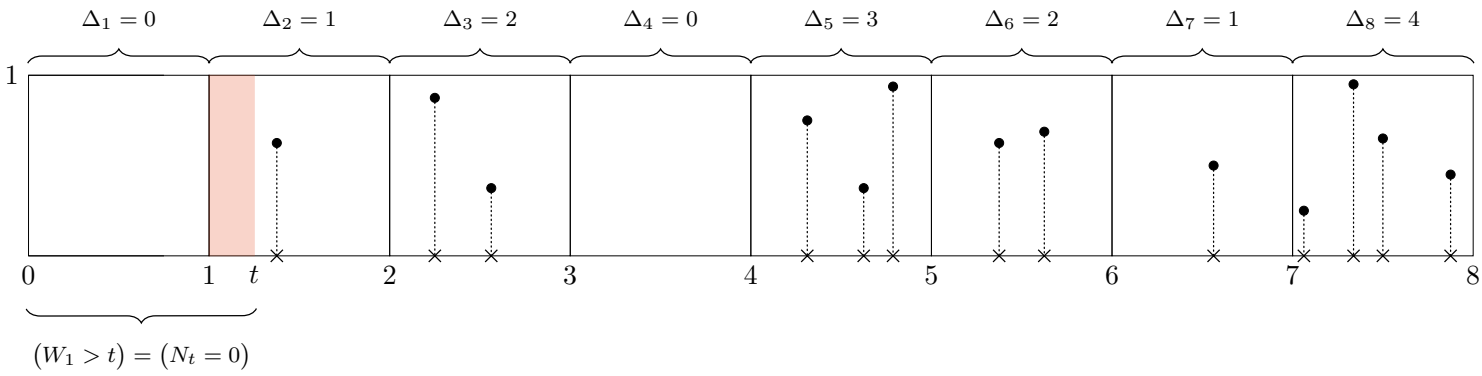

Figure 11.2: Notion of duality. In particular, the first hit arriving after time t, the event $W_t > 0$ is the same as $N_t = 0$.

For example, take independent uniform $[0, 1]$ variables $X_{i,j}$ and $Y_{i,j}$, independent of all the N_i, and for each $i = 1, 2, \ldots$ place points at

$$(i - 1 + X_{i,j}, Y_{i,j}) \text{ for } 1 \leq i \leq N_{i,j}$$

with of course no points at all in the ith square if $N_i = 0$
Now Pitman asks what happens if we project this strip down onto a line. Notice that the probability of projecting two or more points onto the same point on the line is 0. Or in other words, there are no multiple points (repeated X-values) in the strip. Let $W_1, W_2, W_3, \ldots$ be the spacings between points along the X-axis.

Case: $0 < t < 1$

Notice $\mathbb{P}(W_1 > t)$ is simply the probability that there there are no points to the left of t.

$$\mathbb{P}(W_1 > t) = \mathbb{P}(N_{\lambda t} = 0) = e^{-\lambda t},$$

by design and the Poissonization of the binomial.

Case: $1 < t < 2$

Here, we use independence to have:

$$\begin{aligned}
\mathbb{P}(W_i > t) &= \mathbb{P}(N_1 = 0 \text{ and count in } [1, t] \times [0, 1] = 0) \\
&= \mathbb{P}(N_1 = 0)\mathbb{P}(\text{count in } [1, t] \times [0, 1] = 0) \\
&= e^{-\lambda} e^{-\lambda(t-1)} \\
&= e^{-\lambda t}.
\end{aligned}$$

This result makes us very happy, and we claim this can be continued inductively for arbitrary $t > 0$.

Repeating this discussion for counts

Let N_t be the number of points to the left or equal to t. Then we simply have:

$$N_t \sim \textbf{Poisson}\,(\lambda t).$$

This is obvious for integer $t = 1, 2, \ldots$, and true also for non-integer $t > 0$ by design and Poissonization of the binomial. Now consider $0 \le s \le t$. Observe that $N_t - N_s$ is the number of points in $(s, t]$. If s, t are integers, it is obvious that $N_t - N_s$ is **Poisson** $(\lambda(t - s))$ independent of N_s. Now if they are not integers, this still works via the Poissonization of the binomials involved in their fractional parts. That is, $N_t - N_s$ is the number of points in $(s, t]$. We have that

$$N_t - N_s \sim \textbf{Poisson}\,(\lambda(t - s)),$$

and notice that this count is **independent** of the N_s. Continuing this, we recover the usual definition of the Poisson point distribution on the half-line from 0 to ∞. (Text, page 101). Thus we see that the construction of a PPP in the strip, just indicated, projects to a PPP with constant rate λ on $[0, \infty)$, just be ignoring the Y-values. Of course, it was not necessary even to involve the Y-values, but the pictures are nicer with both X and Y values, and this construction also serves to explain many further properties of Poisson processes (text Section 2.4).
Let $0 \le t_1 < t_2 < t_3 \le \cdots \le t_n$. Then

$$N(t_1), N(t_2) - N(t_1), \ldots, N(t_n) - N(t_{n-1})$$

are independent Poisson with parameters

$$\lambda t_1, \lambda(t_2 - t_1), \ldots, \lambda(t_n - t_{n-1}).$$

Theorem

The Poisson finite dimensional distributions indicated above for a counting process $(N(t), t \ge 0)$ are equivalent to spacings $W_1, W_2, \ldots$ iid **Exponential** (λ), where $W_1, W_1 + W_2, \ldots$ are the arrival times on the X-axis.

This picture has a continuous time axis and a discrete count (vertical) axis. There are formulas that correspond to the picture. Because N_t is a count, we write it as a sum of indicators:

$$N_t = \sum_{n=1}^{\infty} \mathbb{1}(T_n \le t),$$

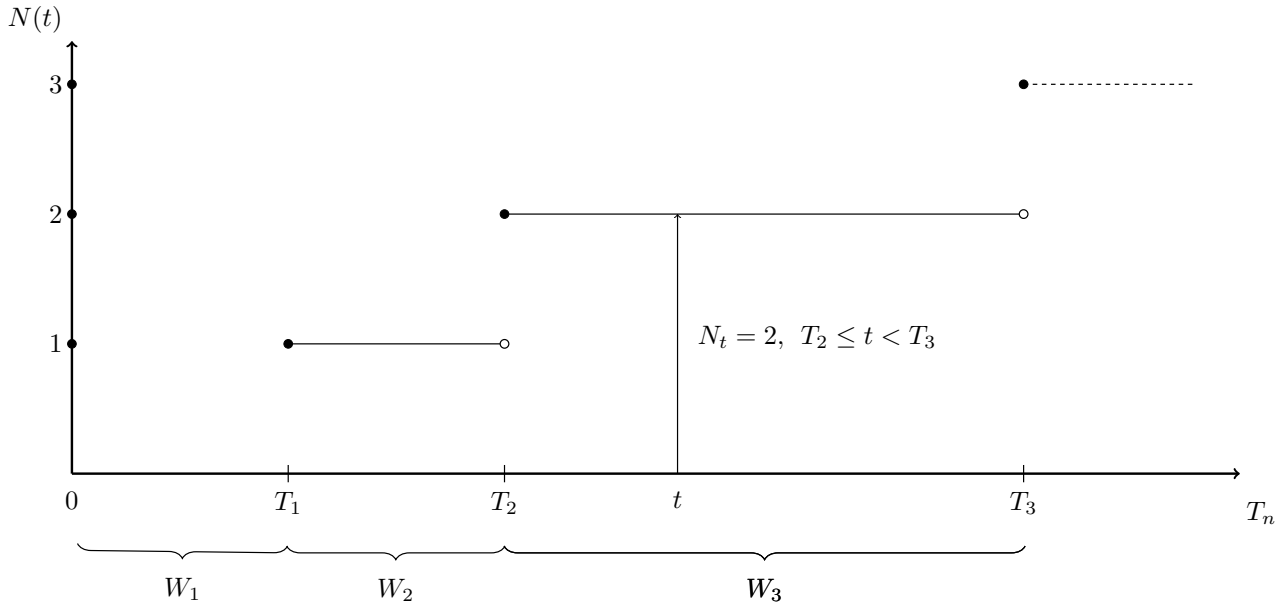

Figure 11.3: Poisson Arrival Process with counting variable $N(t) = N_t$, inter-arrival times $0 < W_1, W_2, \ldots$, and arrivals $T_1 < T_2 < \ldots$

which is simply counting the T_i less than or equal to t. Then for the inverse direction, we have the attained min:

$$T_n = \min\{t : N_t = n\} = \min\{t : N_t \geq n\},$$

where the second expression is generally true even if we jump over a value. Pitman reminds that these are very important formulas. Now we want to notice a key duality implied by the above definitions. Note that this duality is valid for any counting process (N_t) and (T_n) as above, without making any distributional assumptions. It applies for instance also to counting processes, where $T_n = W_1 + \cdots + W_n$ for W_n that might not be independent or not identically distributed.

Key Duality

For counting process N_t and arrival times T_n

$$(T_n > t) = (N_t < n)$$

Equivalently,

$$(T_n \leq t) = (N_t \geq n).$$

To check this, we logically check implication in both ways. In fact, $t \to N_t$ by definition is increasing ($\uparrow$) and right-continuous, and no missing value (that is to

say no repeated points).

11.5 Applications

$$\begin{aligned}\mathbb{P}(T_n \le t) &= \mathbb{P}(N_t \ge n) \\ &= 1 - \mathbb{P}(N_t < n) \\ &= 1 - \sum_{k=0}^{n-1} e^{-\lambda t}\frac{(\lambda t)^k}{k!}.\end{aligned}$$

Pitman asks us to find the density, given the CDF. To do so, we differentiate. This gives

$$\begin{aligned}f_{T_n}(t) &= \frac{d}{dt}\mathbb{P}(T_n \le t) \\ &= \frac{d}{dt}\left(1 - \sum_{k=0}^{n-1} e^{-\lambda t}\frac{(\lambda t)^k}{k!}\right),\end{aligned}$$

which gives a telescoping sum. After evaluating, we get a gamma distribution, $\textsf{Gamma}(n, \lambda)$.

11.6 Secret Method

Pitman gives a secret method so that we are sure to get this right (no book will tell us this). This is much better than performing the telescoping sum.
Look at an infinitesimal interval $[t, t+dt]$, so that

$$\begin{aligned}\mathbb{P}(t \le T_n \le t+dt) &= \mathbb{P}(n-1 \text{ points in } [0,t] \text{ and } \ge 1 \text{ points in } [t, t+dt] + o(dt) \\ &= \left[\frac{e^{-\lambda t}(\lambda t)^{n-1}}{(n-1)!}\right]\left[\frac{\cancel{e^{-\lambda dt}}(\lambda dt)^1}{1!} + o(dt)\right] + o(dt) \\ &= e^{-\lambda t}\frac{(\lambda t)^{n-1}}{(n-1)!}\lambda\, dt + o(dt)\end{aligned}$$

where $o(dt)$ indicates a term which is negligible compared to dt in calculus limit as $dt \to 0$ (for the possibility of more than 1 point in the interval of length dt). So this gives the answer that T_n has probability density

$$f_{T_n}(t) = \frac{e^{-\lambda t}\lambda^n t^{n-1}}{(n-1)!} \qquad (t > 0)$$

which is the $\textsf{Gamma}(n, \lambda)$ density. See also text Theorem 2.2 for the more usual derivation of the same density from the important representation

$$T_n = W_1 + \cdots + W_n$$

for independent exponential λ spacings W_n, denoted τ_n by Durrett.

LECTURE 12

Poisson Processes, Part 2

12.1 Theorem 2.10

Some notation: Pitman will be employing X's instead of Durret's Y's.

Theorem 2.10 (Durrret)

Let X_1, X_2 be a sequence of i.i.d. random variables, each with the same distribution as X, and define

$$S_N = X_1 + X_2 + \ldots + X_N \tag{12.1}$$

where $N \geq 0$ is a random index independent of the sequence $X_1, X_2,$, and summing zero terms yields 0, formally $S_N = 0$ if $N = 0$. Then

(*i*) $\mathbb{E}(S_N) = \mathbb{E}(N)\mathbb{E}(X)$ provided $\mathbb{E}N < \infty$ and $\mathbb{E}|X| < \infty$

(*ii*) $\mathbb{V}\text{ar}(S_N) = \mathbb{E}(N)\mathbb{V}\text{ar}(X) + \mathbb{V}\text{ar}(N)(\mathbb{E}X)^2$
provided $\mathbb{E}N < \infty$ and $\mathbb{E}X^2 < \infty$

(*iii*) If $N \sim$ **Poisson** (μ), then $\mathbb{E}(N) = \mathbb{V}\text{ar}(N) = \mu$. Implying

$$\mathbb{V}\text{ar}(S_N) = \mu\mathbb{V}\text{ar}(X) + \mu(\mathbb{E}X)^2 = \mu\ \mathbb{E}(X^2)$$

Proof. Durret's proof is fine. But here is a shorter derivation. Recall that for any numerical random variable Y with $\mathbb{E}|Y| < \infty$, and any random variable X

$$\mathbb{E}(Y) = \mathbb{E}\Big[\mathbb{E}(Y \,|\, X)\Big] \tag{12.2}$$

In words,

> "The expectation is the expectation of the conditional expectation"

And for any random variable Y with $\mathbb{E}Y^2 < \infty$ and any random variable X

$$\mathbb{V}\text{ar}(Y) = \mathbb{E}\Big[\mathbb{V}\text{ar}(Y \mid X)\Big] + \mathbb{V}\text{ar}\Big[\mathbb{E}(Y \mid X)\Big] \tag{12.3}$$

> "The variance of is the expectation of the conditional variance plus the variance of the conditional expectation"[1]

These formulas apply in an obvious way to our set-up with $Y = S_N$ and $X = N$. Now to the proof.

For conclusion (i)

$$\mathbb{E}(S_N) = \mathbb{E}\Big[\mathbb{E}(S_N \mid N)\Big] = \mathbb{E}\big[N\mathbb{E}(X)\big] = \mathbb{E}(N)\mathbb{E}(X)$$

For conclusion (ii)

$$\begin{aligned}\mathbb{V}\text{ar}(Y) &= \mathbb{E}\Big[\mathbb{V}\text{ar}(S_N \mid N)\Big] + \mathbb{V}\text{ar}\Big[\mathbb{E}(S_N \mid N)\Big] \\ &= \mathbb{E}\Big[N\mathbb{V}\text{ar}(X)\Big] + \mathbb{V}\text{ar}\Big[N\mathbb{E}(X)\Big] \\ &= \mathbb{E}(N)\mathbb{V}\text{ar}(X) + \mathbb{V}\text{ar}(N)(\mathbb{E}X)^2\end{aligned}$$

For conclusion (iii) just plug in the fact that $N \sim$ **Poisson** (μ) has mean and variance both equal to μ. □

12.2 Generalization to a Stopping Time N

In the case where N is not necessarily independent of $X_1, X_2, \ldots$, we study `Wald's Identities`.

12.2.1 Wald's Identities

We have this notion of N being a stopping time for the sequence $X_1, X_2, \ldots$. For example, $N = \min\{n \geq 1 : X_n \in B\}$ or $N = \min\{n \geq 1 : X_n \in S_n\}$. We realize,

- $(N = n)$ is *determined* by $(X_1, X_2, \ldots X_n)$
- $(N \leq n)$ and $(N > n)$ are *determined* by $(X_1, X_2, \ldots X_n)$
- $(N < n)$ and $(N \geq n)$ are *determined* by $(X_1, X_2, \ldots X_{n-1})$

[1] Pitman suggests reciting this 9 times before one sleeps (if you haven't yet internalized it already).

Claim If both $\mathbb{E}N$ and

$$\mathbb{E}|X| < \infty$$

, and N is a stopping time of the sequence $X_1, X_2, \ldots$, then $\mathbb{E}(S_N) = \mathbb{E}(N)\mathbb{E}(X)$ is still true. i.e. It is as if N and X are independent. This is not intuitive at first, we sketch the proof and show the claim is true first for $X \geq 0$.

$$S_N = \sum_{k=1}^{N} X_k = \sum_{k=1}^{\infty} X_k \, \mathbb{1}\{k \leq N\}$$

Notice $(k \leq N) = (N > k-1) = (N \leq k-1)^c$ is determined by $X_1, \ldots, X_{k-1}$, hence independent of X_k. Applying expectation

$$\mathbb{E}(S_N) = \mathbb{E}\left(\sum_{k=1}^{\infty} X_k \, \mathbb{1}\{k \leq N\}\right) = \sum_{k=1}^{\infty} (\mathbb{E}X_k)\mathbb{P}(N \geq k)$$

Since $\mathbb{E}X_k = \mathbb{E}X$ for all x, and exploiting the tail-sum formula, this is

$$(\mathbb{E}X) \underbrace{\sum_{k=1}^{\infty} \mathbb{P}(N \geq k)}_{\text{tail-sum}} = (\mathbb{E}X)(\mathbb{E}N)$$

Here the swap of $\mathbb{E}$ and Σ is justified by non-negativity of all variables. If the X_i are signed, first split each $X_i = X_i^+ - X_i^-$ into the difference of two positive variables, work on the positive and negative parts separately, then argue that both pieces are finite so OK to take their difference.

Note. What really makes this argument work is that for each k the event $N \geq k$ is independent of X_k. One way to achieve this is to assume N is independent of each X_k. Another way is to assume N is a stopping time, so $(N \geq k)$ is determined by $X_1, \ldots, X_{k-1}$. A slight generalization that includes both cases is to assume there is some richer sequence of random variables $(W_0, W_1, \ldots)$ such that

- a) N with values in $\{0, 1, 2, \ldots\}$ is a stopping time relative to $(W_0, W_1, \ldots)$, meaning $(N \leq n)$ is determined by $(W_0, \ldots, W_n)$.
- b) each variable X_n for $n \geq 1$ has the same distribution as X, and X_n is independent of $(W_0, \ldots, W_{n-1})$.

The proof is exactly the same. There is also a companion Wald identity for variances. Under the same assumptions a) and b) above,

- if $\mathbb{E}N < \infty$ and $\mathbb{E}X^2 < \infty$ and $\mathbb{E}X = 0$ then $\mathbb{E}S_N = 0$ and $\mathbb{E}S_N^2 = \mathbb{E}N(EX^2)$

To see this, assume first that N is bounded above by say m, and compute

$$\mathbb{E}(S_N^2) = \mathbb{E}\left(\sum_{j=1}^{m} X_j 1(j \leq N)\right)\left(\sum_{k=1}^{m} X_k 1(k \leq N)\right)$$

be expanding the product. The diagonal terms easily give $\mathbb{E}N(EX^2)$, and the off-diagonal terms are all 0 if $\mathbb{E}X = 0$. Some care is required to pass to the limit of unbounded stopping times N, using convergence in mean square of $S_{N \wedge m}$ to S_N as $m \to \infty$.

12.3 Poisson Thinning

We turn back to looking at the Poisson Point Process (PPP) on the strip $[0, \infty) \times [0, 1]$. Recall that points are distributed with intensity λ per unit area: each square $[n, n+1] \times [0, 1]$ contains a Poisson(λ) number of points, independently from one square to the next, and given m points in one of these squares, these points are distributed i.i.d. according to area in the square.
The general formula for $\mathbb{V}\text{ar}(S_N)$ for a stopping time N when $\mathbb{E}X \neq 0$ is not so simple. Exercise: find the formula, and note it contains an $\mathbb{E}NS_N$ which is easy to evaluate if N is independent of $X_1, X_2, \ldots$, but not always so easy to evaluate.

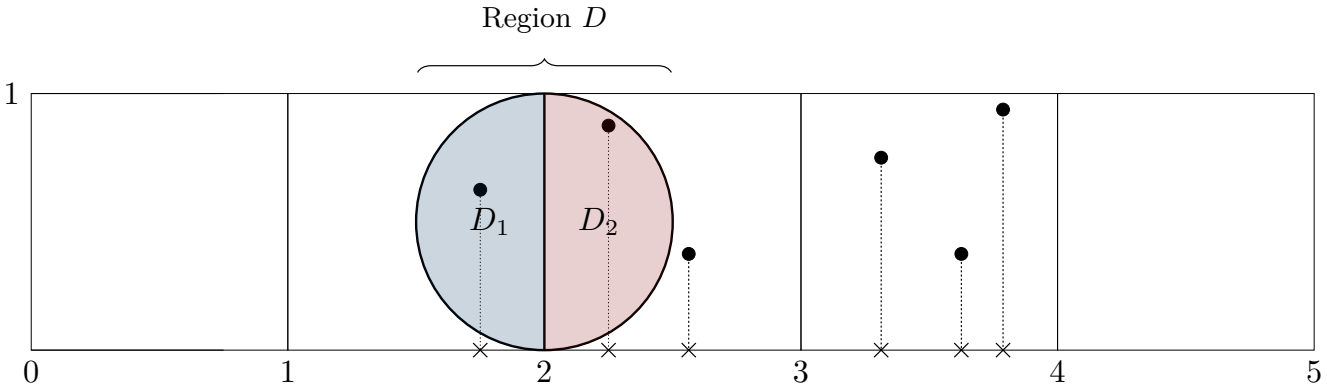

Figure 12.1: PPP on $[0, \infty) \times [0, 1]$ with region D.

Consider the region D, a circle with diameter one shown above. Let $N(D)$ be the number of points in D. So $N(D) = 2$ for the particular realization of the Poisson scatter illustrated. We seek to compute the probability of this event

$$\mathbb{P}\Big(N(D) = 2\Big)$$

We observe

- $N(D) = N(D_1) + N(D_2)$, where D_1 is the left half-disk and D_2 the right half-disk.
- $\text{Area}(D) = \frac{\pi}{4}$, and
- by symmetry, $N(D_1)$ and $N(D_2)$ are i.i.d.

So to re-frame the question,

> Given that a point falls in $(1, 2] \times [0, 1]$, what is the probability that it falls into D?

It is clear that $\text{Area}(D_1) = \frac{1}{2}\left(\frac{\pi}{4}\right)$. We can then see

$$N(D_1) \stackrel{d}{=} N(D_2) \sim \textbf{Poisson}\left(\lambda \times \frac{\pi}{8}\right)$$

That is, we derive the distribution of $N(D_1)$ from the known Poisson (λ) distribution of $N((1, 2] \times [0, 1]$ by Poisson *thinning*. Hence

$$N(D) \sim \textbf{Poisson}\left(\lambda \times \frac{\pi}{8} + \lambda \times \frac{\pi}{8}\right) = \textbf{Poisson}\left(\lambda \times \frac{\pi}{4}\right) = \textbf{Poisson}\left(\lambda \times \text{Area}(D)\right)$$

The conclusion holds when we translate the disk to the left by $1/4$. Or any other amoount. Check it!

12.3.1 Poisson Thinning for a General Region

Let D be any sub-region of the strip with finite area as depicted below.

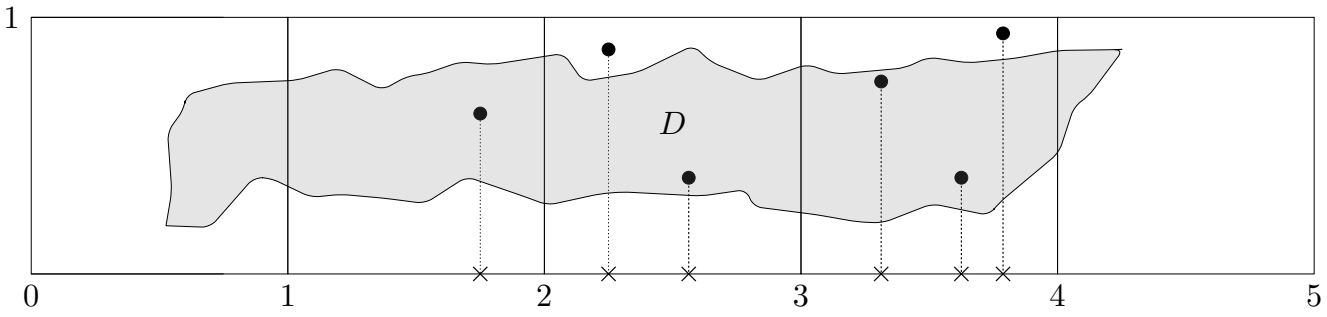

Figure 12.2: PPP on $[0, \infty) \times [0, 1]$ with general region D.

For simplicity, let $m < \infty$. We partition D

$$D = D_1 \cup D_2 \cup \cdots \cup D_m$$

where

$$D_i = D \cap \{(i, i + 1] \times [0, 1]\}$$

and it follows

$$N(D) = \underbrace{N(D_1) + N(D_1) + \ldots + N(D_m)}_{\text{independent Poisson counts}}$$

where by construction and Poisson thining

$$N(D_i) \sim \textbf{Poisson}\,(\lambda \times \text{Area}(D_i))$$

Hence by additing independent Poisson variables, *no matter what the shape of the region D, provided its area is well defined (technically, D is a measurable subset of the strip), the number of Poisson points that fall in D has* **Poisson** $(()\lambda Area(D))$ *distribution*

12.3.2 Poisson Thinning for Two General Regions

Suppose D and F are disjoint regions as depicted below. By Poisson thinning

$$N(D) \sim \textbf{Poisson}\,(\lambda \times \text{Area}(D)) \quad \text{and} \quad N(F) \sim \textbf{Poisson}\,(\lambda \times \text{Area}(F))$$

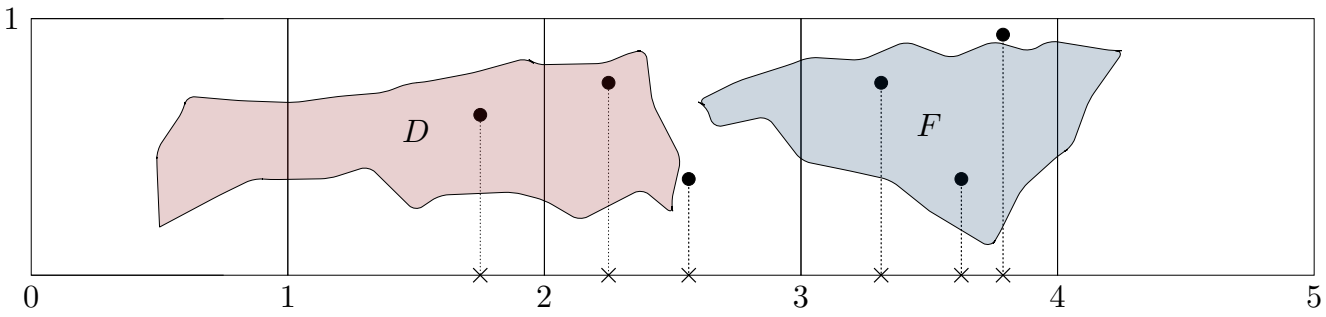

Figure 12.3: PPP on $[0,\infty) \times [0,1]$ with two disjoint regions D and F.

Observe that for each block on the strip we have two independent counts. In total, m independent pairs

$$\begin{aligned} N(D) &= \boxed{N(D_1)} + \boxed{N(D_2)} + \ldots + \boxed{N(D_m)} \\ N(F) &= \boxed{N(F_1)} + \boxed{N(F_2)} + \ldots + \boxed{N(F_m)} \end{aligned}$$

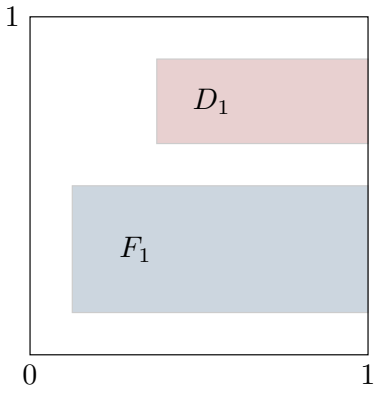

Within each square i there is a count for D_i, a count for F_i and a count for the rest of the square. Hence, by Poissonizaton of the trinomial, the counts D_i and F_i in each square are independent, and they are independent between squares by construction, so we conclude:

- If D and F are disjoint sub-regions of the strip, then the counts $N(D)$ and $N(F)$ are independent Poisson variables with means $\lambda\text{Area}(D)$ and $\lambda\text{Area}(F)$.

The extension of this argument to 3 or more disjoint regions is obvious. The counts in any number of disjoint regions are independent Poisson variables.

12.4 General Measures

A *measure* μ on a space $\mathcal{S}$ is a function of subsets of $\mathcal{S}$. e.g. length, area, volume, probability, subject countable additivity

$$\mu(A) = \mu(A_1) + \mu(A_2) + \dots$$

Define

$$(\mathcal{S}, \mathfrak{S}, \mu) = \text{ measure space}$$

where $\mathfrak{S}$ (curly S) is the domain of measurable sets A, and $\mu(A)$ is a measure. Say a stochastic process

$$\Big(N(A), A \in \mathfrak{S}\Big) \text{ is a PPP with intensity } \mu \text{ iff}$$

- $N(A) \sim \textbf{Poisson}\,(\mu(A))$
- If $A_1, A_2, \dots, A_m$ are disjoint measurable sets, then the $N(A_i)$ are independent.

In the case $\mathcal{S} = [0,\infty) \times [0,1]$ and $\mu = \lambda \times \text{Area}$ we constructed a PPP(μ) by throwing down Poisson numbers of points i.i.d. with probability proportional to area in each of a sequence of unit squares. It was not important in this construction that the squares were unit squares, or even that they were squares. Any way of covering the strip with rectangles would work the same: give each rectangle R a Poisson $\mu(R)$ number of points, and given there are m points in the rectangle, let each of them be uniformly distributed. You could do the same with triangles and get the same Poisson process. More generally, the same argument shows that if μ is any σ-finite measure, meaning

$$\mathcal{S} = \mathcal{S}_1 + \mathcal{S}_2 + \dots \tag{12.4}$$

is the disjoint union of sets $\mathcal{S}_i$ with $\mu(\mathcal{S}_i) < \infty$ for all $i = 1, 2, \dots$, then
then we can construct a PPP(μ), by throwing down independent $\textbf{Poisson}\,(\mu(\mathcal{S}_i))$ numbers of points in $\mathcal{S}_i$, and given n points in $\mathcal{S}_i$, assign them be i.i.d. locations according to $\mu(\,\cdot \,|\, \mathcal{S}_i)$, meaning for $\cdot = A$,

$$\mathbb{P}(\text{point in } A) = \mu(A \,|\, \mathcal{S}_i) := \frac{\mu(A \cap \mathcal{S}_i)}{\mu(\mathcal{S}_i)}.$$

The result will be a PPP(μ), no matter what the choice of partition (12.4) used in this construction. This is a very powerful and general way of thinking about Poisson processes, which is more useful in many respects than the logically equivalent description in the case $\mathcal{S} = [0, \infty)$ and μ is λ times length measure, that the spacings between the points are i.i.d. exponential (λ) variables.

Bibliography

[1] Richard Durrett. *Essentials of Stochastic Processes*. Springer, Reading, Massachusetts, 2012.

[2] Jim Pitman. *Probability*. Springer, Berkeley, California, 1992.

[3] Geoffrey R. Grimmett and David R. Stirzaker. *Probability and Random Processes*. Oxford University Press, New York, third edition, 2001.

[4] Sø ren Asmussen. *Applied Probability and Queues*, volume 51 of *Applications of Mathematics (New York)*. Springer-Verlag, New York, second edition, 2003. Stochastic Modelling and Applied Probability.

[5] J. R. Norris. *Markov Chains*, volume 2 of *Cambridge Series in Statistical and Probabilistic Mathematics*. Cambridge University Press, Cambridge, 1998. Reprint of 1997 original.

[6] William Feller. *An Introduction to Probability Theory and its Applications. Vol. I*. John Wiley and Sons, Inc., New York; Chapman and Hall, Ltd., London, 1957. 2nd ed.

[7] Alvaro Corral and Francesc Font-Clos. Criticality and self-organization in branching processes: application to natural hazards. 07 2012.